Question #799

A company needs to extract the names of ingredients from recipe records that are stored as text files in an Amazon S3 bucket. A web application will use the ingredient names to query an Amazon DynamoDB table and determine a nutrition score.

The application can handle non-food records and errors. The company does not have any employees who have machine learning knowledge to develop this solution.

Which solution will meet these requirements MOST cost-effectively?

A. Use S3 Event Notifications to invoke an AWS Lambda function when PutObject requests occur. Program the Lambda function to analyze the object and extract the ingredient names by using Amazon Comprehend. Store the Amazon Comprehend output in the DynamoDB table.

B. Use an Amazon EventBridge rule to invoke an AWS Lambda function when PutObject requests occur. Program the Lambda function to analyze the object by using Amazon Forecast to extract the ingredient names. Store the Forecast output in the DynamoDB table.

C. Use S3 Event Notifications to invoke an AWS Lambda function when PutObject requests occur. Use Amazon Polly to create audio recordings of the recipe records. Save the audio files in the S3 bucket. Use Amazon Simple Notification Service (Amazon SNS) to send a URL as a message to employees. Instruct the employees to listen to the audio files and calculate the nutrition score. Store the ingredient names in the DynamoDB table.

D. Use an Amazon EventBridge rule to invoke an AWS Lambda function when a PutObject request occurs. Program the Lambda function to analyze the object and extract the ingredient names by using Amazon SageMaker. Store the inference output from the SageMaker endpoint in the DynamoDB table.

Correct Answer: *A*

Community vote distribution

A (100%)

seetpt 2 weeks, 3 days ago

Selected Answer: A

A correct

upvoted 2 times

seetpt 2 weeks, 3 days ago

Selected Answer: A

A is correct

upvoted 2 times

asdfcdsxdfc 2 weeks, 6 days ago

shouldn't it be A?

upvoted 2 times

Question #800

A company needs to create an AWS Lambda function that will run in a VPC in the company's primary AWS account. The Lambda function needs to access files that the company stores in an Amazon Elastic File System (Amazon EFS) file system. The EFS file system is located in a secondary AWS account. As the company adds files to the file system, the solution must scale to meet the demand.

Which solution will meet these requirements MOST cost-effectively?

A. Create a new EFS file system in the primary account. Use AWS DataSync to copy the contents of the original EFS file system to the new EFS file system.

B. Create a VPC peering connection between the VPCs that are in the primary account and the secondary account.

C. Create a second Lambda function in the secondary account that has a mount that is configured for the file system. Use the primary account's Lambda function to invoke the secondary account's Lambda function.

D. Move the contents of the file system to a Lambda layer. Configure the Lambda layer's permissions to allow the company's secondary account to use the Lambda layer.

Correct Answer: *B*

Community vote distribution

B (100%)

lenotc 2 weeks, 1 day ago

Selected Answer: B

B -> VPC peering allows the Lambda access secondary account securely and efficiently
A -> redundancy
C -> additional complexity
D -> sharing code libraries

upvoted 3 times

osmk 2 weeks, 1 day ago

Selected Answer: B

https://docs.aws.amazon.com/efs/latest/ug/efs-different-vpc.html

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Shouldn't it be B?

upvoted 1 times

1dd 2 weeks, 2 days ago

I thinks AWS DataSync less costly

upvoted 1 times

rytizzle 1 week, 2 days ago

setting up a peering connection is free. same for data transfer in the same AZ. data sync at the end of the day cost $$$ to move data.

upvoted 1 times

Question #801

A financial company needs to handle highly sensitive data. The company will store the data in an Amazon S3 bucket. The company needs to ensure that the data is encrypted in transit and at rest. The company must manage the encryption keys outside the AWS Cloud.

Which solution will meet these requirements?

A. Encrypt the data in the S3 bucket with server-side encryption (SSE) that uses an AWS Key Management Service (AWS KMS) customer managed key.

B. Encrypt the data in the S3 bucket with server-side encryption (SSE) that uses an AWS Key Management Service (AWS KMS) AWS managed key.

C. Encrypt the data in the S3 bucket with the default server-side encryption (SSE).

D. Encrypt the data at the company's data center before storing the data in the S3 bucket.

Correct Answer: D

Community vote distribution

D (100%)

Mikado211 1 week, 1 day ago

Selected Answer: D

A, B and C need to have the key stored in AWS cloud.
D is correct.

upvoted 1 times

osmk 2 weeks, 1 day ago

Selected Answer: D

Client-side encryption – You encrypt your data client-side and upload the encrypted data to Amazon S3. In this case, you manage the encryption process, encryption keys, and related tools.https://docs.aws.amazon.com/AmazonS3/latest/userguide/UsingClientSideEncryption.html

upvoted 2 times

giovanna_mag 2 weeks, 1 day ago

Selected Answer: D

For me it's D, it's the only one that provides encryption also in transit

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

A looks correct

upvoted 2 times

Question #802

A company wants to run its payment application on AWS. The application receives payment notifications from mobile devices. Payment notifications require a basic validation before they are sent for further processing.

The backend processing application is long running and requires compute and memory to be adjusted. The company does not want to manage the infrastructure.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Amazon Simple Queue Service (Amazon SQS) queue. Integrate the queue with an Amazon EventBridge rule to receive payment notifications from mobile devices. Configure the rule to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Kubernetes Service (Amazon EKS) Anywhere. Create a standalone cluster.

B. Create an Amazon API Gateway API. Integrate the API with an AWS Step Functions state machine to receive payment notifications from mobile devices. Invoke the state machine to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Kubernetes Service (Amazon EKS). Configure an EKS cluster with self-managed nodes.

C. Create an Amazon Simple Queue Service (Amazon SQS) queue. Integrate the queue with an Amazon EventBridge rule to receive payment notifications from mobile devices. Configure the rule to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon EC2 Spot Instances. Configure a Spot Fleet with a default allocation strategy.

D. Create an Amazon API Gateway API. Integrate the API with AWS Lambda to receive payment notifications from mobile devices. Invoke a Lambda function to validate payment notifications and send the notifications to the backend application. Deploy the backend application on Amazon Elastic Container Service (Amazon ECS). Configure Amazon ECS with an AWS Fargate launch type.

Correct Answer: D

Community vote distribution

D (100%)

Mikado211 1 week, 1 day ago

Selected Answer: D

We want to have least overhead and no infrastructure (aka no server).
So no infrastructure == not C
least overhead == ECS better than EKS == not B and not A

Fargate is serverless so D is still valid.

So the answer is D.

upvoted 2 times

seetpt 2 weeks, 3 days ago

Selected Answer: D

D is correct

upvoted 4 times

asdfcdsxdfc 2 weeks, 6 days ago

shouldn't it be D?

upvoted 3 times

Question #803

A solutions architect is designing a user authentication solution for a company. The solution must invoke two-factor authentication for users that log in from inconsistent geographical locations, IP addresses, or devices. The solution must also be able to scale up to accommodate millions of users.

Which solution will meet these requirements?

A. Configure Amazon Cognito user pools for user authentication. Enable the risk-based adaptive authentication feature with multifactor authentication (MFA).

B. Configure Amazon Cognito identity pools for user authentication. Enable multi-factor authentication (MFA).

C. Configure AWS Identity and Access Management (IAM) users for user authentication. Attach an IAM policy that allows the AllowManageOwnUserMFA action.

D. Configure AWS IAM Identity Center (AWS Single Sign-On) authentication for user authentication. Configure the permission sets to require multi-factor authentication (MFA).

Correct Answer: *A*

Community vote distribution

A (100%)

osmk 2 weeks, 1 day ago

Selected Answer: A

With adaptive authentication, you can configure your user pool to require second factor authentication in response to an increased risk level. To add adaptive authentication to your user pool, see Adding advanced security to a user pool.https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-pool-settings-advanced-security.html

upvoted 2 times

lenotc 2 weeks, 1 day ago

Selected Answer: A

A is correct
B is wrong beacause is desing for temporary credentials

upvoted 1 times

giovanna_mag 2 weeks, 1 day ago

Selected Answer: A

I belive it's A

upvoted 1 times

xBUGx 2 weeks, 3 days ago

accommodate millions of users and GEO, IP, etc. I think A

upvoted 2 times

Question #804

A company has an Amazon S3 data lake. The company needs a solution that transforms the data from the data lake and loads the data into a data warehouse every day. The data warehouse must have massively parallel processing (MPP) capabilities.

Data analysts then need to create and train machine learning (ML) models by using SQL commands on the data. The solution must use serverless AWS services wherever possible.

Which solution will meet these requirements?

A. Run a daily Amazon EMR job to transform the data and load the data into Amazon Redshift. Use Amazon Redshift ML to create and train the ML models.

B. Run a daily Amazon EMR job to transform the data and load the data into Amazon Aurora Serverless. Use Amazon Aurora ML to create and train the ML models.

C. Run a daily AWS Glue job to transform the data and load the data into Amazon Redshift Serverless. Use Amazon Redshift ML to create and train the ML models.

D. Run a daily AWS Glue job to transform the data and load the data into Amazon Athena tables. Use Amazon Athena ML to create and train the ML models.

Correct Answer: *C*

Community vote distribution

C (100%)

Mikado211 1 week, 1 day ago

Selected Answer: C

Data warehouse ==> Redshift
Without additional informations both EMR and Glue Jobs can work.
Since the question asks to use serverless as much as possible, Redshift Serverless is a better solution.
C

upvoted 1 times

1dd 2 weeks, 2 days ago

Selected Answer: C

Option C

upvoted 1 times

1dd 2 weeks, 2 days ago

EMR works with big data transfer

upvoted 1 times

1dd 2 weeks, 2 days ago

MPP --> use Redshift so eliminate B,D
As it required Serverless services --> Glue

upvoted 1 times

1dd 2 weeks, 2 days ago

A have no serverless
C is the answer

upvoted 1 times

seetpt 2 weeks, 3 days ago

Selected Answer: C

C is correct

upvoted 2 times

asdfcdsxdfc 2 weeks, 6 days ago

should be C

upvoted 1 times

Question #805

A company runs containers in a Kubernetes environment in the company's local data center. The company wants to use Amazon Elastic Kubernetes Service (Amazon EKS) and other AWS managed services. Data must remain locally in the company's data center and cannot be stored in any remote site or cloud to maintain compliance.

Which solution will meet these requirements?

A. Deploy AWS Local Zones in the company's data center.

B. Use an AWS Snowmobile in the company's data center.

C. Install an AWS Outposts rack in the company's data center.

D. Install an AWS Snowball Edge Storage Optimized node in the data center.

Correct Answer: *C*

Community vote distribution

C (100%)

Mikado211 1 week, 1 day ago

Selected Answer: C

Outpost is a service where AWS has physical servers in your datacenter.
C

upvoted 1 times

seetpt 2 weeks, 3 days ago

Selected Answer: C

C is correct

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Selected Answer: C

C looks correct

upvoted 3 times

Question #806

A social media company has workloads that collect and process data. The workloads store the data in on-premises NFS storage. The data store cannot scale fast enough to meet the company's expanding business needs. The company wants to migrate the current data store to AWS.

Which solution will meet these requirements MOST cost-effectively?

A. Set up an AWS Storage Gateway Volume Gateway. Use an Amazon S3 Lifecycle policy to transition the data to the appropriate storage class.

B. Set up an AWS Storage Gateway Amazon S3 File Gateway. Use an Amazon S3 Lifecycle policy to transition the data to the appropriate storage class.

C. Use the Amazon Elastic File System (Amazon EFS) Standard-Infrequent Access (Standard-IA) storage class. Activate the infrequent access lifecycle policy.

D. Use the Amazon Elastic File System (Amazon EFS) One Zone-Infrequent Access (One Zone-IA) storage class. Activate the infrequent access lifecycle policy.

Correct Answer: *B*

Community vote distribution

B (100%)

alawada 3 days, 1 hour ago

Selected Answer: B

This solution meets the requirements most cost-effectively because it enables the company to migrate its on-premises NFS data store to AWS without changing the existing applications or workflows. AWS Storage Gateway is a hybrid cloud storage service that provides seamless and secure integration between on-premises and AWS storage. Amazon S3 File Gateway is a type of AWS Storage Gateway that provides a file interface to Amazon S3, with local caching for low-latency access. By setting up an Amazon S3 File Gateway, the company can store and retrieve files as objects in Amazon S3 using standard file protocols such as NFS.

upvoted 1 times

alawada 3 days, 1 hour ago

Selected Answer: B

yeah B

upvoted 1 times

seetpt 2 weeks, 3 days ago

Selected Answer: B

I think B too

upvoted 2 times

asdfcdsxdfc 2 weeks, 6 days ago

Selected Answer: B

B looks correct

upvoted 1 times

Question #807

A company uses high concurrency AWS Lambda functions to process a constantly increasing number of messages in a message queue during marketing events. The Lambda functions use CPU intensive code to process the messages. The company wants to reduce the compute costs and to maintain service latency for its customers.

Which solution will meet these requirements?

A. Configure reserved concurrency for the Lambda functions. Decrease the memory allocated to the Lambda functions.

B. Configure reserved concurrency for the Lambda functions. Increase the memory according to AWS Compute Optimizer recommendations.

C. Configure provisioned concurrency for the Lambda functions. Decrease the memory allocated to the Lambda functions.

D. Configure provisioned concurrency for the Lambda functions. Increase the memory according to AWS Compute Optimizer recommendations.

Correct Answer: D

Community vote distribution

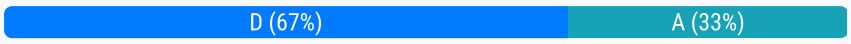

alawada 3 days, 1 hour ago

Selected Answer: D

Provisioned Concurrency keeps the Lambda functions initialized and ready to process incoming events, reducing the cold start latency associated with spinning up new execution environments.

upvoted 1 times

asdfcdsxdfc 2 weeks ago

Selected Answer: D

D is correct

upvoted 1 times

osmk 2 weeks, 1 day ago

Selected Answer: A

When a large number of messages are in the SQS queue, Lambda scales out, adding additional functions to process the messages. The scale out can consume the concurrency quota in the account. To prevent this from happening, you can set reserved concurrency for individual Lambda functions. This ensures that the specified Lambda function can always scale to that much concurrency, but it also cannot exceed this number.
https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html

upvoted 2 times

osmk 2 weeks, 1 day ago

When a large number of messages are in the SQS queue, Lambda scales out, adding additional functions to process the messages. The scale out can consume the concurrency quota in the account. To prevent this from happening, you can set reserved concurrency for individual Lambda functions. This ensures that the specified Lambda function can always scale to that much concurrency, but it also cannot exceed this number.
https://docs.aws.amazon.com/lambda/latest/operatorguide/computing-power.html

upvoted 1 times

Sivaeas 2 weeks, 1 day ago

Selected Answer: D

To reduce compute costs and maintain service latency for customers while using AWS Lambda functions for processing CPU-intensive tasks, you can consider the following strategies:

Optimize Lambda Function Configuration:
Adjust the memory allocation for Lambda functions to better match the CPU requirements of your workload. Higher memory configurations provide more CPU power.
Tune the timeout settings to match the expected processing time of your workload. This prevents unnecessary over-provisioning and reduces costs.
Fine-tune the concurrency settings to control the number of concurrent executions based on your workload's characteristics.
Use Provisioned Concurrency:
AWS Lambda's provisioned concurrency feature allows you to preallocate a number of execution environments to handle incoming requests instantly. This can help reduce cold starts and maintain consistent performance, especially during peak events.

upvoted 2 times

1dd 2 weeks, 2 days ago

Reserved concurrency its no charges reduce the computation cost, "latency for its customer" then I'll go for A

upvoted 1 times

lenotc 2 weeks, 1 day ago

Question #808

A company runs its workloads on Amazon Elastic Container Service (Amazon ECS). The container images that the ECS task definition uses need to be scanned for Common Vulnerabilities and Exposures (CVEs). New container images that are created also need to be scanned.

Which solution will meet these requirements with the FEWEST changes to the workloads?

A. Use Amazon Elastic Container Registry (Amazon ECR) as a private image repository to store the container images. Specify scan on push filters for the ECR basic scan.

B. Store the container images in an Amazon S3 bucket. Use Amazon Macie to scan the images. Use an S3 Event Notification to initiate a Macie scan for every event with an s3:ObjectCreated:Put event type.

C. Deploy the workloads to Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon Elastic Container Registry (Amazon ECR) as a private image repository. Specify scan on push filters for the ECR enhanced scan.

D. Store the container images in an Amazon S3 bucket that has versioning enabled. Configure an S3 Event Notification for s3:ObjectCreated:* events to invoke an AWS Lambda function. Configure the Lambda function to initiate an Amazon Inspector scan.

Correct Answer: *A*

Community vote distribution

A (100%)

1dd 2 weeks, 2 days ago

Selected Answer: A

need less workload changes and CVEs
https://docs.aws.amazon.com/AmazonECR/latest/userguide/image-scanning.html

upvoted 2 times

xBUGx 2 weeks, 3 days ago

Selected Answer: A

FEWEST changes to the workloads and scan CVE is enough. A looks OK.

upvoted 2 times

Question #809

A company uses an AWS Batch job to run its end-of-day sales process. The company needs a serverless solution that will invoke a third-party reporting application when the AWS Batch job is successful. The reporting application has an HTTP API interface that uses username and password authentication.

Which solution will meet these requirements?

A. Configure an Amazon EventBridge rule to match incoming AWS Batch job SUCCEEDED events. Configure the third-party API as an EventBridge API destination with a username and password. Set the API destination as the EventBridge rule target.

B. Configure Amazon EventBridge Scheduler to match incoming AWS Batch job SUCCEEDED events. Configure an AWS Lambda function to invoke the third-party API by using a username and password. Set the Lambda function as the EventBridge rule target.

C. Configure an AWS Batch job to publish job SUCCEEDED events to an Amazon API Gateway REST API. Configure an HTTP proxy integration on the API Gateway REST API to invoke the third-party API by using a username and password.

D. Configure an AWS Batch job to publish job SUCCEEDED events to an Amazon API Gateway REST API. Configure a proxy integration on the API Gateway REST API to an AWS Lambda function. Configure the Lambda function to invoke the third-party API by using a username and password.

Correct Answer: *D*

Community vote distribution

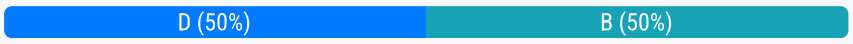

alawada 3 days ago

Selected Answer: D

Create an AWS Lambda function responsible for invoking the third-party reporting application's HTTP API endpoint. The Lambda function will be triggered by the successful completion of the AWS Batch job.

upvoted 1 times

k_k_kkk 6 days, 19 hours ago

Selected Answer: B

AWS Batch sends job status change to EventBridge.

https://docs.aws.amazon.com/batch/latest/userguide/batch_cwe_events.html

upvoted 1 times

osmk 2 weeks, 1 day ago

look like B

upvoted 1 times

Question #810

A company collects and processes data from a vendor. The vendor stores its data in an Amazon RDS for MySQL database in the vendor's own AWS account. The company's VPC does not have an internet gateway, an AWS Direct Connect connection, or an AWS Site-to-Site VPN connection. The company needs to access the data that is in the vendor database.

Which solution will meet this requirement?

A. Instruct the vendor to sign up for the AWS Hosted Connection Direct Connect Program. Use VPC peering to connect the company's VPC and the vendor's VPC.

B. Configure a client VPN connection between the company's VPC and the vendor's VPC. Use VPC peering to connect the company's VPC and the vendor's VPC.

C. Instruct the vendor to create a Network Load Balancer (NLB). Place the NLB in front of the Amazon RDS for MySQL database. Use AWS PrivateLink to integrate the company's VPC and the vendor's VPC.

D. Use AWS Transit Gateway to integrate the company's VPC and the vendor's VPC. Use VPC peering to connect the company's VPC and the vendor's VPC.

Correct Answer: *C*

Community vote distribution

C (80%) A (20%)

xBUGx 5 days, 22 hours ago

D does not involve internet. But TGW is unnecessary.
A is more simple and clear.

upvoted 1 times

Sivaeas 2 weeks, 1 day ago

Selected Answer: C

AWS PrivateLink:
AWS PrivateLink enables you to privately access services hosted on AWS in a highly available and scalable manner. With PrivateLink, you can access the vendor's RDS for MySQL instance securely without exposing it to the public internet.
The vendor can create a VPC endpoint for RDS within their own VPC, which acts as an entry point for accessing the RDS instance. This endpoint can then be shared with the company.
The company can create a VPC endpoint service in their VPC and accept the endpoint connection request from the vendor. This allows the company's resources to communicate with the RDS instance securely through PrivateLink.

upvoted 1 times

lenotc 2 weeks, 1 day ago

Selected Answer: C

C is correct:
https://aws.amazon.com/blogs/networking-and-content-delivery/how-to-securely-publish-internet-applications-at-scale-using-application-load-balancer-and-aws-privatelink/

upvoted 1 times

1dd 2 weeks, 2 days ago

Selected Answer: C

Plz commit the previous comment,
A involve- Direct connect
B involve - peering required same region
D involve - uses internet gateway

upvoted 2 times

1dd 2 weeks, 2 days ago

Selected Answer: A

No internet gateway XD
No Direct connect XC
No Peering XB

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Shouldn't it be D?

upvoted 2 times

1dd 2 weeks, 2 days ago

Question #811

A company wants to set up Amazon Managed Grafana as its visualization tool. The company wants to visualize data from its Amazon RDS database as one data source. The company needs a secure solution that will not expose the data over the internet.

Which solution will meet these requirements?

A. Create an Amazon Managed Grafana workspace without a VPC. Create a public endpoint for the RDS database. Configure the public endpoint as a data source in Amazon Managed Grafana.

B. Create an Amazon Managed Grafana workspace in a VPC. Create a private endpoint for the RDS database. Configure the private endpoint as a data source in Amazon Managed Grafana.

C. Create an Amazon Managed Grafana workspace without a VPCreate an AWS PrivateLink endpoint to establish a connection between Amazon Managed Grafana and Amazon RDS. Set up Amazon RDS as a data source in Amazon Managed Grafana.

D. Create an Amazon Managed Grafana workspace in a VPC. Create a public endpoint for the RDS database. Configure the public endpoint as a data source in Amazon Managed Grafana.

Correct Answer: *B*

Community vote distribution

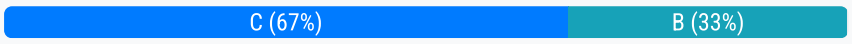

Bazzix 4 days, 14 hours ago

Selected Answer: B

B is correct

upvoted 1 times

osmk 1 week, 6 days ago

Selected Answer: C

cccc ccc

upvoted 2 times

Question #812

A company hosts a data lake on Amazon S3. The data lake ingests data in Apache Parquet format from various data sources. The company uses multiple transformation steps to prepare the ingested data. The steps include filtering of anomalies, normalizing of data to standard date and time values, and generation of aggregates for analyses.

The company must store the transformed data in S3 buckets that data analysts access. The company needs a prebuilt solution for data transformation that does not require code. The solution must provide data lineage and data profiling. The company needs to share the data transformation steps with employees throughout the company.

Which solution will meet these requirements?

A. Configure an AWS Glue Studio visual canvas to transform the data. Share the transformation steps with employees by using AWS Glue jobs.

B. Configure Amazon EMR Serverless to transform the data. Share the transformation steps with employees by using EMR Serverless jobs.

C. Configure AWS Glue DataBrew to transform the data. Share the transformation steps with employees by using DataBrew recipes.

D. Create Amazon Athena tables for the data. Write Athena SQL queries to transform the data. Share the Athena SQL queries with employees.

Correct Answer: *C*

Community vote distribution

C (100%)

seetpt 2 weeks, 3 days ago

Selected Answer: C

Agree with C

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Selected Answer: C

Should be C

upvoted 3 times

Question #813

A solutions architect runs a web application on multiple Amazon EC2 instances that are in individual target groups behind an Application Load Balancer (ALB). Users can reach the application through a public website.

The solutions architect wants to allow engineers to use a development version of the website to access one specific development EC2 instance to test new features for the application. The solutions architect wants to use an Amazon Route 53 hosted zone to give the engineers access to the development instance. The solution must automatically route to the development instance even if the development instance is replaced.

Which solution will meet these requirements?

A. Create an A Record for the development website that has the value set to the ALB. Create a listener rule on the ALB that forwards requests for the development website to the target group that contains the development instance.

B. Recreate the development instance with a public IP address. Create an A Record for the development website that has the value set to the public IP address of the development instance.

C. Create an A Record for the development website that has the value set to the ALB. Create a listener rule on the ALB to redirect requests for the development website to the public IP address of the development instance.

D. Place all the instances in the same target group. Create an A Record for the development website. Set the value to the ALB. Create a listener rule on the ALB that forwards requests for the development website to the target group.

Correct Answer: *A*

Community vote distribution

A (100%)

Mikado211 1 week, 1 day ago

Both A and C look correct but with the C you pass through the ALB to be redirected to a public IP (so go outside) to come back again through this public IP which is not ideal.

The answer A is much cleaner and simplier with a dedicated target group and a listener rule pointing it.

upvoted 1 times

gdf54634 2 weeks ago

Selected Answer: A

Should be A as it points to the target group for easy replacement etc

upvoted 2 times

asdfcdsxdfc 2 weeks ago

Selected Answer: A

I think its A

upvoted 1 times

Question #814

A company runs a container application on a Kubernetes cluster in the company's data center. The application uses Advanced Message Queuing Protocol (AMQP) to communicate with a message queue. The data center cannot scale fast enough to meet the company's expanding business needs. The company wants to migrate the workloads to AWS.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the container application to Amazon Elastic Container Service (Amazon ECS). Use Amazon Simple Queue Service (Amazon SQS) to retrieve the messages.

B. Migrate the container application to Amazon Elastic Kubernetes Service (Amazon EKS). Use Amazon MQ to retrieve the messages.

C. Use highly available Amazon EC2 instances to run the application. Use Amazon MQ to retrieve the messages.

D. Use AWS Lambda functions to run the application. Use Amazon Simple Queue Service (Amazon SQS) to retrieve the messages.

Correct Answer: *B*

Community vote distribution

Mikado211 1 week, 1 day ago

Selected Answer: B

This question is a trap because A is definitely the answer for a Least overhead (ECS + SQS) and in a real life scenario could be good in 99% of cases.

However SQS do not implement AMQP (SQS is only a simple queueing system very basic) so we have to use Amazon MQ.

In terms of containers EKS will always be a better solution than a manual setup of Docker.

Good solution would have been ECS+AmazonMQ not given here

Lambda can work with containers, but since there are limitations like 15 minutes limit we can't really consider it as a good solution.

So B is the least bad solution.

upvoted 3 times

seetpt 2 weeks, 3 days ago

Selected Answer: B

B because only solution with Kubernetes

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Selected Answer: B

Should be B

upvoted 2 times

Question #815

An online gaming company hosts its platform on Amazon EC2 instances behind Network Load Balancers (NLBs) across multiple AWS Regions. The NLBs can route requests to targets over the internet. The company wants to improve the customer playing experience by reducing end-to-end load time for its global customer base.

Which solution will meet these requirements?

A. Create Application Load Balancers (ALBs) in each Region to replace the existing NLBs. Register the existing EC2 instances as targets for the ALBs in each Region.

B. Configure Amazon Route 53 to route equally weighted traffic to the NLBs in each Region.

C. Create additional NLBs and EC2 instances in other Regions where the company has large customer bases.

D. Create a standard accelerator in AWS Global Accelerator. Configure the existing NLBs as target endpoints.

Correct Answer: D

Community vote distribution

Mikado211 1 week, 1 day ago

Selected Answer: D

In such situation if you had an ALB you would use Cloudfront
Since you have a NLB you use AWS Global Accelerator
So D.

upvoted 3 times

seetpt 2 weeks, 3 days ago

Selected Answer: D

Agree with D

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Selected Answer: D

Should be D

upvoted 4 times

Question #816

A company has an on-premises application that uses SFTP to collect financial data from multiple vendors. The company is migrating to the AWS Cloud. The company has created an application that uses Amazon S3 APIs to upload files from vendors.

Some vendors run their systems on legacy applications that do not support S3 APIs. The vendors want to continue to use SFTP-based applications to upload data. The company wants to use managed services for the needs of the vendors that use legacy applications.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an AWS Database Migration Service (AWS DMS) instance to replicate data from the storage of the vendors that use legacy applications to Amazon S3. Provide the vendors with the credentials to access the AWS DMS instance.

B. Create an AWS Transfer Family endpoint for vendors that use legacy applications.

C. Configure an Amazon EC2 instance to run an SFTP server. Instruct the vendors that use legacy applications to use the SFTP server to upload data.

D. Configure an Amazon S3 File Gateway for vendors that use legacy applications to upload files to an SMB file share.

Correct Answer: *B*

Community vote distribution

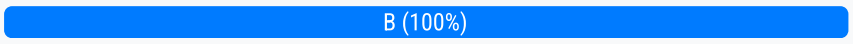

asdfcdsxdfc Highly Voted 2 weeks, 6 days ago

Selected Answer: B

B is correct

upvoted 5 times

seetpt Most Recent 2 weeks, 3 days ago

Selected Answer: B

B is correct

upvoted 1 times

Question #817

A marketing team wants to build a campaign for an upcoming multi-sport event. The team has news reports from the past five years in PDF format. The team needs a solution to extract insights about the content and the sentiment of the news reports. The solution must use Amazon Textract to process the news reports.

Which solution will meet these requirements with the LEAST operational overhead?

A. Provide the extracted insights to Amazon Athena for analysis. Store the extracted insights and analysis in an Amazon S3 bucket.

B. Store the extracted insights in an Amazon DynamoDB table. Use Amazon SageMaker to build a sentiment model.

C. Provide the extracted insights to Amazon Comprehend for analysis. Save the analysis to an Amazon S3 bucket.

D. Store the extracted insights in an Amazon S3 bucket. Use Amazon QuickSight to visualize and analyze the data.

Correct Answer: *C*

Community vote distribution

C (100%)

alawada 2 days, 23 hours ago

Selected Answer: C

Whenever new PDF files are uploaded to the designated S3 bucket, the Lambda function will be triggered to extract insights using Textract and Comprehend.

upvoted 1 times

Mikado211 1 week, 1 day ago

Selected Answer: C

When you have words like "sentiment" in a sentence, it's related to Comprehend
So C.

upvoted 1 times

seetpt 2 weeks, 3 days ago

Selected Answer: C

Maybe C?

upvoted 1 times

asdfcdsxdfc 2 weeks, 6 days ago

Shouldn't it be C?

upvoted 3 times

Question #818

A company's application runs on Amazon EC2 instances that are in multiple Availability Zones. The application needs to ingest real-time data from third-party applications.

The company needs a data ingestion solution that places the ingested raw data in an Amazon S3 bucket.

Which solution will meet these requirements?

A. Create Amazon Kinesis data streams for data ingestion. Create Amazon Kinesis Data Firehose delivery streams to consume the Kinesis data streams. Specify the S3 bucket as the destination of the delivery streams.

B. Create database migration tasks in AWS Database Migration Service (AWS DMS). Specify replication instances of the EC2 instances as the source endpoints. Specify the S3 bucket as the target endpoint. Set the migration type to migrate existing data and replicate ongoing changes.

C. Create and configure AWS DataSync agents on the EC2 instances. Configure DataSync tasks to transfer data from the EC2 instances to the S3 bucket.

D. Create an AWS Direct Connect connection to the application for data ingestion. Create Amazon Kinesis Data Firehose delivery streams to consume direct PUT operations from the application. Specify the S3 bucket as the destination of the delivery streams.

Correct Answer: *A*

Community vote distribution

A (100%)

asdfcdsxdfc Highly Voted 2 weeks, 6 days ago

Selected Answer: A

A is correct

upvoted 5 times

seetpt Most Recent 2 weeks, 3 days ago

Agree with A

upvoted 1 times

Question #819

A company's application is receiving data from multiple data sources. The size of the data varies and is expected to increase over time. The current maximum size is 700 KB. The data volume and data size continue to grow as more data sources are added.

The company decides to use Amazon DynamoDB as the primary database for the application. A solutions architect needs to identify a solution that handles the large data sizes.

Which solution will meet these requirements in the MOST operationally efficient way?

A. Create an AWS Lambda function to filter the data that exceeds DynamoDB item size limits. Store the larger data in an Amazon DocumentDB (with MongoDB compatibility) database.

B. Store the large data as objects in an Amazon S3 bucket. In a DynamoDB table, create an item that has an attribute that points to the S3 URL of the data.

C. Split all incoming large data into a collection of items that have the same partition key. Write the data to a DynamoDB table in a single operation by using the BatchWriteItem API operation.

D. Create an AWS Lambda function that uses gzip compression to compress the large objects as they are written to a DynamoDB table.

Correct Answer: *B*

Community vote distribution

B (100%)

Neung983 Highly Voted 2 weeks, 5 days ago

Selected Answer: B

option B is the most operationally efficient solution for handling large data sizes in Amazon DynamoDB.

upvoted 5 times

seetpt Most Recent 2 weeks, 3 days ago

Selected Answer: B

B is correct

upvoted 4 times

Question #820

A company is migrating a legacy application from an on-premises data center to AWS. The application relies on hundreds of cron jobs that run between 1 and 20 minutes on different recurring schedules throughout the day.

The company wants a solution to schedule and run the cron jobs on AWS with minimal refactoring. The solution must support running the cron jobs in response to an event in the future.

Which solution will meet these requirements?

A. Create a container image for the cron jobs. Use Amazon EventBridge Scheduler to create a recurring schedule. Run the cron job tasks as AWS Lambda functions.

B. Create a container image for the cron jobs. Use AWS Batch on Amazon Elastic Container Service (Amazon ECS) with a scheduling policy to run the cron jobs.

C. Create a container image for the cron jobs. Use Amazon EventBridge Scheduler to create a recurring schedule. Run the cron job tasks on AWS Fargate.

D. Create a container image for the cron jobs. Create a workflow in AWS Step Functions that uses a Wait state to run the cron jobs at a specified time. Use the RunTask action to run the cron job tasks on AWS Fargate.

Correct Answer: *C*

Community vote distribution

C (100%)

Kezuko 5 days, 14 hours ago

Give yourself a pat on the back when you reach this question, its been a long run

upvoted 4 times

Drew3000 4 days, 12 hours ago

I finally managed to get through the last question, then refreshed the page , and they have added more questions.

upvoted 1 times

cvoiceip 2 weeks ago

Ans : C

https://aws.amazon.com/blogs/containers/migrate-cron-jobs-to-event-driven-architectures-using-amazon-elastic-container-service-and-amazon-eventbridge/

upvoted 1 times

seetpt 2 weeks, 3 days ago

Selected Answer: C

C because lambda has 15min time limit.

upvoted 2 times

asdfcdsxdfc 2 weeks, 6 days ago

Selected Answer: C

its either A or C. C looks correct because lambda works for 15 mins and the question says between 1-20

upvoted 3 times

Question #821

A company uses Salesforce. The company needs to load existing data and ongoing data changes from Salesforce to Amazon Redshift for analysis. The company does not want the data to travel over the public internet.

Which solution will meet these requirements with the LEAST development effort?

A. Establish a VPN connection from the VPC to Salesforce. Use AWS Glue DataBrew to transfer data.

B. Establish an AWS Direct Connect connection from the VPC to Salesforce. Use AWS Glue DataBrew to transfer data.

C. Create an AWS PrivateLink connection in the VPC to Salesforce. Use Amazon AppFlow to transfer data.

D. Create a VPC peering connection to Salesforce. Use Amazon AppFlow to transfer data.

Correct Answer: *C*

Kaula 2 days, 11 hours ago

C
https://docs.aws.amazon.com/connect/latest/adminguide/integrate-salesforce-tasks.html
https://docs.aws.amazon.com/connect/latest/adminguide/vpc-interface-endpoints.html
upvoted 1 times

Question #822

A company recently migrated its application to AWS. The application runs on Amazon EC2 Linux instances in an Auto Scaling group across multiple Availability Zones. The application stores data in an Amazon Elastic File System (Amazon EFS) file system that uses EFS Standard-Infrequent Access storage. The application indexes the company's files. The index is stored in an Amazon RDS database.

The company needs to optimize storage costs with some application and services changes.

Which solution will meet these requirements MOST cost-effectively?

A. Create an Amazon S3 bucket that uses an Intelligent-Tiering lifecycle policy. Copy all files to the S3 bucket. Update the application to use Amazon S3 API to store and retrieve files.

B. Deploy Amazon FSx for Windows File Server file shares. Update the application to use CIFS protocol to store and retrieve files.

C. Deploy Amazon FSx for OpenZFS file system shares. Update the application to use the new mount point to store and retrieve files.

D. Create an Amazon S3 bucket that uses S3 Glacier Flexible Retrieval. Copy all files to the S3 bucket. Update the application to use Amazon S3 API to store and retrieve files as standard retrievals.

Correct Answer: *A*

TruthWS 9 hours, 55 minutes ago
A is correct
upvoted 1 times

Kenneth99 1 day, 16 hours ago
should be A?
upvoted 1 times

Question #823

A robotics company is designing a solution for medical surgery. The robots will use advanced sensors, cameras, and AI algorithms to perceive their environment and to complete surgeries.

The company needs a public load balancer in the AWS Cloud that will ensure seamless communication with backend services. The load balancer must be capable of routing traffic based on the query strings to different target groups. The traffic must also be encrypted.

Which solution will meet these requirements?

A. Use a Network Load Balancer with a certificate attached from AWS Certificate Manager (ACM). Use query parameter-based routing.

B. Use a Gateway Load Balancer. Import a generated certificate in AWS Identity and Access Management (IAM). Attach the certificate to the load balancer. Use HTTP path-based routing.

C. Use an Application Load Balancer with a certificate attached from AWS Certificate Manager (ACM). Use query parameter-based routing.

D. Use a Network Load Balancer. Import a generated certificate in AWS Identity and Access Management (IAM). Attach the certificate to the load balancer. Use query parameter-based routing.

Correct Answer: *C*

Community vote distribution

C (100%)

alawada 2 days, 23 hours ago

Selected Answer: C

Provision an Application Load Balancer (ALB) in the AWS Cloud. ALB is a Layer 7 load balancer that supports advanced routing features, including path-based routing.

upvoted 1 times

Question #824

A company has an application that runs on a single Amazon EC2 instance. The application uses a MySQL database that runs on the same EC2 instance. The company needs a highly available and automatically scalable solution to handle increased traffic.

Which solution will meet these requirements?

A. Deploy the application to EC2 instances that run in an Auto Scaling group behind an Application Load Balancer. Create an Amazon Redshift cluster that has multiple MySQL-compatible nodes.

B. Deploy the application to EC2 instances that are configured as a target group behind an Application Load Balancer. Create an Amazon RDS for MySQL cluster that has multiple instances.

C. Deploy the application to EC2 instances that run in an Auto Scaling group behind an Application Load Balancer. Create an Amazon Aurora Serverless MySQL cluster for the database layer.

D. Deploy the application to EC2 instances that are configured as a target group behind an Application Load Balancer. Create an Amazon ElastiCache for Redis cluster that uses the MySQL connector.

Correct Answer: *C*

Community vote distribution

C (100%)

alawada 2 days, 23 hours ago

Selected Answer: C

C Is what I will go for

upvoted 1 times

haci 4 days, 17 hours ago

Selected Answer: C

Target groups are just a group of Ec2 instances. Target groups are closely associated with ELB and not ASG. We can just use ELB and Target groups to route requests to EC2 instances. With this setup, there is no autoscaling which means instances cannot be added or removed when your load increases/decreases.

upvoted 2 times

Question #825

A company is planning to migrate data to an Amazon S3 bucket. The data must be encrypted at rest within the S3 bucket. The encryption key must be rotated automatically every year.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the data to the S3 bucket. Use server-side encryption with Amazon S3 managed keys (SSE-S3). Use the built-in key rotation behavior of SSE-S3 encryption keys.

B. Create an AWS Key Management Service (AWS KMS) customer managed key. Enable automatic key rotation. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Migrate the data to the S3 bucket.

C. Create an AWS Key Management Service (AWS KMS) customer managed key. Set the S3 bucket's default encryption behavior to use the customer managed KMS key. Migrate the data to the S3 bucket. Manually rotate the KMS key every year.

D. Use customer key material to encrypt the data. Migrate the data to the S3 bucket. Create an AWS Key Management Service (AWS KMS) key without key material. Import the customer key material into the KMS key. Enable automatic key rotation.

Correct Answer: B

Community vote distribution

B (100%)

Yushib 2 days, 4 hours ago

Selected Answer: B

B is the right one

upvoted 1 times

haci 4 days, 17 hours ago

Same with Question #202, I'll go with B but not sure

upvoted 1 times

Question #826

A company is migrating applications from an on-premises Microsoft Active Directory that the company manages to AWS. The company deploys the applications in multiple AWS accounts. The company uses AWS Organizations to manage the accounts centrally.

The company's security team needs a single sign-on solution across all the company's AWS accounts. The company must continue to manage users and groups that are in the on-premises Active Directory.

Which solution will meet these requirements?

A. Create an Enterprise Edition Active Directory in AWS Directory Service for Microsoft Active Directory. Configure the Active Directory to be the identity source for AWS IAM Identity Center.

B. Enable AWS IAM Identity Center. Configure a two-way forest trust relationship to connect the company's self-managed Active Directory with IAM Identity Center by using AWS Directory Service for Microsoft Active Directory.

C. Use AWS Directory Service and create a two-way trust relationship with the company's self-managed Active Directory.

D. Deploy an identity provider (IdP) on Amazon EC2. Link the IdP as an identity source within AWS IAM Identity Center.

Correct Answer: *B*

Community vote distribution

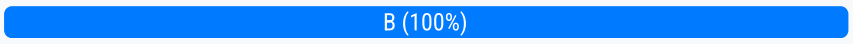

Kaula 2 days, 11 hours ago

Selected Answer: B

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/ms_ad_setup_trust.html

upvoted 1 times

haci 4 days, 17 hours ago

Selected Answer: B

Same with Q-28

upvoted 1 times

Question #827

A company is planning to deploy its application on an Amazon Aurora PostgreSQL Serverless v2 cluster. The application will receive large amounts of traffic. The company wants to optimize the storage performance of the cluster as the load on the application increases.

Which solution will meet these requirements MOST cost-effectively?

A. Configure the cluster to use the Aurora Standard storage configuration.

B. Configure the cluster storage type as Provisioned IOPS.

C. Configure the cluster storage type as General Purpose.

D. Configure the cluster to use the Aurora I/O-Optimized storage configuration.

Correct Answer: *C*

Community vote distribution

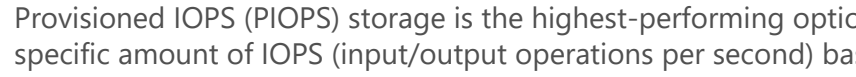

TruthWS 11 hours, 21 minutes ago

I think A is true answer

upvoted 1 times

xBUGx 1 day, 7 hours ago

Selected Answer: D

https://aws.amazon.com/about-aws/whats-new/2023/05/amazon-aurora-i-o-optimized/
Aurora I/O-Optimized offers up to 40% cost savings for I/O-intensive applications where I/O charges exceed 25% of the total Aurora database spend.

upvoted 1 times

Kaula 2 days, 11 hours ago

Selected Answer: C

Agree with haci

upvoted 1 times

haci 4 days, 17 hours ago

Selected Answer: C

The traffic load is not defined well enough to decide which storage type to use.

General Purpose (SSD) storage suits many workloads, including small to medium-sized databases and it is cost-effective.

Provisioned IOPS (PIOPS) storage is the highest-performing option available for RDS instances. With Provisioned IOPS storage, you can provision a specific amount of IOPS (input/output operations per second) based on your application’s needs. But here we don't know the amount of requests.

So since the question is asking for cost-effective I'll go with C

upvoted 1 times

Question #828

A financial services company that runs on AWS has designed its security controls to meet industry standards. The industry standards include the National Institute of Standards and Technology (NIST) and the Payment Card Industry Data Security Standard (PCI DSS).

The company's third-party auditors need proof that the designed controls have been implemented and are functioning correctly. The company has hundreds of AWS accounts in a single organization in AWS Organizations. The company needs to monitor the current state of the controls across accounts.

Which solution will meet these requirements?

A. Designate one account as the Amazon Inspector delegated administrator account from the Organizations management account. Integrate Inspector with Organizations to discover and scan resources across all AWS accounts. Enable Inspector industry standards for NIST and PCI DSS.

B. Designate one account as the Amazon GuardDuty delegated administrator account from the Organizations management account. In the designated GuardDuty administrator account, enable GuardDuty to protect all member accounts. Enable GuardDuty industry standards for NIST and PCI DSS.

C. Configure an AWS CloudTrail organization trail in the Organizations management account. Designate one account as the compliance account. Enable CloudTrail security standards for NIST and PCI DSS in the compliance account.

D. Designate one account as the AWS Security Hub delegated administrator account from the Organizations management account. In the designated Security Hub administrator account, enable Security Hub for all member accounts. Enable Security Hub standards for NIST and PCI DSS.

Correct Answer: *D*

Community vote distribution

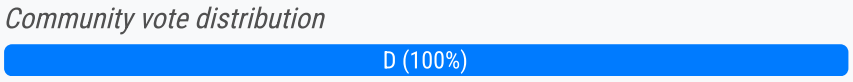

Kaula 2 days, 10 hours ago

Selected Answer: D

https://docs.aws.amazon.com/securityhub/latest/userguide/what-is-securityhub.html

upvoted 1 times

Question #829

A company uses an Amazon S3 bucket as its data lake storage platform. The S3 bucket contains a massive amount of data that is accessed randomly by multiple teams and hundreds of applications. The company wants to reduce the S3 storage costs and provide immediate availability for frequently accessed objects.

What is the MOST operationally efficient solution that meets these requirements?

A. Create an S3 Lifecycle rule to transition objects to the S3 Intelligent-Tiering storage class.

B. Store objects in Amazon S3 Glacier. Use S3 Select to provide applications with access to the data.

C. Use data from S3 storage class analysis to create S3 Lifecycle rules to automatically transition objects to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class.

D. Transition objects to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class. Create an AWS Lambda function to transition objects to the S3 Standard storage class when they are accessed by an application.

Correct Answer: *A*

Community vote distribution

A (100%)

Kaula 2 days, 10 hours ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/intelligent-tiering-managing.html

upvoted 1 times

Question #830

A company has 5 TB of datasets. The datasets consist of 1 million user profiles and 10 million connections. The user profiles have connections as many-to-many relationships. The company needs a performance efficient way to find mutual connections up to five levels.

Which solution will meet these requirements?

A. Use an Amazon S3 bucket to store the datasets. Use Amazon Athena to perform SQL JOIN queries to find connections.

B. Use Amazon Neptune to store the datasets with edges and vertices. Query the data to find connections.

C. Use an Amazon S3 bucket to store the datasets. Use Amazon QuickSight to visualize connections.

D. Use Amazon RDS to store the datasets with multiple tables. Perform SQL JOIN queries to find connections.

Correct Answer: *B*

Community vote distribution

B (100%)

Kaula 2 days, 10 hours ago

Selected Answer: B

https://docs.aws.amazon.com/neptune/latest/userguide/notebooks-visualization.html

upvoted 1 times

alawada 2 days, 23 hours ago

Selected Answer: B

Neptune automatically scales storage and compute resources based on workload demands, ensuring optimal performance even as the dataset grows over time.

upvoted 1 times

Question #831

A company needs a secure connection between its on-premises environment and AWS. This connection does not need high bandwidth and will handle a small amount of traffic. The connection should be set up quickly.

What is the MOST cost-effective method to establish this type of connection?

A. Implement a client VPN.

B. Implement AWS Direct Connect.

C. Implement a bastion host on Amazon EC2.

D. Implement an AWS Site-to-Site VPN connection.

Correct Answer: *D*

Community vote distribution

D (100%)

Kaula 2 days, 10 hours ago

Selected Answer: D

https://docs.aws.amazon.com/vpn/latest/s2svpn/VPC_VPN.html

upvoted 1 times

Question #832

A company has an on-premises SFTP file transfer solution. The company is migrating to the AWS Cloud to scale the file transfer solution and to optimize costs by using Amazon S3. The company's employees will use their credentials for the on-premises Microsoft Active Directory (AD) to access the new solution. The company wants to keep the current authentication and file access mechanisms.

Which solution will meet these requirements with the LEAST operational overhead?

A. Configure an S3 File Gateway. Create SMB file shares on the file gateway that use the existing Active Directory to authenticate.

B. Configure an Auto Scaling group with Amazon EC2 instances to run an SFTP solution. Configure the group to scale up at 60% CPU utilization.

C. Create an AWS Transfer Family server with SFTP endpoints. Choose the AWS Directory Service option as the identity provider. Use AD Connector to connect the on-premises Active Directory.

D. Create an AWS Transfer Family SFTP endpoint. Configure the endpoint to use the AWS Directory Service option as the identity provider to connect to the existing Active Directory.

Correct Answer: *C*

Community vote distribution

C (100%)

Kaula 2 days, 10 hours ago

Selected Answer: C

https://docs.aws.amazon.com/directoryservice/latest/admin-guide/directory_ad_connector.html

upvoted 1 times

Question #833

A company is designing an event-driven order processing system. Each order requires multiple validation steps after the order is created. An idempotent AWS Lambda function performs each validation step. Each validation step is independent from the other validation steps. Individual validation steps need only a subset of the order event information.

The company wants to ensure that each validation step Lambda function has access to only the information from the order event that the function requires. The components of the order processing system should be loosely coupled to accommodate future business changes.

Which solution will meet these requirements?

A. Create an Amazon Simple Queue Service (Amazon SQS) queue for each validation step. Create a new Lambda function to transform the order data to the format that each validation step requires and to publish the messages to the appropriate SQS queues. Subscribe each validation step Lambda function to its corresponding SQS queue.

B. Create an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe the validation step Lambda functions to the SNS topic. Use message body filtering to send only the required data to each subscribed Lambda function.

C. Create an Amazon EventBridge event bus. Create an event rule for each validation step. Configure the input transformer to send only the required data to each target validation step Lambda function.

D. Create an Amazon Simple Queue Service (Amazon SQS) queue. Create a new Lambda function to subscribe to the SQS queue and to transform the order data to the format that each validation step requires. Use the new Lambda function to perform synchronous invocations of the validation step Lambda functions in parallel on separate threads.

Correct Answer: *C*

Community vote distribution

C (100%)

Kaula 2 days, 10 hours ago

Selected Answer: C

https://docs.aws.amazon.com/eventbridge/latest/userguide/eb-event-bus.html

upvoted 1 times

Question #834

A company is migrating a three-tier application to AWS. The application requires a MySQL database. In the past, the application users reported poor application performance when creating new entries. These performance issues were caused by users generating different real-time reports from the application during working hours.

Which solution will improve the performance of the application when it is moved to AWS?

A. Import the data into an Amazon DynamoDB table with provisioned capacity. Refactor the application to use DynamoDB for reports.

B. Create the database on a compute optimized Amazon EC2 instance. Ensure compute resources exceed the on-premises database.

C. Create an Amazon Aurora MySQL Multi-AZ DB cluster with multiple read replicas. Configure the application to use the reader endpoint for reports.

D. Create an Amazon Aurora MySQL Multi-AZ DB cluster. Configure the application to use the backup instance of the cluster as an endpoint for the reports.

Correct Answer: *C*

Question #835

A company is expanding a secure on-premises network to the AWS Cloud by using an AWS Direct Connect connection. The on-premises network has no direct internet access. An application that runs on the on-premises network needs to use an Amazon S3 bucket.

Which solution will meet these requirements MOST cost-effectively?

A. Create a public virtual interface (VIF). Route the AWS traffic over the public VIF.

B. Create a VPC and a NAT gateway. Route the AWS traffic from the on-premises network to the NAT gateway.

C. Create a VPC and an Amazon S3 interface endpoint. Route the AWS traffic from the on-premises network to the S3 interface endpoint.

D. Create a VPC peering connection between the on-premises network and Direct Connect. Route the AWS traffic over the peering connection.

Correct Answer: *C*

Community vote distribution

C (100%)

Kaula 2 days, 10 hours ago

Selected Answer: C

https://docs.aws.amazon.com/AmazonS3/latest/userguide/privatelink-interface-endpoints.html

upvoted 1 times

Question #836

A company serves its website by using an Auto Scaling group of Amazon EC2 instances in a single AWS Region. The website does not require a database.

The company is expanding, and the company's engineering team deploys the website to a second Region. The company wants to distribute traffic across both Regions to accommodate growth and for disaster recovery purposes. The solution should not serve traffic from a Region in which the website is unhealthy.

Which policy or resource should the company use to meet these requirements?

A. An Amazon Route 53 simple routing policy

B. An Amazon Route 53 multivalue answer routing policy

C. An Application Load Balancer in one Region with a target group that specifies the EC2 instance IDs from both Regions

D. An Application Load Balancer in one Region with a target group that specifies the IP addresses of the EC2 instances from both Regions

Correct Answer: *B*

Community vote distribution

B (100%)

Kaula 2 days, 10 hours ago

Selected Answer: B

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-multivalue.html

upvoted 1 times

Question #837

A company runs its applications on Amazon EC2 instances that are backed by Amazon Elastic Block Store (Amazon EBS). The EC2 instances run the most recent Amazon Linux release. The applications are experiencing availability issues when the company's employees store and retrieve files that are 25 GB or larger. The company needs a solution that does not require the company to transfer files between EC2 instances. The files must be available across many EC2 instances and across multiple Availability Zones.

Which solution will meet these requirements?

A. Migrate all the files to an Amazon S3 bucket. Instruct the employees to access the files from the S3 bucket.

B. Take a snapshot of the existing EBS volume. Mount the snapshot as an EBS volume across the EC2 instances. Instruct the employees to access the files from the EC2 instances.

C. Mount an Amazon Elastic File System (Amazon EFS) file system across all the EC2 instances. Instruct the employees to access the files from the EC2 instances.

D. Create an Amazon Machine Image (AMI) from the EC2 instances. Configure new EC2 instances from the AMI that use an instance store volume. Instruct the employees to access the files from the EC2 instances.

Correct Answer: *C*

Question #838

A company is running a highly sensitive application on Amazon EC2 backed by an Amazon RDS database. Compliance regulations mandate that all personally identifiable information (PII) be encrypted at rest.

Which solution should a solutions architect recommend to meet this requirement with the LEAST amount of changes to the infrastructure?

A. Deploy AWS Certificate Manager to generate certificates. Use the certificates to encrypt the database volume.

B. Deploy AWS CloudHSM, generate encryption keys, and use the keys to encrypt database volumes.

C. Configure SSL encryption using AWS Key Management Service (AWS KMS) keys to encrypt database volumes.

D. Configure Amazon Elastic Block Store (Amazon EBS) encryption and Amazon RDS encryption with AWS Key Management Service (AWS KMS) keys to encrypt instance and database volumes.

Correct Answer: *D*

Question #839

A company runs an AWS Lambda function in private subnets in a VPC. The subnets have a default route to the internet through an Amazon EC2 NAT instance. The Lambda function processes input data and saves its output as an object to Amazon S3.

Intermittently, the Lambda function times out while trying to upload the object because of saturated traffic on the NAT instance's network. The company wants to access Amazon S3 without traversing the internet.

Which solution will meet these requirements?

A. Replace the EC2 NAT instance with an AWS managed NAT gateway.

B. Increase the size of the EC2 NAT instance in the VPC to a network optimized instance type.

C. Provision a gateway endpoint for Amazon S3 in the VPUpdate the route tables of the subnets accordingly.

D. Provision a transit gateway. Place transit gateway attachments in the private subnets where the Lambda function is running.

Correct Answer: *C*

Question #840

A news company that has reporters all over the world is hosting its broadcast system on AWS. The reporters send live broadcasts to the broadcast system. The reporters use software on their phones to send live streams through the Real Time Messaging Protocol (RTMP).

A solutions architect must design a solution that gives the reporters the ability to send the highest quality streams. The solution must provide accelerated TCP connections back to the broadcast system.

What should the solutions architect use to meet these requirements?

A. Amazon CloudFront

B. AWS Global Accelerator

C. AWS Client VPN

D. Amazon EC2 instances and AWS Elastic IP addresses

Correct Answer: *B*

Community vote distribution

B (100%)

Kaula 2 days, 9 hours ago

Selected Answer: B

B makes sense not A since CloudFront is CDN

upvoted 1 times

dds69 4 days, 14 hours ago

Selected Answer: B

Global accelerator provides the acceleration for TCP

upvoted 3 times