Question #685

A company's ecommerce website has unpredictable traffic and uses AWS Lambda functions to directly access a private Amazon RDS for PostgreSQL DB instance. The company wants to maintain predictable database performance and ensure that the Lambda invocations do not overload the database with too many connections.

What should a solutions architect do to meet these requirements?

A. Point the client driver at an RDS custom endpoint. Deploy the Lambda functions inside a VPC.

B. Point the client driver at an RDS proxy endpoint. Deploy the Lambda functions inside a VPC.

C. Point the client driver at an RDS custom endpoint. Deploy the Lambda functions outside a VPC.

D. Point the client driver at an RDS proxy endpoint. Deploy the Lambda functions outside a VPC.

Correct Answer: *B*

Community vote distribution

B (100%)

ogerber 4 days, 19 hours ago

Selected Answer: B

Option B.
Reduce number of connection to RDS -> RDS Proxy.
"A Lambda function that's outside of a VPC can't access an RDS instance that's inside a VPC."

https://repost.aws/knowledge-center/connect-lambda-to-an-rds-instance

upvoted 1 times

ogerber 4 days, 19 hours ago

Option B.
Reduce number of connection to RDS -> RDS Proxy.
"A Lambda function that's outside of a VPC can't access an RDS instance that's inside a VPC."

https://repost.aws/knowledge-center/connect-lambda-to-an-rds-instance

upvoted 1 times

Moon239 1 week, 5 days ago

Same as question 802 in SAA-C02

upvoted 2 times

Question #686

A company is creating an application. The company stores data from tests of the application in multiple on-premises locations.

The company needs to connect the on-premises locations to VPCs in an AWS Region in the AWS Cloud. The number of accounts and VPCs will increase during the next year. The network architecture must simplify the administration of new connections and must provide the ability to scale.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Create a peering connection between the VPCs. Create a VPN connection between the VPCs and the on-premises locations.

B. Launch an Amazon EC2 instance. On the instance, include VPN software that uses a VPN connection to connect all VPCs and on-premises locations.

C. Create a transit gateway. Create VPC attachments for the VPC connections. Create VPN attachments for the on-premises connections.

D. Create an AWS Direct Connect connection between the on-premises locations and a central VPC. Connect the central VPC to other VPCs by using peering connections.

Correct Answer: *C*

Community vote distribution

C (100%)

ogerber 4 days, 19 hours ago

Selected Answer: C

high number of accounts and VPC to conect to on prem _> exactly the transit gateway use case

upvoted 1 times

1Alpha1 1 week, 2 days ago

Selected Answer: C

multiple on-premises locations + increasing number of accounts and VPCs --> connections using *transit gateway*

upvoted 2 times

PZ29 1 week, 4 days ago

Hi,
Seems like after question 684, the discussion are quite less and seems recent comments. Are these new sets of questions updated?
Anyone having any idea around this?

upvoted 1 times

Cali182 1 week, 5 days ago

Selected Answer: C

vote for C

upvoted 2 times

EZforeverman 1 week, 6 days ago

I think its C. LEAST administrative overhead. D can work but AWS direct connection and VPC peering configure require too much administrative overhead

upvoted 2 times

Andy_09 2 weeks ago

Think C would be the correct answer here.

upvoted 4 times

Question #687

A company that uses AWS needs a solution to predict the resources needed for manufacturing processes each month. The solution must use historical values that are currently stored in an Amazon S3 bucket. The company has no machine learning (ML) experience and wants to use a managed service for the training and predictions.

Which combination of steps will meet these requirements? (Choose two.)

A. Deploy an Amazon SageMaker model. Create a SageMaker endpoint for inference.

B. Use Amazon SageMaker to train a model by using the historical data in the S3 bucket.

C. Configure an AWS Lambda function with a function URL that uses Amazon SageMaker endpoints to create predictions based on the inputs.

D. Configure an AWS Lambda function with a function URL that uses an Amazon Forecast predictor to create a prediction based on the inputs.

E. Train an Amazon Forsecast predictor by using the historical data in the S3 bucket.

Correct Answer: *DE*

Community vote distribution

DE (67%) BD (33%)

1Alpha1 1 week, 1 day ago

Selected Answer: DE

E: Amazon Forecast is a fully managed service that uses machine learning (ML) to generate highly accurate forecasts without requiring any prior ML experience. Forecast is applicable in a wide variety of use cases, including estimating product demand, energy demand, workforce planning, computing cloud infrastructure usage, traffic demand, supply chain optimization, and financial planning.

D: Publish demand using AWS Lambda, AWS Step Functions, and Amazon CloudWatch Events rule to periodically (hourly) query the database and write the past X-months (count from the current timestamp) demand data into the source Amazon S3.

https://aws.amazon.com/blogs/machine-learning/automating-your-amazon-forecast-workflow-with-lambda-step-functions-and-cloudwatch-events-rule/

upvoted 3 times

Cali182 1 week, 5 days ago

Selected Answer: BD

B & D is the right choice

upvoted 2 times

anikolov 1 week, 6 days ago

Selected Answer: DE

My votes are for DE based on statement from AWS site:
"Alternatively, if you are looking for a fully managed service to deliver highly accurate forecasts, without writing code, we recommend checking out Amazon Forecast. Amazon Forecast is a time-series forecasting service based on machine learning (ML) and built for business metrics analysis."
https://aws.amazon.com/blogs/machine-learning/deep-demand-forecasting-with-amazon-sagemaker/

upvoted 1 times

jaswantn 1 week, 2 days ago

Why E?

upvoted 1 times

betttty 2 weeks ago

Explanation:

Training the Model with SageMaker (Option B):

Use Amazon SageMaker to train a machine learning model based on historical data. SageMaker simplifies the process of training, deploying, and managing machine learning models.
Creating Predictions with Amazon Forecast (Option D):

Use Amazon Forecast to create a predictor based on historical data. Forecast is designed for time-series forecasting, making it suitable for predicting resources needed for manufacturing processes each month.
Combining SageMaker for training and Amazon Forecast for predictions provides a comprehensive solution, and AWS Lambda can be used to integrate these services into your workflow.

upvoted 3 times

Andy_09 2 weeks ago

BE looks correct

upvoted 1 times

Question #688

A company manages AWS accounts in AWS Organizations. AWS IAM Identity Center (AWS Single Sign-On) and AWS Control Tower are configured for the accounts. The company wants to manage multiple user permissions across all the accounts.

The permissions will be used by multiple IAM users and must be split between the developer and administrator teams. Each team requires different permissions. The company wants a solution that includes new users that are hired on both teams.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create individual users in IAM Identity Center for each account. Create separate developer and administrator groups in IAM Identity Center. Assign the users to the appropriate groups. Create a custom IAM policy for each group to set fine-grained permissions.

B. Create individual users in IAM Identity Center for each account. Create separate developer and administrator groups in IAM Identity Center. Assign the users to the appropriate groups. Attach AWS managed IAM policies to each user as needed for fine-grained permissions.

C. Create individual users in IAM Identity Center. Create new developer and administrator groups in IAM Identity Center. Create new permission sets that include the appropriate IAM policies for each group. Assign the new groups to the appropriate accounts. Assign the new permission sets to the new groups. When new users are hired, add them to the appropriate group.

D. Create individual users in IAM Identity Center. Create new permission sets that include the appropriate IAM policies for each user. Assign the users to the appropriate accounts. Grant additional IAM permissions to the users from within specific accounts. When new users are hired, add them to IAM Identity Center and assign them to the accounts.

Correct Answer: C

Community vote distribution

C (100%)

1Alpha1 1 week, 1 day ago

Selected Answer: C

Check out this one. https://www.youtube.com/watch?v=y_n9xN5mg1g

upvoted 1 times

Moon239 1 week, 5 days ago

Selected Answer: C

https://docs.aws.amazon.com/controltower/latest/userguide/sso.html

upvoted 1 times

Cali182 1 week, 5 days ago

Selected Answer: C

Correct is C

upvoted 1 times

Andy_09 2 weeks ago

The correct answer should be C

upvoted 3 times

Question #689

A company wants to standardize its Amazon Elastic Block Store (Amazon EBS) volume encryption strategy. The company also wants to minimize the cost and configuration effort required to operate the volume encryption check.

Which solution will meet these requirements?

A. Write API calls to describe the EBS volumes and to confirm the EBS volumes are encrypted. Use Amazon EventBridge to schedule an AWS Lambda function to run the API calls.

B. Write API calls to describe the EBS volumes and to confirm the EBS volumes are encrypted. Run the API calls on an AWS Fargate task.

C. Create an AWS Identity and Access Management (IAM) policy that requires the use of tags on EBS volumes. Use AWS Cost Explorer to display resources that are not properly tagged. Encrypt the untagged resources manually.

D. Create an AWS Config rule for Amazon EBS to evaluate if a volume is encrypted and to flag the volume if it is not encrypted.

Correct Answer: D

Community vote distribution

D (100%)

mestule 1 week, 6 days ago

Selected Answer: D

AWS Config is a service that enables you to assess, audit, and evaluate the configurations of your AWS resources. It can check whether your resources comply with certain conditions (such as being encrypted), and it can flag or take action on resources that do not comply.

upvoted 3 times

betttty 2 weeks ago

D :
you could use a managed rule to quickly start assessing whether your Amazon Elastic Block Store (Amazon EBS) volumes are encrypted or whether specific tags are applied to your resources.
https://docs.aws.amazon.com/config/latest/developerguide/evaluate-config_use-managed-rules.html

upvoted 2 times

Andy_09 2 weeks ago

Correct answer is D

upvoted 2 times

Question #690

A company regularly uploads GB-sized files to Amazon S3. After the company uploads the files, the company uses a fleet of Amazon EC2 Spot Instances to transcode the file format. The company needs to scale throughput when the company uploads data from the on-premises data center to Amazon S3 and when the company downloads data from Amazon S3 to the EC2 instances.

Which solutions will meet these requirements? (Choose two.)

A. Use the S3 bucket access point instead of accessing the S3 bucket directly.

B. Upload the files into multiple S3 buckets.

C. Use S3 multipart uploads.

D. Fetch multiple byte-ranges of an object in parallel.

E. Add a random prefix to each object when uploading the files.

Correct Answer: *CD*

Community vote distribution

CD (100%)

Darshan07 1 week, 1 day ago

Selected Answer: CD

CD are the correct options

upvoted 1 times

Cali182 1 week, 5 days ago

Selected Answer: CD

CD is the correct for me

upvoted 1 times

betttty 2 weeks ago

CD
C: Increase the file upload throughput
D: increase the file download throughput

upvoted 3 times

Andy_09 2 weeks ago

Correct answer is CD

upvoted 1 times

Question #691

A solutions architect is designing a shared storage solution for a web application that is deployed across multiple Availability Zones. The web application runs on Amazon EC2 instances that are in an Auto Scaling group. The company plans to make frequent changes to the content. The solution must have strong consistency in returning the new content as soon as the changes occur.

Which solutions meet these requirements? (Choose two.)

A. Use AWS Storage Gateway Volume Gateway Internet Small Computer Systems Interface (iSCSI) block storage that is mounted to the individual EC2 instances.

B. Create an Amazon Elastic File System (Amazon EFS) file system. Mount the EFS file system on the individual EC2 instances.

C. Create a shared Amazon Elastic Block Store (Amazon EBS) volume. Mount the EBS volume on the individual EC2 instances.

D. Use AWS DataSync to perform continuous synchronization of data between EC2 hosts in the Auto Scaling group.

E. Create an Amazon S3 bucket to store the web content. Set the metadata for the Cache-Control header to no-cache. Use Amazon CloudFront to deliver the content.

Correct Answer: *BE*

Community vote distribution

BE (100%)

Cali182 1 week, 5 days ago

Selected Answer: BE

B & E seems to be the most logic

upvoted 3 times

Andy_09 2 weeks ago

Correct answer BE

upvoted 4 times

Question #692

A company is deploying an application in three AWS Regions using an Application Load Balancer. Amazon Route 53 will be used to distribute traffic between these Regions.

Which Route 53 configuration should a solutions architect use to provide the MOST high-performing experience?

A. Create an A record with a latency policy.

B. Create an A record with a geolocation policy.

C. Create a CNAME record with a failover policy.

D. Create a CNAME record with a geoproximity policy.

Correct Answer: *A*

Community vote distribution

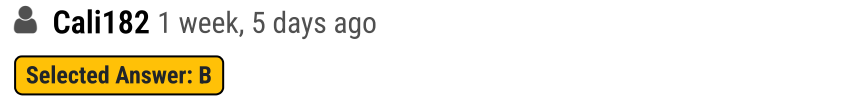

osmk 3 days, 9 hours ago

A
https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-latency.html

upvoted 1 times

haci 3 days, 19 hours ago

Selected Answer: A

Based on previous questions, I believe A is correct. Because; the closest geolocated server doesn't necessarily provide the best performance. Geolocated load balancing is mostly used for serving location-specific content.

upvoted 2 times

1Alpha1 1 week, 1 day ago

Selected Answer: A

Q. What is Amazon Route 53's Latency Based Routing (LBR) feature?

LBR (Latency Based Routing) is a new feature for Amazon Route 53 that helps you improve your application's performance for a global audience. You can run applications in multiple AWS regions and Amazon Route 53, using dozens of edge locations worldwide, will route end users to the AWS region that provides the lowest latency.

https://aws.amazon.com/route53/faqs/

upvoted 1 times

Cali182 1 week, 5 days ago

Selected Answer: B

Why would you use a CNAME record?? Most suitable seems to be option B

upvoted 1 times

osmk 3 days, 9 hours ago

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-latency.html

upvoted 1 times

Andy_09 1 week, 6 days ago

Sorry changing to B.

upvoted 1 times

Andy_09 2 weeks ago

D looks correct.

upvoted 1 times

Question #693

A company has a web application that includes an embedded NoSQL database. The application runs on Amazon EC2 instances behind an Application Load Balancer (ALB). The instances run in an Amazon EC2 Auto Scaling group in a single Availability Zone.

A recent increase in traffic requires the application to be highly available and for the database to be eventually consistent.

Which solution will meet these requirements with the LEAST operational overhead?

A. Replace the ALB with a Network Load Balancer. Maintain the embedded NoSQL database with its replication service on the EC2 instances.

B. Replace the ALB with a Network Load Balancer. Migrate the embedded NoSQL database to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS).

C. Modify the Auto Scaling group to use EC2 instances across three Availability Zones. Maintain the embedded NoSQL database with its replication service on the EC2 instances.

D. Modify the Auto Scaling group to use EC2 instances across three Availability Zones. Migrate the embedded NoSQL database to Amazon DynamoDB by using AWS Database Migration Service (AWS DMS).

Correct Answer: *D*

Community vote distribution

D (100%)

NayeraB 21 hours, 17 minutes ago

But wouldn't migrating an embedded database to a new one introduce operational overhead now and in the future?

upvoted 1 times

1Alpha1 1 week, 1 day ago

Selected Answer: D

DynamoDB + Modifying the Auto Scaling group

upvoted 2 times

Cali182 1 week, 5 days ago

Selected Answer: D

Dynamo DB presents more advantages, because it would need less administrative effort

upvoted 2 times

Andy_09 2 weeks ago

The correct option should be D

upvoted 4 times

Question #694

A company is building a shopping application on AWS. The application offers a catalog that changes once each month and needs to scale with traffic volume. The company wants the lowest possible latency from the application. Data from each user's shopping cart needs to be highly available. User session data must be available even if the user is disconnected and reconnects.

What should a solutions architect do to ensure that the shopping cart data is preserved at all times?

A. Configure an Application Load Balancer to enable the sticky sessions feature (session affinity) for access to the catalog in Amazon Aurora.

B. Configure Amazon ElastiCache for Redis to cache catalog data from Amazon DynamoDB and shopping cart data from the user's session.

C. Configure Amazon OpenSearch Service to cache catalog data from Amazon DynamoDB and shopping cart data from the user's session.

D. Configure an Amazon EC2 instance with Amazon Elastic Block Store (Amazon EBS) storage for the catalog and shopping cart. Configure automated snapshots.

Correct Answer: *B*

Community vote distribution

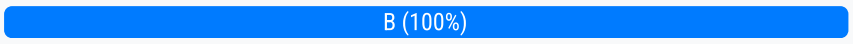

1Alpha1 1 week, 1 day ago

Selected Answer: B

B: ELB <--> ASG <--> ElastiCache <--> DynamoDB

upvoted 1 times

Andy_09 2 weeks ago

B looks correct

upvoted 3 times

Question #695

A company is building a microservices-based application that will be deployed on Amazon Elastic Kubernetes Service (Amazon EKS). The microservices will interact with each other. The company wants to ensure that the application is observable to identify performance issues in the future.

Which solution will meet these requirements?

A. Configure the application to use Amazon ElastiCache to reduce the number of requests that are sent to the microservices.

B. Configure Amazon CloudWatch Container Insights to collect metrics from the EKS clusters. Configure AWS X-Ray to trace the requests between the microservices.

C. Configure AWS CloudTrail to review the API calls. Build an Amazon QuickSight dashboard to observe the microservice interactions.

D. Use AWS Trusted Advisor to understand the performance of the application.

Correct Answer: *B*

Community vote distribution

B (100%)

Cali182 1 week, 5 days ago

Selected Answer: B

Option B
Amazon CloudWatch Container Insights: This service provides monitoring and troubleshooting capabilities for containerized applications. It collects and aggregates metrics, logs, and events from Amazon EKS clusters and containers. This helps in monitoring the performance and health of microservices.

upvoted 2 times

Andy_09 2 weeks ago

Correct answer is B

upvoted 3 times

Question #696

A company needs to provide customers with secure access to its data. The company processes customer data and stores the results in an Amazon S3 bucket.

All the data is subject to strong regulations and security requirements. The data must be encrypted at rest. Each customer must be able to access only their data from their AWS account. Company employees must not be able to access the data.

Which solution will meet these requirements?

A. Provision an AWS Certificate Manager (ACM) certificate for each customer. Encrypt the data client-side. In the private certificate policy, deny access to the certificate for all principals except an IAM role that the customer provides.

B. Provision a separate AWS Key Management Service (AWS KMS) key for each customer. Encrypt the data server-side. In the S3 bucket policy, deny decryption of data for all principals except an IAM role that the customer provides.

C. Provision a separate AWS Key Management Service (AWS KMS) key for each customer. Encrypt the data server-side. In each KMS key policy, deny decryption of data for all principals except an IAM role that the customer provides.

D. Provision an AWS Certificate Manager (ACM) certificate for each customer. Encrypt the data client-side. In the public certificate policy, deny access to the certificate for all principals except an IAM role that the customer provides.

Correct Answer: *C*

Community vote distribution

C (100%)

Cali182 1 week, 5 days ago

Selected Answer: C

Option C
From Chapt
Option A is incorrect because using ACM certificates is typically for establishing secure communication over HTTPS and doesn't directly relate to encrypting data at rest in S3.

Option B is incorrect because while it suggests using AWS KMS keys for encryption, it mentions using S3 bucket policies for access control, which would not be appropriate for controlling decryption permissions.

Option D is incorrect because it suggests using ACM certificates for client-side encryption, which is not typically used for encrypting data at rest in S3, and the approach described would not effectively control access to the encrypted data.

upvoted 1 times

Andy_09 2 weeks ago

Correct answer should be C

upvoted 2 times

Question #697

A solutions architect creates a VPC that includes two public subnets and two private subnets. A corporate security mandate requires the solutions architect to launch all Amazon EC2 instances in a private subnet. However, when the solutions architect launches an EC2 instance that runs a web server on ports 80 and 443 in a private subnet, no external internet traffic can connect to the server.

What should the solutions architect do to resolve this issue?

A. Attach the EC2 instance to an Auto Scaling group in a private subnet. Ensure that the DNS record for the website resolves to the Auto Scaling group identifier.

B. Provision an internet-facing Application Load Balancer (ALB) in a public subnet. Add the EC2 instance to the target group that is associated with the ALEnsure that the DNS record for the website resolves to the ALB.

C. Launch a NAT gateway in a private subnet. Update the route table for the private subnets to add a default route to the NAT gateway. Attach a public Elastic IP address to the NAT gateway.

D. Ensure that the security group that is attached to the EC2 instance allows HTTP traffic on port 80 and HTTPS traffic on port 443. Ensure that the DNS record for the website resolves to the public IP address of the EC2 instance.

Correct Answer: *B*

Community vote distribution

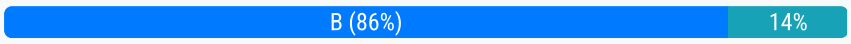

Cali182 1 week, 5 days ago

Selected Answer: C

Option C from Chatgt

upvoted 1 times

jaswantn 1 week, 2 days ago

NAT Gateway stays in public subnet, not in private subnet. So, C can't be.

upvoted 1 times

anikolov 1 week, 6 days ago

Selected Answer: B

B: Provision an internet-facing Application Load Balancer (ALB) in a public subnet makes more sense

upvoted 3 times

mestule 1 week, 6 days ago

Selected Answer: B

B makes most sense

upvoted 3 times

Andy_09 1 week, 6 days ago

Changing to option D

upvoted 1 times

Andy_09 2 weeks ago

C should be the correct answer

upvoted 1 times

Question #698

A company is deploying a new application to Amazon Elastic Kubernetes Service (Amazon EKS) with an AWS Fargate cluster. The application needs a storage solution for data persistence. The solution must be highly available and fault tolerant. The solution also must be shared between multiple application containers.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create Amazon Elastic Block Store (Amazon EBS) volumes in the same Availability Zones where EKS worker nodes are placed. Register the volumes in a StorageClass object on an EKS cluster. Use EBS Multi-Attach to share the data between containers.

B. Create an Amazon Elastic File System (Amazon EFS) file system. Register the file system in a StorageClass object on an EKS cluster. Use the same file system for all containers.

C. Create an Amazon Elastic Block Store (Amazon EBS) volume. Register the volume in a StorageClass object on an EKS cluster. Use the same volume for all containers.

D. Create Amazon Elastic File System (Amazon EFS) file systems in the same Availability Zones where EKS worker nodes are placed. Register the file systems in a StorageClass object on an EKS cluster. Create an AWS Lambda function to synchronize the data between file systems.

Correct Answer: *B*

Community vote distribution

B (100%)

Marunio 6 days, 21 hours ago

Selected Answer: B

B is correct answer because it is high available - EBS isnt HA for that so A isn't dealing with request.

upvoted 1 times

jaswantn 1 week, 2 days ago

Option A... EBS with multi attach does not provide HA so option B is more appropriate.

upvoted 1 times

Andy_09 2 weeks ago

Correct answer is B

upvoted 3 times

Question #699

A company has an application that uses Docker containers in its local data center. The application runs on a container host that stores persistent data in a volume on the host. The container instances use the stored persistent data.

The company wants to move the application to a fully managed service because the company does not want to manage any servers or storage infrastructure.

Which solution will meet these requirements?

A. Use Amazon Elastic Kubernetes Service (Amazon EKS) with self-managed nodes. Create an Amazon Elastic Block Store (Amazon EBS) volume attached to an Amazon EC2 instance. Use the EBS volume as a persistent volume mounted in the containers.

B. Use Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type. Create an Amazon Elastic File System (Amazon EFS) volume. Add the EFS volume as a persistent storage volume mounted in the containers.

C. Use Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type. Create an Amazon S3 bucket. Map the S3 bucket as a persistent storage volume mounted in the containers.

D. Use Amazon Elastic Container Service (Amazon ECS) with an Amazon EC2 launch type. Create an Amazon Elastic File System (Amazon EFS) volume. Add the EFS volume as a persistent storage volume mounted in the containers.

Correct Answer: *B*

Community vote distribution

B (100%)

Marunio 6 days, 21 hours ago

Selected Answer: B

Mounting S3 in Fargate is not supported commonly. You'd have to make it manually. EFS is very well supported with Fargate.
https://stackoverflow.com/questions/66391791/how-to-mount-s3-bucket-to-ecs-fargate-container

https://docs.aws.amazon.com/AmazonECS/latest/bestpracticesguide/storage.html

upvoted 1 times

Andy_09 2 weeks ago

B looks correct

upvoted 4 times

Question #700

A gaming company wants to launch a new internet-facing application in multiple AWS Regions. The application will use the TCP and UDP protocols for communication. The company needs to provide high availability and minimum latency for global users.

Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

A. Create internal Network Load Balancers in front of the application in each Region.

B. Create external Application Load Balancers in front of the application in each Region.

C. Create an AWS Global Accelerator accelerator to route traffic to the load balancers in each Region.

D. Configure Amazon Route 53 to use a geolocation routing policy to distribute the traffic.

E. Configure Amazon CloudFront to handle the traffic and route requests to the application in each Region

Correct Answer: AC

Community vote distribution

AC (100%)

Andy_09 Highly Voted 2 weeks ago

Correct answer should be AC

upvoted 8 times

mestule 1 week, 6 days ago

Agreed.

When you add an internal Load Balancer or an Amazon EC2 instance endpoint in AWS Global Accelerator, you enable internet traffic to flow directly to and from the endpoint in Virtual Private Clouds (VPCs) by targeting it in a private subnet. The VPC that contains the load balancer or EC2 instance must have an internet gateway attached to it, to indicate that the VPC accepts internet traffic. However, you don't need public IP addresses on the load balancer or EC2 instance. You also don't need an associated internet gateway route for the subnet.

upvoted 4 times

ogerber Most Recent 4 days, 19 hours ago

Selected Answer: AC

Gaming + TCP / UDP => always think NLB and global accelerator

upvoted 2 times

1Alpha1 1 week, 1 day ago

Selected Answer: AC

AC - the app is using TCP & UDP

upvoted 2 times

jaswantn 1 week, 2 days ago

For global user where TCP and UDP protocols are used and HA with minimum latency is needed.... Global Accelerator with NLB is the solution combination .

upvoted 2 times

Question #701

A city has deployed a web application running on Amazon EC2 instances behind an Application Load Balancer (ALB). The application's users have reported sporadic performance, which appears to be related to DDoS attacks originating from random IP addresses. The city needs a solution that requires minimal configuration changes and provides an audit trail for the DDoS sources.

Which solution meets these requirements?

A. Enable an AWS WAF web ACL on the ALB, and configure rules to block traffic from unknown sources.

B. Subscribe to Amazon Inspector. Engage the AWS DDoS Response Team (DRT) to integrate mitigating controls into the service.

C. Subscribe to AWS Shield Advanced. Engage the AWS DDoS Response Team (DRT) to integrate mitigating controls into the service.

D. Create an Amazon CloudFront distribution for the application, and set the ALB as the origin. Enable an AWS WAF web ACL on the distribution, and configure rules to block traffic from unknown sources

Correct Answer: *C*

Community vote distribution

C (100%)

Andy_09 Highly Voted 2 weeks ago

C is the correct answer

upvoted 6 times

Darshan07 Most Recent 1 week, 1 day ago

Selected Answer: C

C is the correct answer

upvoted 1 times

Question #702

A company copies 200 TB of data from a recent ocean survey onto AWS Snowball Edge Storage Optimized devices. The company has a high performance computing (HPC) cluster that is hosted on AWS to look for oil and gas deposits. A solutions architect must provide the cluster with consistent sub-millisecond latency and high-throughput access to the data on the Snowball Edge Storage Optimized devices. The company is sending the devices back to AWS.

Which solution will meet these requirements?

A. Create an Amazon S3 bucket. Import the data into the S3 bucket. Configure an AWS Storage Gateway file gateway to use the S3 bucket. Access the file gateway from the HPC cluster instances.

B. Create an Amazon S3 bucket. Import the data into the S3 bucket. Configure an Amazon FSx for Lustre file system, and integrate it with the S3 bucket. Access the FSx for Lustre file system from the HPC cluster instances.

C. Create an Amazon S3 bucket and an Amazon Elastic File System (Amazon EFS) file system. Import the data into the S3 bucket. Copy the data from the S3 bucket to the EFS file system. Access the EFS file system from the HPC cluster instances.

D. Create an Amazon FSx for Lustre file system. Import the data directly into the FSx for Lustre file system. Access the FSx for Lustre file system from the HPC cluster instances.

Correct Answer: *D*

Community vote distribution

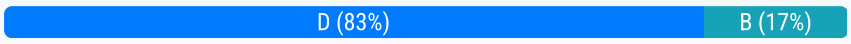

67a3f49 22 hours, 24 minutes ago

Cali182 you cannot directly copy from Snowball Edge to FSx for luster

upvoted 1 times

1Alpha1 2 days, 23 hours ago

Selected Answer: B

Its B.
Snowball Edge (Storage Optimized) --> S3 --integrate--> FSx for Lustre

upvoted 1 times

Darshan07 1 week, 1 day ago

Selected Answer: D

D is the correct answer

upvoted 1 times

Cali182 1 week, 4 days ago

Selected Answer: D

Option D
Option A, B, and C involve using Amazon S3 or Amazon EFS as an intermediary storage layer, which may introduce additional latency and overhead, not meeting the requirement of consistent sub-millisecond latency. Therefore, Option D is the most suitable solution for this scenario.

upvoted 4 times

Andy_09 2 weeks ago

My bad...it should be B

upvoted 4 times

Andy_09 2 weeks ago

Correct answer D

upvoted 1 times

Question #703

A company has NFS servers in an on-premises data center that need to periodically back up small amounts of data to Amazon S3.

Which solution meets these requirements and is MOST cost-effective?

A. Set up AWS Glue to copy the data from the on-premises servers to Amazon S3.

B. Set up an AWS DataSync agent on the on-premises servers, and sync the data to Amazon S3.

C. Set up an SFTP sync using AWS Transfer for SFTP to sync data from on premises to Amazon S3.

D. Set up an AWS Direct Connect connection between the on-premises data center and a VPC, and copy the data to Amazon S3.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

B is the correct option

upvoted 8 times

BillaRanga Most Recent 1 week ago

Selected Answer: B

A -> Used for ETL not copying
B -> Works
C -> Works, but overkill for the described scenario of periodic small backups, high cost
D -> Works but it may not be necessary for transferring small amounts of data periodically. High setup cost

upvoted 3 times

Darshan07 1 week, 1 day ago

Selected Answer: B

B is the correct option

upvoted 1 times

Question #704

An online video game company must maintain ultra-low latency for its game servers. The game servers run on Amazon EC2 instances. The company needs a solution that can handle millions of UDP internet traffic requests each second.

Which solution will meet these requirements MOST cost-effectively?

A. Configure an Application Load Balancer with the required protocol and ports for the internet traffic. Specify the EC2 instances as the targets.

B. Configure a Gateway Load Balancer for the internet traffic. Specify the EC2 instances as the targets.

C. Configure a Network Load Balancer with the required protocol and ports for the internet traffic. Specify the EC2 instances as the targets.

D. Launch an identical set of game servers on EC2 instances in separate AWS Regions. Route internet traffic to both sets of EC2 instances.

Correct Answer: *C*

Community vote distribution

C (100%)

Andy_09 Highly Voted 2 weeks ago

UDP needs NLB

upvoted 6 times

osmk Most Recent 6 days, 17 hours ago

C ->https://docs.aws.amazon.com/elasticloadbalancing/latest/network/introduction.html

upvoted 1 times

Marunio 6 days, 21 hours ago

Selected Answer: C

UDP -> NLB.

ALB is for HTTP/HTTPS.

Gateway Load Balancer is for 3rd party virtual appliances like Firewalls etc not the traffic distribution.

https://aws.amazon.com/compare/the-difference-between-the-difference-between-application-network-and-gateway-load-balancing/#:~:text=An%20NLB%20operates%20on%20layer,level%20along%20with%20gateway%20functionality.

upvoted 2 times

Gagg 6 days, 21 hours ago

Selected Answer: C

UDP, should use network load balancer

upvoted 1 times

nj1999 1 week, 4 days ago

C, NLB

upvoted 4 times

Question #705

A company runs a three-tier application in a VPC. The database tier uses an Amazon RDS for MySQL DB instance.

The company plans to migrate the RDS for MySQL DB instance to an Amazon Aurora PostgreSQL DB cluster. The company needs a solution that replicates the data changes that happen during the migration to the new database.

Which combination of steps will meet these requirements? (Choose two.)

A. Use AWS Database Migration Service (AWS DMS) Schema Conversion to transform the database objects.

B. Use AWS Database Migration Service (AWS DMS) Schema Conversion to create an Aurora PostgreSQL read replica on the RDS for MySQL DB instance.

C. Configure an Aurora MySQL read replica for the RDS for MySQL DB instance.

D. Define an AWS Database Migration Service (AWS DMS) task with change data capture (CDC) to migrate the data.

E. Promote the Aurora PostgreSQL read replica to a standalone Aurora PostgreSQL DB cluster when the replica lag is zero.

Correct Answer: *AD*

Community vote distribution

AD (100%)

haci 3 days, 18 hours ago

Selected Answer: AD

It's quite similar with Q.235, based on that discussion A-D makes more sense.

upvoted 1 times

mestule 1 week, 6 days ago

AD makes sense to me, but I am not sure if that's the best answer.

upvoted 3 times

Andy_09 1 week, 4 days ago

Agreed. AD makes more sense !!

upvoted 3 times

Andy_09 2 weeks ago

Correct answer BE

upvoted 4 times

Question #706

A company hosts a database that runs on an Amazon RDS instance that is deployed to multiple Availability Zones. The company periodically runs a script against the database to report new entries that are added to the database. The script that runs against the database negatively affects the performance of a critical application. The company needs to improve application performance with minimal costs.

Which solution will meet these requirements with the LEAST operational overhead?

A. Add functionality to the script to identify the instance that has the fewest active connections. Configure the script to read from that instance to report the total new entries.

B. Create a read replica of the database. Configure the script to query only the read replica to report the total new entries.

C. Instruct the development team to manually export the new entries for the day in the database at the end of each day.

D. Use Amazon ElastiCache to cache the common queries that the script runs against the database.

Correct Answer: *B*

Community vote distribution

B (100%)

Moon239 1 week, 4 days ago

Selected Answer: B

Read replica

upvoted 3 times

mestule 1 week, 6 days ago

Selected Answer: B

B looks correct

upvoted 2 times

Question #707

A company is using an Application Load Balancer (ALB) to present its application to the internet. The company finds abnormal traffic access patterns across the application. A solutions architect needs to improve visibility into the infrastructure to help the company understand these abnormalities better.

What is the MOST operationally efficient solution that meets these requirements?

A. Create a table in Amazon Athena for AWS CloudTrail logs. Create a query for the relevant information.

B. Enable ALB access logging to Amazon S3. Create a table in Amazon Athena, and query the logs.

C. Enable ALB access logging to Amazon S3. Open each file in a text editor, and search each line for the relevant information.

D. Use Amazon EMR on a dedicated Amazon EC2 instance to directly query the ALB to acquire traffic access log information.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

B is the correct answer

upvoted 7 times

c48b4e2 Most Recent 5 days, 5 hours ago

Why there is a "Correct answer" (the green bordered one) at all while most of the time the community thinks (correctly) otherwise?

upvoted 1 times

Marunio 6 days, 21 hours ago

Selected Answer: B

A - Cloudtrail is for API Calls and changes on AWS account.
B - Going for athena in S3. - Correct
C - Manual work
D - Distractor

upvoted 2 times

betttty 1 week ago

why not A?

upvoted 1 times

Question #708

A company wants to use NAT gateways in its AWS environment. The company's Amazon EC2 instances in private subnets must be able to connect to the public internet through the NAT gateways.

Which solution will meet these requirements?

A. Create public NAT gateways in the same private subnets as the EC2 instances.

B. Create private NAT gateways in the same private subnets as the EC2 instances.

C. Create public NAT gateways in public subnets in the same VPCs as the EC2 instances.

D. Create private NAT gateways in public subnets in the same VPCs as the EC2 instances.

Correct Answer: *C*

Community vote distribution

C (100%)

anikolov Highly Voted 1 week, 6 days ago

Selected Answer: C

Should be C: Public NAT GW in Public Subnet to have access to internet. Private NAT GW is used for VPC or on-prem

upvoted 5 times

mestule Most Recent 1 week, 6 days ago

Selected Answer: C

I think the correct is C, because D would require more than just private NAT gateway.

Private – Instances in private subnets can connect to other VPCs or your on-premises network through a private NAT gateway. You can route traffic from the NAT gateway through a transit gateway or a virtual private gateway. You cannot associate an elastic IP address with a private NAT gateway. You can attach an internet gateway to a VPC with a private NAT gateway, but if you route traffic from the private NAT gateway to the internet gateway, the internet gateway drops the traffic.

https://docs.aws.amazon.com/vpc/latest/userguide/vpc-nat-gateway.html

upvoted 3 times

Andy_09 2 weeks ago

Looks correct

upvoted 1 times

Question #709

A company has an organization in AWS Organizations. The company runs Amazon EC2 instances across four AWS accounts in the root organizational unit (OU). There are three nonproduction accounts and one production account. The company wants to prohibit users from launching EC2 instances of a certain size in the nonproduction accounts. The company has created a service control policy (SCP) to deny access to launch instances that use the prohibited types.

Which solutions to deploy the SCP will meet these requirements? (Choose two.)

A. Attach the SCP to the root OU for the organization.

B. Attach the SCP to the three nonproduction Organizations member accounts.

C. Attach the SCP to the Organizations management account.

D. Create an OU for the production account. Attach the SCP to the OU. Move the production member account into the new OU.

E. Create an OU for the required accounts. Attach the SCP to the OU. Move the nonproduction member accounts into the new OU.

Correct Answer: *BE*

Community vote distribution

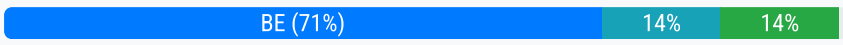

67a3f49 21 hours, 49 minutes ago

According to GPT-4 it's AE:
A. Attach the SCP to the root OU for the organization. This approach will apply the SCP to all accounts under the organization, including both nonproduction and production accounts. However, without additional context or actions, this does not meet the requirement to exclude the production account from the restrictions.

E. Create an OU for the required accounts. Attach the SCP to the OU. Move the nonproduction member accounts into the new OU. This is the correct approach as it directly addresses the requirement. By creating a separate OU for nonproduction accounts and attaching the SCP to this OU, you can specifically target the policy to only those accounts, effectively exempting the production account from the restrictions.

upvoted 1 times

1Alpha1 2 days, 22 hours ago

Selected Answer: AC

AC - same answer

https://docs.aws.amazon.com/organizations/latest/userguide/orgs_getting-started_concepts.html

upvoted 1 times

Cali182 1 week, 4 days ago

Selected Answer: AD

From Chat

A. Attach the SCP to the root OU for the organization: Attaching the SCP to the root OU ensures that it applies to all member accounts within the organization, including both nonproduction and production accounts.

D. Create an OU for the production account. Attach the SCP to the OU. Move the production member account into the new OU: By creating a separate OU for the production account and attaching the SCP to that OU, you can ensure that the SCP only affects the nonproduction accounts while allowing the production account to operate without restrictions.

upvoted 1 times

anikolov 1 week, 6 days ago

Selected Answer: BE

My vote is for BE

upvoted 3 times

mestule 1 week, 6 days ago

Selected Answer: BE

I think it's B (directly attach) and E (attach via OU).

upvoted 2 times

Andy_09 2 weeks ago

CE should be the correct answer

upvoted 1 times

Question #710

A company's website hosted on Amazon EC2 instances processes classified data stored in Amazon S3. Due to security concerns, the company requires a private and secure connection between its EC2 resources and Amazon S3.

Which solution meets these requirements?

A. Set up S3 bucket policies to allow access from a VPC endpoint.

B. Set up an IAM policy to grant read-write access to the S3 bucket.

C. Set up a NAT gateway to access resources outside the private subnet.

D. Set up an access key ID and a secret access key to access the S3 bucket.

Correct Answer: *A*

Community vote distribution

A (100%)

Darshan07 6 days, 23 hours ago

Selected Answer: A

A is the correct answer

upvoted 1 times

Ashy1313 1 week, 3 days ago

Selected Answer: A

A VPC endpoint enables customers to privately connect to supported AWS services .

upvoted 3 times

Question #711

An ecommerce company runs its application on AWS. The application uses an Amazon Aurora PostgreSQL cluster in Multi-AZ mode for the underlying database. During a recent promotional campaign, the application experienced heavy read load and write load. Users experienced timeout issues when they attempted to access the application.

A solutions architect needs to make the application architecture more scalable and highly available.

Which solution will meet these requirements with the LEAST downtime?

A. Create an Amazon EventBridge rule that has the Aurora cluster as a source. Create an AWS Lambda function to log the state change events of the Aurora cluster. Add the Lambda function as a target for the EventBridge rule. Add additional reader nodes to fail over to.

B. Modify the Aurora cluster and activate the zero-downtime restart (ZDR) feature. Use Database Activity Streams on the cluster to track the cluster status.

C. Add additional reader instances to the Aurora cluster. Create an Amazon RDS Proxy target group for the Aurora cluster.

D. Create an Amazon ElastiCache for Redis cache. Replicate data from the Aurora cluster to Redis by using AWS Database Migration Service (AWS DMS) with a write-around approach.

Correct Answer: *C*

Community vote distribution

C (100%)

Marunio 6 days, 21 hours ago

Selected Answer: C

Only C is real viable option - Adding Reader replica for handling Read load and RDS Proxy for connections.

upvoted 2 times

jaswantn 1 week, 2 days ago

RDX proxy to handle timeout issue. option C

upvoted 1 times

Andy_09 2 weeks ago

I would go for option C

upvoted 4 times

Question #712

A company is designing a web application on AWS. The application will use a VPN connection between the company's existing data centers and the company's VPCs.

The company uses Amazon Route 53 as its DNS service. The application must use private DNS records to communicate with the on-premises services from a VPC.

Which solution will meet these requirements in the MOST secure manner?

A. Create a Route 53 Resolver outbound endpoint. Create a resolver rule. Associate the resolver rule with the VPC.

B. Create a Route 53 Resolver inbound endpoint. Create a resolver rule. Associate the resolver rule with the VPC.

C. Create a Route 53 private hosted zone. Associate the private hosted zone with the VPC.

D. Create a Route 53 public hosted zone. Create a record for each service to allow service communication

Correct Answer: *A*

Community vote distribution

A (100%)

haci 3 days, 16 hours ago

Selected Answer: A

If you have workloads that leverage both VPCs and on-premises resources, you also need to resolve DNS records hosted on-premises. Similarly, these on-premises resources may need to resolve names hosted on AWS. Through Resolver endpoints and conditional forwarding rules, you can resolve DNS queries between your on-premises resources and VPCs to create a hybrid cloud setup over VPN or Direct Connect (DX). Specifically:

Inbound Resolver endpoints allow DNS queries to your VPC from your on-premises network or another VPC.

Outbound Resolver endpoints allow DNS queries from your VPC to your on-premises network or another VPC.

Reference: https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/resolver.html

upvoted 2 times

anikolov 1 week, 6 days ago

Selected Answer: A

Should be A "Create a Route 53 Resolver outbound endpoint."

upvoted 4 times

Andy_09 2 weeks ago

Looks correct

upvoted 2 times

Question #713

A company is running a photo hosting service in the us-east-1 Region. The service enables users across multiple countries to upload and view photos. Some photos are heavily viewed for months, and others are viewed for less than a week. The application allows uploads of up to 20 MB for each photo. The service uses the photo metadata to determine which photos to display to each user.

Which solution provides the appropriate user access MOST cost-effectively?

A. Store the photos in Amazon DynamoDB. Turn on DynamoDB Accelerator (DAX) to cache frequently viewed items.

B. Store the photos in the Amazon S3 Intelligent-Tiering storage class. Store the photo metadata and its S3 location in DynamoDB.

C. Store the photos in the Amazon S3 Standard storage class. Set up an S3 Lifecycle policy to move photos older than 30 days to the S3 Standard-Infrequent Access (S3 Standard-IA) storage class. Use the object tags to keep track of metadata.

D. Store the photos in the Amazon S3 Glacier storage class. Set up an S3 Lifecycle policy to move photos older than 30 days to the S3 Glacier Deep Archive storage class. Store the photo metadata and its S3 location in Amazon OpenSearch Service.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

B is the correct option

upvoted 7 times

Typewriter101 Most Recent 5 days, 1 hour ago

Selected Answer: B

The Intelligent-Tiering storage class automatically moves objects between two access tiers (frequent access and infrequent access) based on their access patterns, which aligns well with the varying view frequencies of the photos. Storing metadata in DynamoDB allows for efficient querying and retrieval of photo metadata.

upvoted 1 times

Question #714

A company runs a highly available web application on Amazon EC2 instances behind an Application Load Balancer. The company uses Amazon CloudWatch metrics.

As the traffic to the web application increases, some EC2 instances become overloaded with many outstanding requests. The CloudWatch metrics show that the number of requests processed and the time to receive the responses from some EC2 instances are both higher compared to other EC2 instances. The company does not want new requests to be forwarded to the EC2 instances that are already overloaded.

Which solution will meet these requirements?

A. Use the round robin routing algorithm based on the RequestCountPerTarget and ActiveConnectionCount CloudWatch metrics.

B. Use the least outstanding requests algorithm based on the RequestCountPerTarget and ActiveConnectionCount CloudWatch metrics.

C. Use the round robin routing algorithm based on the RequestCount and TargetResponseTime CloudWatch metrics.

D. Use the least outstanding requests algorithm based on the RequestCount and TargetResponseTime CloudWatch metrics.

Correct Answer: *D*

Andy_09 Highly Voted 2 weeks ago

Option B would be the correct choice

upvoted 5 times

osmk Most Recent 2 days, 15 hours ago

D>>> The least outstanding requests routing algorithm routes requests to the targets with the lowest number of in progress requests > https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html

upvoted 1 times

Moon239 1 week, 4 days ago

Why not D?

upvoted 2 times

jaswantn 1 week, 2 days ago

With Least outstanding requests algorithm, new request will send it to the "target" with least number of outstanding requests. Targets processing long-standing requests or having lower processing capabilities are not burdened with more requests. That's why option B is correct and not option D.

upvoted 2 times

Question #715

A company uses Amazon EC2, AWS Fargate, and AWS Lambda to run multiple workloads in the company's AWS account. The company wants to fully make use of its Compute Savings Plans. The company wants to receive notification when coverage of the Compute Savings Plans drops.

Which solution will meet these requirements with the MOST operational efficiency?

A. Create a daily budget for the Savings Plans by using AWS Budgets. Configure the budget with a coverage threshold to send notifications to the appropriate email message recipients.

B. Create a Lambda function that runs a coverage report against the Savings Plans. Use Amazon Simple Email Service (Amazon SES) to email the report to the appropriate email message recipients.

C. Create an AWS Budgets report for the Savings Plans budget. Set the frequency to daily.

D. Create a Savings Plans alert subscription. Enable all notification options. Enter an email address to receive notifications.

Correct Answer: *A*

Community vote distribution

A (100%)

anikolov Highly Voted 1 week, 6 days ago

Selected Answer: A

My vote is for A : https://docs.aws.amazon.com/savingsplans/latest/userguide/sp-usingBudgets.html

upvoted 6 times

jaswantn Most Recent 1 week, 2 days ago

Option D...In the Savings Plans Overview page indicate how many days in advance you would like to receive Savings Plans Alerts for Plan's expiration and upcoming queued purchase notifications.

upvoted 1 times

Andy_09 2 weeks ago

Option D

upvoted 1 times

hajra313 1 week, 1 day ago

alert subscription will notify u before ending saving plan

upvoted 1 times

Question #716

A company runs a real-time data ingestion solution on AWS. The solution consists of the most recent version of Amazon Managed Streaming for Apache Kafka (Amazon MSK). The solution is deployed in a VPC in private subnets across three Availability Zones.

A solutions architect needs to redesign the data ingestion solution to be publicly available over the internet. The data in transit must also be encrypted.

Which solution will meet these requirements with the MOST operational efficiency?

A. Configure public subnets in the existing VPC. Deploy an MSK cluster in the public subnets. Update the MSK cluster security settings to enable mutual TLS authentication.

B. Create a new VPC that has public subnets. Deploy an MSK cluster in the public subnets. Update the MSK cluster security settings to enable mutual TLS authentication.

C. Deploy an Application Load Balancer (ALB) that uses private subnets. Configure an ALB security group inbound rule to allow inbound traffic from the VPC CIDR block for HTTPS protocol.

D. Deploy a Network Load Balancer (NLB) that uses private subnets. Configure an NLB listener for HTTPS communication over the internet.

Correct Answer: *A*

Community vote distribution

A (100%)

haci 3 days, 15 hours ago

Selected Answer: A

Since we are talking about real-time data (UDP packets) ALB is not a viable solution. You don't need to listen HTTPS, so D is eliminated. If you create a new VPC, you must create link between the old one and this is not mentioned in B. So It is A for me.

upvoted 1 times

Marunio 6 days, 20 hours ago

Selected Answer: A

A, since Kafka is loadbalancing itself. - https://dattell.com/data-architecture-blog/load-balancing-with-kafka/#:~:text=Load%20balancing%20with%20Kafka%20is,partitions%20while%20preserving%20message%20ordering.

B - why create new VPC?

C / D - Kafka is loadbalacing itself, also NLB can't handle HTTPS.

upvoted 1 times

Andy_09 2 weeks ago

Option A

upvoted 3 times

Question #717

A company wants to migrate an on-premises legacy application to AWS. The application ingests customer order files from an on-premises enterprise resource planning (ERP) system. The application then uploads the files to an SFTP server. The application uses a scheduled job that checks for order files every hour.

The company already has an AWS account that has connectivity to the on-premises network. The new application on AWS must support integration with the existing ERP system. The new application must be secure and resilient and must use the SFTP protocol to process orders from the ERP system immediately.

Which solution will meet these requirements?

A. Create an AWS Transfer Family SFTP internet-facing server in two Availability Zones. Use Amazon S3 storage. Create an AWS Lambda function to process order files. Use S3 Event Notifications to send s3:ObjectCreated:* events to the Lambda function.

B. Create an AWS Transfer Family SFTP internet-facing server in one Availability Zone. Use Amazon Elastic File System (Amazon EFS) storage. Create an AWS Lambda function to process order files. Use a Transfer Family managed workflow to invoke the Lambda function.

C. Create an AWS Transfer Family SFTP internal server in two Availability Zones. Use Amazon Elastic File System (Amazon EFS) storage. Create an AWS Step Functions state machine to process order files. Use Amazon EventBridge Scheduler to invoke the state machine to periodically check Amazon EFS for order files.

D. Create an AWS Transfer Family SFTP internal server in two Availability Zones. Use Amazon S3 storage. Create an AWS Lambda function to process order files. Use a Transfer Family managed workflow to invoke the Lambda function.

Correct Answer: *D*

Community vote distribution

D (83%) A (17%)

anikolov Highly Voted 1 week, 6 days ago

Selected Answer: D

D looks more secure over existing on-prem to AWS connection
-Transfer Family SFTP internal server in two Availability Zones.
-Use Amazon S3 storage.
-Use a Transfer Family managed workflow to invoke the Lambda function"

upvoted 5 times

hajra313 1 week, 1 day ago

If the legacy application needs to ingest customer order files from an on-premises ERP system and upload them to an SFTP server, an internet-facing AWS Transfer Family SFTP server would be the appropriate choice.

In this scenario, the SFTP server needs to be accessible from the internet to facilitate the file transfer between the on-premises system and AWS. Therefore, an internet-facing server is required to securely receive the files.

upvoted 1 times

67a3f49 Most Recent 20 hours, 42 minutes ago

I would go in D as it's internal network.

upvoted 1 times

NayeraB 22 hours, 43 minutes ago

Selected Answer: A

I think A makes more sense

upvoted 1 times

Andy_09 2 weeks ago

A is the correct option

upvoted 2 times

Question #718

A company's applications use Apache Hadoop and Apache Spark to process data on premises. The existing infrastructure is not scalable and is complex to manage.

A solutions architect must design a scalable solution that reduces operational complexity. The solution must keep the data processing on premises.

Which solution will meet these requirements?

A. Use AWS Site-to-Site VPN to access the on-premises Hadoop Distributed File System (HDFS) data and application. Use an Amazon EMR cluster to process the data.

B. Use AWS DataSync to connect to the on-premises Hadoop Distributed File System (HDFS) cluster. Create an Amazon EMR cluster to process the data.

C. Migrate the Apache Hadoop application and the Apache Spark application to Amazon EMR clusters on AWS Outposts. Use the EMR clusters to process the data.

D. Use an AWS Snowball device to migrate the data to an Amazon S3 bucket. Create an Amazon EMR cluster to process the data.

Correct Answer: *C*

Community vote distribution

C (100%)

anikolov Highly Voted 1 week, 6 days ago

Selected Answer: C

C cover requirement: The solution must keep the data processing on premises

upvoted 5 times

Andy_09 Most Recent 2 weeks ago

I would go for option C, as data processing has to be done on premise.

upvoted 4 times

Question #719

A company is migrating a large amount of data from on-premises storage to AWS. Windows, Mac, and Linux based Amazon EC2 instances in the same AWS Region will access the data by using SMB and NFS storage protocols. The company will access a portion of the data routinely. The company will access the remaining data infrequently.

The company needs to design a solution to host the data.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Amazon Elastic File System (Amazon EFS) volume that uses EFS Intelligent-Tiering. Use AWS DataSync to migrate the data to the EFS volume.

B. Create an Amazon FSx for ONTAP instance. Create an FSx for ONTAP file system with a root volume that uses the auto tiering policy. Migrate the data to the FSx for ONTAP volume.

C. Create an Amazon S3 bucket that uses S3 Intelligent-Tiering. Migrate the data to the S3 bucket by using an AWS Storage Gateway Amazon S3 File Gateway.

D. Create an Amazon FSx for OpenZFS file system. Migrate the data to the new volume.

Correct Answer: *B*

Community vote distribution

B (67%) C (33%)

Appon 23 hours, 22 minutes ago

Selected Answer: B

option B

upvoted 1 times

ogerber 2 days, 4 hours ago

Selected Answer: B

Amazon FsX for NetAPP ONTAP feature: Multi-protocol access to data using the Network File System (NFS), Server Message Block (SMB), and Internet Small Computer Systems Interface (iSCSI) protocols

upvoted 3 times

MattBJ 1 week ago

Selected Answer: C

C is correct

upvoted 2 times

hajra313 1 week, 1 day ago

Option A and D do not support SMB and NFS file system . Option b looks correvt

upvoted 4 times

jaswantn 1 week, 2 days ago

option C .... SMB and NFS storage protocols ->S3 file gateway

upvoted 2 times

Andy_09 2 weeks ago

Option with S3 usage looks corrcet

upvoted 1 times

Question #720

A manufacturing company runs its report generation application on AWS. The application generates each report in about 20 minutes. The application is built as a monolith that runs on a single Amazon EC2 instance. The application requires frequent updates to its tightly coupled modules. The application becomes complex to maintain as the company adds new features.

Each time the company patches a software module, the application experiences downtime. Report generation must restart from the beginning after any interruptions. The company wants to redesign the application so that the application can be flexible, scalable, and gradually improved. The company wants to minimize application downtime.

Which solution will meet these requirements?

A. Run the application on AWS Lambda as a single function with maximum provisioned concurrency.

B. Run the application on Amazon EC2 Spot Instances as microservices with a Spot Fleet default allocation strategy.

C. Run the application on Amazon Elastic Container Service (Amazon ECS) as microservices with service auto scaling.

D. Run the application on AWS Elastic Beanstalk as a single application environment with an all-at-once deployment strategy.

Correct Answer: *C*

Community vote distribution

C (100%)

Andy_09 Highly Voted 2 weeks ago

Microservices using ECS

upvoted 6 times

Typewriter101 Most Recent 4 days, 23 hours ago

Selected Answer: C

B will not help

spot instances provide cost savings but using it for a stateful task isn't right cause spot instances can be interrupted

upvoted 1 times

Question #721

A company wants to rearchitect a large-scale web application to a serverless microservices architecture. The application uses Amazon EC2 instances and is written in Python.

The company selected one component of the web application to test as a microservice. The component supports hundreds of requests each second. The company wants to create and test the microservice on an AWS solution that supports Python. The solution must also scale automatically and require minimal infrastructure and minimal operational support.

Which solution will meet these requirements?

A. Use a Spot Fleet with auto scaling of EC2 instances that run the most recent Amazon Linux operating system.

B. Use an AWS Elastic Beanstalk web server environment that has high availability configured.

C. Use Amazon Elastic Kubernetes Service (Amazon EKS). Launch Auto Scaling groups of self-managed EC2 instances.

D. Use an AWS Lambda function that runs custom developed code.

Correct Answer: *D*

Community vote distribution

D (100%)

Andy_09 Highly Voted 2 weeks ago

Lambda looks like a better option

upvoted 7 times

Typewriter101 Most Recent 4 days, 23 hours ago

Selected Answer: D

Lambda
serverless, scalable, minimal infrastructure, handling hundreds of requests per second

upvoted 1 times

Umuntu 1 week, 5 days ago

C is the correct answer. The best way to deploy microservice is to use container-based service such as EKS or ECS. So C is great

upvoted 3 times

Typewriter101 4 days, 23 hours ago

Using ECS or EKS involves managing cluster and ec2 which will increase the infrastructure and operational overhead compared to lambda which is serverless.

upvoted 1 times

Andy_09 2 weeks ago

EBS for minimal infra maintenance

upvoted 1 times

Question #722

A company has an AWS Direct Connect connection from its on-premises location to an AWS account. The AWS account has 30 different VPCs in the same AWS Region. The VPCs use private virtual interfaces (VIFs). Each VPC has a CIDR block that does not overlap with other networks under the company's control.

The company wants to centrally manage the networking architecture while still allowing each VPC to communicate with all other VPCs and on-premises networks.

Which solution will meet these requirements with the LEAST amount of operational overhead?

A. Create a transit gateway, and associate the Direct Connect connection with a new transit VIF. Turn on the transit gateway's route propagation feature.

B. Create a Direct Connect gateway. Recreate the private VIFs to use the new gateway. Associate each VPC by creating new virtual private gateways.

C. Create a transit VPConnect the Direct Connect connection to the transit VPCreate a peering connection between all other VPCs in the Region. Update the route tables.

D. Create AWS Site-to-Site VPN connections from on premises to each VPC. Ensure that both VPN tunnels are UP for each connection. Turn on the route propagation feature.

Correct Answer: *A*

Community vote distribution

A (100%)

Andy_09 Highly Voted 2 weeks ago

Option A

upvoted 6 times

Typewriter101 Most Recent 4 days, 23 hours ago

Selected Answer: A

transit gateway -> hub and spoke

upvoted 2 times

Umuntu 1 week, 5 days ago

A is the best solution

upvoted 4 times

Question #723

A company has applications that run on Amazon EC2 instances. The EC2 instances connect to Amazon RDS databases by using an IAM role that has associated policies. The company wants to use AWS Systems Manager to patch the EC2 instances without disrupting the running applications.

Which solution will meet these requirements?

A. Create a new IAM role. Attach the AmazonSSMManagedInstanceCore policy to the new IAM role. Attach the new IAM role to the EC2 instances and the existing IAM role.

B. Create an IAM user. Attach the AmazonSSMManagedInstanceCore policy to the IAM user. Configure Systems Manager to use the IAM user to manage the EC2 instances.

C. Enable Default Host Configuration Management in Systems Manager to manage the EC2 instances.

D. Remove the existing policies from the existing IAM role. Add the AmazonSSMManagedInstanceCore policy to the existing IAM role.

Correct Answer: *C*

NayeraB 1 day ago

So is C same as A, but automated?

upvoted 1 times

osmk 2 days, 14 hours ago

C is fine

upvoted 1 times

jaswantn 1 week, 2 days ago

option C....Default Host Management Configuration creates and applies a default IAM role to ensure that Systems Manager has permissions to manage all instances in the Region and perform automated patch scans using Patch Manager.

upvoted 2 times

Andy_09 1 week, 6 days ago

C is a better option

upvoted 2 times

Andy_09 2 weeks ago

Correct answer A

upvoted 3 times

Question #724

A company runs container applications by using Amazon Elastic Kubernetes Service (Amazon EKS) and the Kubernetes Horizontal Pod Autoscaler. The workload is not consistent throughout the day. A solutions architect notices that the number of nodes does not automatically scale out when the existing nodes have reached maximum capacity in the cluster, which causes performance issues.

Which solution will resolve this issue with the LEAST administrative overhead?

A. Scale out the nodes by tracking the memory usage.

B. Use the Kubernetes Cluster Autoscaler to manage the number of nodes in the cluster.

C. Use an AWS Lambda function to resize the EKS cluster automatically.

D. Use an Amazon EC2 Auto Scaling group to distribute the workload.

Correct Answer: *B*

jaswantn 1 week, 2 days ago

option B.

upvoted 1 times

Andy_09 2 weeks ago

Kubernetes Cluster Autoscaler looks correct

upvoted 3 times

Question #725

A company maintains about 300 TB in Amazon S3 Standard storage month after month. The S3 objects are each typically around 50 GB in size and are frequently replaced with multipart uploads by their global application. The number and size of S3 objects remain constant, but the company's S3 storage costs are increasing each month.

How should a solutions architect reduce costs in this situation?

A. Switch from multipart uploads to Amazon S3 Transfer Acceleration.

B. Enable an S3 Lifecycle policy that deletes incomplete multipart uploads.

C. Configure S3 inventory to prevent objects from being archived too quickly.

D. Configure Amazon CloudFront to reduce the number of objects stored in Amazon S3.

Correct Answer: *B*

Community vote distribution

B (100%)

Typewriter101 4 days, 22 hours ago

Selected Answer: B

when primary concern is cost and the data transfer multipart upload may be the more cost-effective than S3 transfer acceleration. So switching to s3 TA is won't be reasonable.

upvoted 1 times

Umuntu 1 week, 5 days ago

Option B is correct

upvoted 1 times

Andy_09 2 weeks ago

Option B

upvoted 3 times

Question #726

A company has deployed a multiplayer game for mobile devices. The game requires live location tracking of players based on latitude and longitude. The data store for the game must support rapid updates and retrieval of locations.

The game uses an Amazon RDS for PostgreSQL DB instance with read replicas to store the location data. During peak usage periods, the database is unable to maintain the performance that is needed for reading and writing updates. The game's user base is increasing rapidly.

What should a solutions architect do to improve the performance of the data tier?

A. Take a snapshot of the existing DB instance. Restore the snapshot with Multi-AZ enabled.

B. Migrate from Amazon RDS to Amazon OpenSearch Service with OpenSearch Dashboards.

C. Deploy Amazon DynamoDB Accelerator (DAX) in front of the existing DB instance. Modify the game to use DAX.

D. Deploy an Amazon ElastiCache for Redis cluster in front of the existing DB instance. Modify the game to use Redis.

Correct Answer: *D*

Umuntu 1 week, 5 days ago

D looks correct

upvoted 3 times

Andy_09 2 weeks ago

Looks correct

upvoted 1 times

Question #727

A company stores critical data in Amazon DynamoDB tables in the company's AWS account. An IT administrator accidentally deleted a DynamoDB table. The deletion caused a significant loss of data and disrupted the company's operations. The company wants to prevent this type of disruption in the future.

Which solution will meet this requirement with the LEAST operational overhead?

A. Configure a trail in AWS CloudTrail. Create an Amazon EventBridge rule for delete actions. Create an AWS Lambda function to automatically restore deleted DynamoDB tables.

B. Create a backup and restore plan for the DynamoDB tables. Recover the DynamoDB tables manually.

C. Configure deletion protection on the DynamoDB tables.

D. Enable point-in-time recovery on the DynamoDB tables.

Correct Answer: *C*

Community vote distribution

C (100%)

BillaRanga 3 days, 22 hours ago

Selected Answer: C

https://aws.amazon.com/about-aws/whats-new/2023/03/amazon-dynamodb-table-deletion-protection/

Deletion protection is now available for Amazon DynamoDB tables in all AWS Regions. DynamoDB now makes it possible for you to protect your tables from accidental deletion when performing regular table management operations. When creating new tables or managing existing tables, authorized administrators can set the deletion protection property for each table, which will govern whether a table can be deleted.

upvoted 3 times

BillaRanga 3 days, 21 hours ago

Option B and D talks about recovering but not preventing. A is tooooo much work

upvoted 1 times

Typewriter101 4 days, 22 hours ago

Selected Answer: C

B involves more operations.

upvoted 1 times

Andy_09 2 weeks ago

Option C

upvoted 4 times

Question #728

A company has an on-premises data center that is running out of storage capacity. The company wants to migrate its storage infrastructure to AWS while minimizing bandwidth costs. The solution must allow for immediate retrieval of data at no additional cost.

How can these requirements be met?

A. Deploy Amazon S3 Glacier Vault and enable expedited retrieval. Enable provisioned retrieval capacity for the workload.

B. Deploy AWS Storage Gateway using cached volumes. Use Storage Gateway to store data in Amazon S3 while retaining copies of frequently accessed data subsets locally.

C. Deploy AWS Storage Gateway using stored volumes to store data locally. Use Storage Gateway to asynchronously back up point-in-time snapshots of the data to Amazon S3.

D. Deploy AWS Direct Connect to connect with the on-premises data center. Configure AWS Storage Gateway to store data locally. Use Storage Gateway to asynchronously back up point-in-time snapshots of the data to Amazon S3.

Correct Answer: *C*

Community vote distribution

C (100%)

osmk 2 days, 20 hours ago

C>>>
Cached Mode: In this mode, your primary data resides in Amazon S3, while frequently accessed data is cached locally for low-latency access.
Stored Mode: Here, your entire dataset is stored locally, allowing low-latency access on premises. Simultaneously, the data is asynchronously backed up to Amazon S3.

upvoted 1 times

BillaRanga 3 days, 21 hours ago

Selected Answer: C

D -> It takes One month to set up AWS Direct Connect setup
A -> No sense as it talks nothing about On-Prem
B -> Cached volume only stores frequently access data On-Prem, But requirement tells "Data" so we assume it tells All data

C -> Correct, as Stored volumes stores everything in Storage Gateway On-Prem while asynchronously backing up to the cloud

upvoted 1 times

jaswantn 1 week, 2 days ago

option C... data being accessible through stored volume reduces bandwidth cost and provides immediate retrieval of data.

upvoted 1 times

Andy_09 2 weeks ago

Option C, as it makes all the data available for low-latency access.

upvoted 1 times

Question #729

A company runs a three-tier web application in a VPC across multiple Availability Zones. Amazon EC2 instances run in an Auto Scaling group for the application tier.

The company needs to make an automated scaling plan that will analyze each resource's daily and weekly historical workload trends. The configuration must scale resources appropriately according to both the forecast and live changes in utilization.

Which scaling strategy should a solutions architect recommend to meet these requirements?

A. Implement dynamic scaling with step scaling based on average CPU utilization from the EC2 instances.

B. Enable predictive scaling to forecast and scale. Configure dynamic scaling with target tracking

C. Create an automated scheduled scaling action based on the traffic patterns of the web application.

D. Set up a simple scaling policy. Increase the cooldown period based on the EC2 instance startup time.

Correct Answer: B

Community vote distribution

B (100%)

BillaRanga 3 days, 21 hours ago

Selected Answer: B

By configuring dynamic scaling with target tracking, the company can automatically adjust resources based on the forecasted demand while also responding to live changes in utilization

upvoted 1 times

Andy_09 2 weeks ago

Option B

upvoted 4 times

Question #730

A package delivery company has an application that uses Amazon EC2 instances and an Amazon Aurora MySQL DB cluster. As the application becomes more popular, EC2 instance usage increases only slightly. DB cluster usage increases at a much faster rate.

The company adds a read replica, which reduces the DB cluster usage for a short period of time. However, the load continues to increase. The operations that cause the increase in DB cluster usage are all repeated read statements that are related to delivery details. The company needs to alleviate the effect of repeated reads on the DB cluster.

Which solution will meet these requirements MOST cost-effectively?

A. Implement an Amazon ElastiCache for Redis cluster between the application and the DB cluster.

B. Add an additional read replica to the DB cluster.

C. Configure Aurora Auto Scaling for the Aurora read replicas.

D. Modify the DB cluster to have multiple writer instances.

Correct Answer: *A*

Community vote distribution

A (100%)

BillaRanga 3 days, 21 hours ago

Selected Answer: A

The question says, "The operations that cause the increase in DB cluster usage are all **repeated read statements** that are related to delivery details." - Read statements mean we can cache the results - hence, we need No read-replicas; we need only a cache layer to improve the performance.. Also, Adding read replicas costs money. The requirement is to meet them MOST cost-effectively

upvoted 1 times

Andy_09 2 weeks ago

Option A

upvoted 4 times

Question #731

A company has an application that uses an Amazon DynamoDB table for storage. A solutions architect discovers that many requests to the table are not returning the latest data. The company's users have not reported any other issues with database performance. Latency is in an acceptable range.

Which design change should the solutions architect recommend?

A. Add read replicas to the table.

B. Use a global secondary index (GSI).

C. Request strongly consistent reads for the table.

D. Request eventually consistent reads for the table.

Correct Answer: *C*

Community vote distribution

C (100%)

BillaRanga 3 days, 21 hours ago

Selected Answer: C

Both tables and LSIs provide two read consistency options: eventually consistent (default) and strongly consistent reads.
1) Eventually Consistent Reads
Eventually consistent is the default read consistent model for all read operations. When issuing eventually consistent reads to a DynamoDB table or an index, the responses may not reflect the results of a recently completed write operation. If you repeat your read request after a short time, the response should eventually return the more recent item.

upvoted 1 times

BillaRanga 3 days, 21 hours ago

2) Strongly Consistent Reads
Read operations such as GetItem, Query, and Scan provide an optional ConsistentRead parameter. If you set ConsistentRead to true, DynamoDB returns a response with the most up-to-date data, reflecting the updates from all prior write operations that were successful.

Hence it is C
A) Read-replicas are Async again, Which will persist the same problem.
B) Indexing will further cause latency, this has nothing to do with the question

upvoted 1 times

Andy_09 2 weeks ago

Option C

upvoted 3 times

Question #732

A company has deployed its application on Amazon EC2 instances with an Amazon RDS database. The company used the principle of least privilege to configure the database access credentials. The company's security team wants to protect the application and the database from SQL injection and other web-based attacks.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use security groups and network ACLs to secure the database and application servers.

B. Use AWS WAF to protect the application. Use RDS parameter groups to configure the security settings.

C. Use AWS Network Firewall to protect the application and the database.

D. Use different database accounts in the application code for different functions. Avoid granting excessive privileges to the database users.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 6 times

BillaRanga Most Recent 3 days, 21 hours ago

Selected Answer: B

protect the application and the database from SQL injection and other web-based attacks. -> WAF

upvoted 2 times

Typewriter101 4 days, 21 hours ago

Selected Answer: B

SQL injection -> WAF

upvoted 1 times

Question #733

An ecommerce company runs applications in AWS accounts that are part of an organization in AWS Organizations. The applications run on Amazon Aurora PostgreSQL databases across all the accounts. The company needs to prevent malicious activity and must identify abnormal failed and incomplete login attempts to the databases.

Which solution will meet these requirements in the MOST operationally efficient way?

A. Attach service control policies (SCPs) to the root of the organization to identity the failed login attempts.

B. Enable the Amazon RDS Protection feature in Amazon GuardDuty for the member accounts of the organization.

C. Publish the Aurora general logs to a log group in Amazon CloudWatch Logs. Export the log data to a central Amazon S3 bucket.

D. Publish all the Aurora PostgreSQL database events in AWS CloudTrail to a central Amazon S3 bucket.

Correct Answer: *B*

Community vote distribution

B (100%)

BillaRanga 3 days, 21 hours ago

Selected Answer: B

A -> SCPs are not for monitoring or logging

B-> correct
After you enable the RDS Protection feature, GuardDuty immediately starts monitoring RDS login activity from Aurora databases in your account. GuardDuty continuously monitors and profiles RDS login activity for suspicious activity, for example, unauthorized access to Aurora database in your account, from a previously unseen external actor.

upvoted 1 times

Andy_09 2 weeks ago

Option B

upvoted 2 times

Question #734

A company has an AWS Direct Connect connection from its corporate data center to its VPC in the us-east-1 Region. The company recently acquired a corporation that has several VPCs and a Direct Connect connection between its on-premises data center and the eu-west-2 Region. The CIDR blocks for the VPCs of the company and the corporation do not overlap. The company requires connectivity between two Regions and the data centers. The company needs a solution that is scalable while reducing operational overhead.

What should a solutions architect do to meet these requirements?

A. Set up inter-Region VPC peering between the VPC in us-east-1 and the VPCs in eu-west-2.

B. Create private virtual interfaces from the Direct Connect connection in us-east-1 to the VPCs in eu-west-2.

C. Establish VPN appliances in a fully meshed VPN network hosted by Amazon EC2. Use AWS VPN CloudHub to send and receive data between the data centers and each VPC.

D. Connect the existing Direct Connect connection to a Direct Connect gateway. Route traffic from the virtual private gateways of the VPCs in each Region to the Direct Connect gateway.

Correct Answer: *D*

Community vote distribution

D (100%)

BillaRanga 3 days, 21 hours ago

Selected Answer: D

"If you want to set up a Direct Connect to one or more VPC in many different regions (same account), you must use a Direct Connect Gateway."

upvoted 1 times

BillaRanga 3 days, 21 hours ago

CloudHub is a VPN (encrypted) connection, so it goes over the public Internet., Whereas DirectConnect is Private (but not encrypted). So CloudHub is not suited for this useCase

upvoted 1 times

jaswantn 1 week, 2 days ago

option D

upvoted 1 times

Andy_09 1 week, 6 days ago

Changing to Option D for simpler implementation.

upvoted 2 times

Andy_09 2 weeks ago

Option A

upvoted 1 times

Question #735

A company is developing a mobile game that streams score updates to a backend processor and then posts results on a leaderboard. A solutions architect needs to design a solution that can handle large traffic spikes, process the mobile game updates in order of receipt, and store the processed updates in a highly available database. The company also wants to minimize the management overhead required to maintain the solution.

What should the solutions architect do to meet these requirements?

A. Push score updates to Amazon Kinesis Data Streams. Process the updates in Kinesis Data Streams with AWS Lambda. Store the processed updates in Amazon DynamoDB.

B. Push score updates to Amazon Kinesis Data Streams. Process the updates with a fleet of Amazon EC2 instances set up for Auto Scaling. Store the processed updates in Amazon Redshift.

C. Push score updates to an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe an AWS Lambda function to the SNS topic to process the updates. Store the processed updates in a SQL database running on Amazon EC2.

D. Push score updates to an Amazon Simple Queue Service (Amazon SQS) queue. Use a fleet of Amazon EC2 instances with Auto Scaling to process the updates in the SQS queue. Store the processed updates in an Amazon RDS Multi-AZ DB instance.

Correct Answer: *A*

Community vote distribution

A (100%)

Andy_09 Highly Voted 2 weeks ago

Option A

upvoted 5 times

BillaRanga Most Recent 3 days, 21 hours ago

Selected Answer: A

requirement -1: "Stream + process in order + Minimum Overhead" = Kinesis Data Stream + Lambda
requirement-2: "Highly available database + Min Management overhead" = DynamoDb

Setting Up Ec2 instance or MultiAZ DB = overhead

upvoted 2 times

Question #736

A company has multiple AWS accounts with applications deployed in the us-west-2 Region. Application logs are stored within Amazon S3 buckets in each account. The company wants to build a centralized log analysis solution that uses a single S3 bucket. Logs must not leave us-west-2, and the company wants to incur minimal operational overhead.

Which solution meets these requirements and is MOST cost-effective?

A. Create an S3 Lifecycle policy that copies the objects from one of the application S3 buckets to the centralized S3 bucket.

B. Use S3 Same-Region Replication to replicate logs from the S3 buckets to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.

C. Write a script that uses the PutObject API operation every day to copy the entire contents of the buckets to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.

D. Write AWS Lambda functions in these accounts that are triggered every time logs are delivered to the S3 buckets (s3:ObjectCreated:* event). Copy the logs to another S3 bucket in us-west-2. Use this S3 bucket for log analysis.

Correct Answer: *B*

Community vote distribution

B (100%)

BillaRanga 3 days, 20 hours ago

Selected Answer: B

The main Use case of S3 same region replication is "log aggregation, live replication between production and test accounts".

upvoted 2 times

Andy_09 2 weeks ago

Option B

upvoted 4 times

Question #737

A company has an application that delivers on-demand training videos to students around the world. The application also allows authorized content developers to upload videos. The data is stored in an Amazon S3 bucket in the us-east-2 Region.

The company has created an S3 bucket in the eu-west-2 Region and an S3 bucket in the ap-southeast-1 Region. The company wants to replicate the data to the new S3 buckets. The company needs to minimize latency for developers who upload videos and students who stream videos near eu-west-2 and ap-southeast-1.

Which combination of steps will meet these requirements with the FEWEST changes to the application? (Choose two.)

A. Configure one-way replication from the us-east-2 S3 bucket to the eu-west-2 S3 bucket. Configure one-way replication from the us-east-2 S3 bucket to the ap-southeast-1 S3 bucket.

B. Configure one-way replication from the us-east-2 S3 bucket to the eu-west-2 S3 bucket. Configure one-way replication from the eu-west-2 S3 bucket to the ap-southeast-1 S3 bucket.

C. Configure two-way (bidirectional) replication among the S3 buckets that are in all three Regions.

D. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming. Do not modify the application for video uploads.

E. Create an S3 Multi-Region Access Point. Modify the application to use the Amazon Resource Name (ARN) of the Multi-Region Access Point for video streaming and uploads.

Correct Answer: *CE*

Community vote distribution

CE (100%)

BillaRanga 3 days, 20 hours ago

Selected Answer: CE

To keep replication in SYNC across all three regions, we use Bi-directional.
Multi-Region Access Point for video streaming and uploads. -> uploads to nearest Low latency region and Bi-directional replication will keep other two regions in SYNC this reducing the upload and streaming latency

upvoted 1 times

Andy_09 2 weeks ago

Correct answer CE

upvoted 3 times

Question #738

A company has a new mobile app. Anywhere in the world, users can see local news on topics they choose. Users also can post photos and videos from inside the app.

Users access content often in the first minutes after the content is posted. New content quickly replaces older content, and then the older content disappears. The local nature of the news means that users consume 90% of the content within the AWS Region where it is uploaded.

Which solution will optimize the user experience by providing the LOWEST latency for content uploads?

A. Upload and store content in Amazon S3. Use Amazon CloudFront for the uploads.

B. Upload and store content in Amazon S3. Use S3 Transfer Acceleration for the uploads.

C. Upload content to Amazon EC2 instances in the Region that is closest to the user. Copy the data to Amazon S3.

D. Upload and store content in Amazon S3 in the Region that is closest to the user. Use multiple distributions of Amazon CloudFront.

Correct Answer: *B*

Community vote distribution

B (100%)

BillaRanga 3 days, 18 hours ago

Selected Answer: B

Question says - " LOWEST latency for content uploads"
Hence Use S3 Transfer Acceleration for the uploads.

upvoted 3 times

Cali182 1 week, 3 days ago

Selected Answer: B

Cloudfront is for reading not for uploading
Option B

upvoted 3 times

Andy_09 2 weeks ago

Option D

upvoted 1 times

jaswantn 1 week, 2 days ago

option B... S3 transfer acceleration for LOWEST latency for content uploads. question is not asking for low latency for content retrieval. Happy to be corrected

upvoted 1 times

Question #739

A company is building a new application that uses serverless architecture. The architecture will consist of an Amazon API Gateway REST API and AWS Lambda functions to manage incoming requests.

The company wants to add a service that can send messages received from the API Gateway REST API to multiple target Lambda functions for processing. The service must offer message filtering that gives the target Lambda functions the ability to receive only the messages the functions need.

Which solution will meet these requirements with the LEAST operational overhead?

A. Send the requests from the API Gateway REST API to an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe Amazon Simple Queue Service (Amazon SQS) queues to the SNS topic. Configure the target Lambda functions to poll the different SQS queues.

B. Send the requests from the API Gateway REST API to Amazon EventBridge. Configure EventBridge to invoke the target Lambda functions.

C. Send the requests from the API Gateway REST API to Amazon Managed Streaming for Apache Kafka (Amazon MSK). Configure Amazon MSK to publish the messages to the target Lambda functions.

D. Send the requests from the API Gateway REST API to multiple Amazon Simple Queue Service (Amazon SQS) queues. Configure the target Lambda functions to poll the different SQS queues.

Correct Answer: *B*

Community vote distribution

B (100%)

BillaRanga 3 days, 18 hours ago

Selected Answer: B

to multiple target = SNS, EventBridge.

Also, SNS has to use SQS to send filtered content, and Lambda has to poll the SQS to get the message, which is clearly an Overhead. Meanwhile, EventBridge can invoke a Lambda function, which reduces the Operational Overhead.

upvoted 1 times

jaswantn 1 week, 2 days ago

option A.. SNS message filtering

upvoted 2 times

Andy_09 2 weeks ago

Option A

upvoted 1 times

Question #740

A company migrated millions of archival files to Amazon S3. A solutions architect needs to implement a solution that will encrypt all the archival data by using a customer-provided key. The solution must encrypt existing unencrypted objects and future objects.

Which solution will meet these requirements?

A. Create a list of unencrypted objects by filtering an Amazon S3 Inventory report. Configure an S3 Batch Operations job to encrypt the objects from the list with a server-side encryption with a customer-provided key (SSE-C). Configure the S3 default encryption feature to use a server-side encryption with a customer-provided key (SSE-C).

B. Use S3 Storage Lens metrics to identify unencrypted S3 buckets. Configure the S3 default encryption feature to use a server-side encryption with AWS KMS keys (SSE-KMS).

C. Create a list of unencrypted objects by filtering the AWS usage report for Amazon S3. Configure an AWS Batch job to encrypt the objects from the list with a server-side encryption with AWS KMS keys (SSE-KMS). Configure the S3 default encryption feature to use a server-side encryption with AWS KMS keys (SSE-KMS).

D. Create a list of unencrypted objects by filtering the AWS usage report for Amazon S3. Configure the S3 default encryption feature to use a server-side encryption with a customer-provided key (SSE-C).

Correct Answer: *A*

Community vote distribution

A (100%)

BillaRanga 3 days, 17 hours ago

Selected Answer: A

S3 inventory list has "Encryption status" field so you can use this to filter the unencrypted objects. and use S3 batch to encrypt it with SSE-C key.

AWS Usage report does not provide details about encryption status of individual objects

upvoted 1 times

jaswantn 1 week, 2 days ago

option B.... S3 Inventory report to check for unencrypted objects in s3 and then using Batch operation.

upvoted 1 times

OX_HDR 1 week, 5 days ago

Selected Answer: A

A seems correct here.

https://aws.amazon.com/blogs/storage/encrypting-objects-with-amazon-s3-batch-operations/

upvoted 4 times

mestule 1 week, 6 days ago

Selected Answer: A

The solution must encrypt existing unencrypted objects. Batch will do that.

upvoted 3 times

Andy_09 2 weeks ago

Option B

upvoted 1 times

Question #741

The DNS provider that hosts a company's domain name records is experiencing outages that cause service disruption for a website running on AWS. The company needs to migrate to a more resilient managed DNS service and wants the service to run on AWS.

What should a solutions architect do to rapidly migrate the DNS hosting service?

A. Create an Amazon Route 53 public hosted zone for the domain name. Import the zone file containing the domain records hosted by the previous provider.

B. Create an Amazon Route 53 private hosted zone for the domain name. Import the zone file containing the domain records hosted by the previous provider.

C. Create a Simple AD directory in AWS. Enable zone transfer between the DNS provider and AWS Directory Service for Microsoft Active Directory for the domain records.

D. Create an Amazon Route 53 Resolver inbound endpoint in the VPC. Specify the IP addresses that the provider's DNS will forward DNS queries to. Configure the provider's DNS to forward DNS queries for the domain to the IP addresses that are specified in the inbound endpoint.

Correct Answer: *A*

Community vote distribution

A (100%)

Andy_09 Highly Voted 2 weeks ago

Option A

upvoted 6 times

BillaRanga Most Recent 3 days, 17 hours ago

Selected Answer: A

A -> Correct as we need to route to a Company in public network.
B -> No, because it routes only within one or more VPC
C -> Added as a distractor
D -> Inbound resolver is for traffic from On-Prem to VPC

upvoted 1 times

Question #742

A company is building an application on AWS that connects to an Amazon RDS database. The company wants to manage the application configuration and to securely store and retrieve credentials for the database and other services.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Use AWS AppConfig to store and manage the application configuration. Use AWS Secrets Manager to store and retrieve the credentials.

B. Use AWS Lambda to store and manage the application configuration. Use AWS Systems Manager Parameter Store to store and retrieve the credentials.

C. Use an encrypted application configuration file. Store the file in Amazon S3 for the application configuration. Create another S3 file to store and retrieve the credentials.

D. Use AWS AppConfig to store and manage the application configuration. Use Amazon RDS to store and retrieve the credentials.

Correct Answer: *A*

Community vote distribution

A (100%)

Andy_09 Highly Voted 2 weeks ago

Option A

upvoted 6 times

BillaRanga Most Recent 3 days, 17 hours ago

Selected Answer: A

AppConfig useCase = You can use AWS AppConfig to deploy configuration data stored in the AWS AppConfig hosted configuration store, AWS Secrets Manager, Systems Manager Parameter Store, or Amazon S3.
So B and C are out.

use RDS to store credentials is not a good practise. So D is out.

Ans is A

upvoted 2 times

Question #743

To meet security requirements, a company needs to encrypt all of its application data in transit while communicating with an Amazon RDS MySQL DB instance. A recent security audit revealed that encryption at rest is enabled using AWS Key Management Service (AWS KMS), but data in transit is not enabled.

What should a solutions architect do to satisfy the security requirements?

A. Enable IAM database authentication on the database.

B. Provide self-signed certificates. Use the certificates in all connections to the RDS instance.

C. Take a snapshot of the RDS instance. Restore the snapshot to a new instance with encryption enabled.

D. Download AWS-provided root certificates. Provide the certificates in all connections to the RDS instance.

Correct Answer: D

Community vote distribution

D (100%)

BillaRanga 3 days, 17 hours ago

Selected Answer: D

Amazon RDS creates an SSL certificate and installs the certificate on the DB instance when the instance is provisioned. So it is AWS provided.

upvoted 2 times

Andy_09 2 weeks ago

Option D

upvoted 3 times

Question #744

A company is designing a new web service that will run on Amazon EC2 instances behind an Elastic Load Balancing (ELB) load balancer. However, many of the web service clients can only reach IP addresses authorized on their firewalls.

What should a solutions architect recommend to meet the clients' needs?

A. A Network Load Balancer with an associated Elastic IP address.

B. An Application Load Balancer with an associated Elastic IP address.

C. An A record in an Amazon Route 53 hosted zone pointing to an Elastic IP address.

D. An EC2 instance with a public IP address running as a proxy in front of the load balancer.

Correct Answer: *C*

Community vote distribution

C (50%) A (50%)

PolarFox 2 days, 19 hours ago

Selected Answer: C

Option C

upvoted 1 times

BillaRanga 3 days, 17 hours ago

Selected Answer: A

B -> Application Load Balancer cannot be assigned an Elastic IP address (static IP address).
C -> Its DNS after all, "Associated elastic IP" is what IP? Makes no sense
D -> "If you require a persistent public IP address that can be associated to and from instances as you require, use an Elastic IP address instead." PUBLIC IP of an EC2 is not persistent, although we can give an Elastic Ip, Using EC2 in front of a Load Balancer is tooooo much. What if it gets a million request? So to scale that EC2 you use another LB and an ASG>? This makes no sense

A is correct because a NLB can have an elastic IP and we can use this in our firewall as per the use case

upvoted 1 times

hajra313 1 week, 2 days ago

Setting up an EC2 instance with a public IP address to act as a proxy in front of the load balancer allows clients with restricted IP access to connect to the web service. The EC2 instance can handle IP address whitelisting and proxy requests to the ELB load balancer, ensuring that only authorized clients can access the service. This solution provides flexibility and control over access while leveraging the scalability and availability benefits of ELB.

upvoted 1 times

BillaRanga 3 days, 17 hours ago

Is this ChatGPT answer? Can you provide the AWS documentation link?

upvoted 1 times

Andy_09 2 weeks ago

Option C

upvoted 2 times

jaswantn 3 days, 20 hours ago

is there any valid justification for opting C? Glad to be informed, as these questions are tricky to answer.

upvoted 1 times

jaswantn 3 days, 20 hours ago

My inclination is for Option D, but not 100 % sure

upvoted 1 times

Question #745

A company has established a new AWS account. The account is newly provisioned and no changes have been made to the default settings. The company is concerned about the security of the AWS account root user.

What should be done to secure the root user?

A. Create IAM users for daily administrative tasks. Disable the root user.

B. Create IAM users for daily administrative tasks. Enable multi-factor authentication on the root user.

C. Generate an access key for the root user. Use the access key for daily administration tasks instead of the AWS Management Console.

D. Provide the root user credentials to the most senior solutions architect. Have the solutions architect use the root user for daily administration tasks.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 5 times

BillaRanga Most Recent 3 days, 17 hours ago

Selected Answer: B

"As a best practice, do not use the AWS account root user for any task where it's not required. Instead, create a new IAM user for each person that requires administrator access."
It's B :)

upvoted 3 times

Question #746

A company is deploying an application that processes streaming data in near-real time. The company plans to use Amazon EC2 instances for the workload. The network architecture must be configurable to provide the lowest possible latency between nodes.

Which combination of network solutions will meet these requirements? (Choose two.)

A. Enable and configure enhanced networking on each EC2 instance.

B. Group the EC2 instances in separate accounts.

C. Run the EC2 instances in a cluster placement group.

D. Attach multiple elastic network interfaces to each EC2 instance.

E. Use Amazon Elastic Block Store (Amazon EBS) optimized instance types.

Correct Answer: *AC*

Community vote distribution

AC (100%)

osmk 2 days, 19 hours ago

what's AM?

upvoted 1 times

jaswantn 1 week, 2 days ago

option C & E.
Option A is not viable as ..EC2 provides enhanced networking capabilities using single root I/O virtualization (SR-IOV) only on supported instance types.

upvoted 1 times

jaswantn 1 week, 2 days ago

option E... EBS–optimized instance uses an optimized configuration

upvoted 1 times

mestule 1 week, 6 days ago

Selected Answer: AC

A. Enable and configure enhanced networking on each EC2 instance. Enhanced networking provides higher bandwidth, higher packet per second (PPS) performance, and consistently lower inter-instance latencies.

C. Run the EC2 instances in a cluster placement group. A cluster placement group is a logical grouping of instances within a single Availability Zone. This configuration is recommended for applications that need low network latency, high network throughput, or both.

upvoted 4 times

Andy_09 2 weeks ago

Correct option should be CD

upvoted 1 times

Question #747

A financial services company wants to shut down two data centers and migrate more than 100 TB of data to AWS. The data has an intricate directory structure with millions of small files stored in deep hierarchies of subfolders. Most of the data is unstructured, and the company's file storage consists of SMB-based storage types from multiple vendors. The company does not want to change its applications to access the data after migration.

What should a solutions architect do to meet these requirements with the LEAST operational overhead?

A. Use AWS Direct Connect to migrate the data to Amazon S3.

B. Use AWS DataSync to migrate the data to Amazon FSx for Lustre.

C. Use AWS DataSync to migrate the data to Amazon FSx for Windows File Server.

D. Use AWS Direct Connect to migrate the data on-premises file storage to an AWS Storage Gateway volume gateway.

Correct Answer: *C*

Community vote distribution

C (100%)

Andy_09 Highly Voted 2 weeks ago

Option C

upvoted 5 times

ogerber Most Recent 2 days, 3 hours ago

Selected Answer: C

Option C since its SMB (windows) , and low operational effort so DataSync over Direct Connect

upvoted 1 times

osmk 2 days, 7 hours ago

Selected Answer: C

https://docs.aws.amazon.com/datasync/latest/userguide/create-fsx-location.html

upvoted 1 times

Question #748

A company uses an organization in AWS Organizations to manage AWS accounts that contain applications. The company sets up a dedicated monitoring member account in the organization. The company wants to query and visualize observability data across the accounts by using Amazon CloudWatch.

Which solution will meet these requirements?

A. Enable CloudWatch cross-account observability for the monitoring account. Deploy an AWS CloudFormation template provided by the monitoring account in each AWS account to share the data with the monitoring account.

B. Set up service control policies (SCPs) to provide access to CloudWatch in the monitoring account under the Organizations root organizational unit (OU).

C. Configure a new IAM user in the monitoring account. In each AWS account, configure an IAM policy to have access to query and visualize the CloudWatch data in the account. Attach the new IAM policy to the new IAM user.

D. Create a new IAM user in the monitoring account. Create cross-account IAM policies in each AWS account. Attach the IAM policies to the new IAM user.

Correct Answer: *A*

Community vote distribution

A (100%)

osmk 2 days, 7 hours ago

Selected Answer: A

https://docs.amazonaws.cn/en_us/AmazonCloudWatch/latest/monitoring/cloudwatch_crossaccount_dashboard.html

upvoted 1 times

jaswantn 3 days, 20 hours ago

option A
below are the links to check both parts of option A.
https://docs.amazonaws.cn/en_us/AmazonCloudWatch/latest/monitoring/cloudwatch_crossaccount_dashboard.html

https://docs.aws.amazon.com/AmazonCloudWatch/latest/monitoring/CloudWatch-Unified-Cross-Account-Setup.html#Unified-Cross-Account-SetupSource-SingleTemplate

upvoted 2 times

Andy_09 2 weeks ago

Option A

upvoted 2 times

Question #749

A company's website is used to sell products to the public. The site runs on Amazon EC2 instances in an Auto Scaling group behind an Application Load Balancer (ALB). There is also an Amazon CloudFront distribution, and AWS WAF is being used to protect against SQL injection attacks. The ALB is the origin for the CloudFront distribution. A recent review of security logs revealed an external malicious IP that needs to be blocked from accessing the website.

What should a solutions architect do to protect the application?

A. Modify the network ACL on the CloudFront distribution to add a deny rule for the malicious IP address.

B. Modify the configuration of AWS WAF to add an IP match condition to block the malicious IP address.

C. Modify the network ACL for the EC2 instances in the target groups behind the ALB to deny the malicious IP address.

D. Modify the security groups for the EC2 instances in the target groups behind the ALB to deny the malicious IP address.

Correct Answer: B

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 7 times

Question #750

A company sets up an organization in AWS Organizations that contains 10 AWS accounts. A solutions architect must design a solution to provide access to the accounts for several thousand employees. The company has an existing identity provider (IdP). The company wants to use the existing IdP for authentication to AWS.

Which solution will meet these requirements?

A. Create IAM users for the employees in the required AWS accounts. Connect IAM users to the existing IdP. Configure federated authentication for the IAM users.

B. Set up AWS account root users with user email addresses and passwords that are synchronized from the existing IdP.

C. Configure AWS IAM Identity Center (AWS Single Sign-On). Connect IAM Identity Center to the existing IdP. Provision users and groups from the existing IdP.

D. Use AWS Resource Access Manager (AWS RAM) to share access to the AWS accounts with the users in the existing IdP.

Correct Answer: *C*

Community vote distribution

C (100%)

ogerber 2 days, 3 hours ago

Selected Answer: C

Option C
https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html

upvoted 1 times

osmk 2 days, 7 hours ago

Selected Answer: C

https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html#provisioning-when-external-idp

upvoted 1 times

osmk 2 days, 10 hours ago

c--> Regardless of how you provision users, IAM Identity Center redirects the AWS Management Console, command line interface, and application authentication to your external IdP. IAM Identity Center then grants access to those resources based on policies you create in IAM Identity Center
https://docs.aws.amazon.com/singlesignon/latest/userguide/manage-your-identity-source-idp.html#provisioning-when-external-idp

upvoted 1 times

Andy_09 2 weeks ago

Option C

upvoted 2 times

Question #751

A solutions architect is designing an AWS Identity and Access Management (IAM) authorization model for a company's AWS account. The company has designated five specific employees to have full access to AWS services and resources in the AWS account.

The solutions architect has created an IAM user for each of the five designated employees and has created an IAM user group.

Which solution will meet these requirements?

A. Attach the AdministratorAccess resource-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

B. Attach the SystemAdministrator identity-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

C. Attach the AdministratorAccess identity-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

D. Attach the SystemAdministrator resource-based policy to the IAM user group. Place each of the five designated employee IAM users in the IAM user group.

Correct Answer: *C*

Community vote distribution

C (100%)

osmk 2 days, 9 hours ago

Selected Answer: C

C>>>https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_manage-attach-detach.html

upvoted 1 times

osmk 2 days, 10 hours ago

C>>>https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_manage-attach-detach.html

upvoted 1 times

Umuntu 1 week, 4 days ago

C looks correct

upvoted 2 times

Andy_09 1 week, 6 days ago

Option C

upvoted 3 times

Question #752

A company has a multi-tier payment processing application that is based on virtual machines (VMs). The communication between the tiers occurs asynchronously through a third-party middleware solution that guarantees exactly-once delivery.

The company needs a solution that requires the least amount of infrastructure management. The solution must guarantee exactly-once delivery for application messaging.

Which combination of actions will meet these requirements? (Choose two.)

A. Use AWS Lambda for the compute layers in the architecture.

B. Use Amazon EC2 instances for the compute layers in the architecture.

C. Use Amazon Simple Notification Service (Amazon SNS) as the messaging component between the compute layers.

D. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues as the messaging component between the compute layers.

E. Use containers that are based on Amazon Elastic Kubernetes Service (Amazon EKS) for the compute layers in the architecture.

Correct Answer: *AD*

Community vote distribution

AD (100%)

PolarFox 1 day ago

someone please explain why the combination of D and E is not the correct?

upvoted 1 times

osmk 2 days, 10 hours ago

Selected Answer: AD

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/FIFO-queues-exactly-once-processing.html

upvoted 1 times

jaswantn 3 days, 13 hours ago

option A for payment processing.
option D for exactly once delivery.

upvoted 1 times

Umuntu 1 week, 4 days ago

CD IS THE BEST ANSWER

upvoted 1 times

hajra313 1 week, 4 days ago

a and d

upvoted 2 times

Question #753

A company has a nightly batch processing routine that analyzes report files that an on-premises file system receives daily through SFTP. The company wants to move the solution to the AWS Cloud. The solution must be highly available and resilient. The solution also must minimize operational effort.

Which solution meets these requirements?

A. Deploy AWS Transfer for SFTP and an Amazon Elastic File System (Amazon EFS) file system for storage. Use an Amazon EC2 instance in an Auto Scaling group with a scheduled scaling policy to run the batch operation.

B. Deploy an Amazon EC2 instance that runs Linux and an SFTP service. Use an Amazon Elastic Block Store (Amazon EBS) volume for storage. Use an Auto Scaling group with the minimum number of instances and desired number of instances set to 1.

C. Deploy an Amazon EC2 instance that runs Linux and an SFTP service. Use an Amazon Elastic File System (Amazon EFS) file system for storage. Use an Auto Scaling group with the minimum number of instances and desired number of instances set to 1.

D. Deploy AWS Transfer for SFTP and an Amazon S3 bucket for storage. Modify the application to pull the batch files from Amazon S3 to an Amazon EC2 instance for processing. Use an EC2 instance in an Auto Scaling group with a scheduled scaling policy to run the batch operation.

Correct Answer: *D*

Community vote distribution

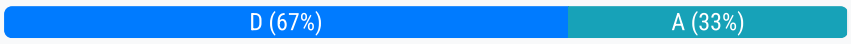

PolarFox 1 day ago

Selected Answer: D

trasnfer + S3 = HA, scheduled scaling = resilient

upvoted 1 times

NayeraB 1 day, 3 hours ago

Selected Answer: D

I'm not 100% sure, but D looks like the right flow to me

upvoted 1 times

osmk 2 days, 10 hours ago

Selected Answer: A

The service is designed to be highly scalable, highly available, and highly durable. Amazon EFS offers the following file system types to meet your availability and durability needs
->https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html
Amazon S3 achieves high availability by replicating data across multiple servers within AWS data centers->https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html

upvoted 1 times

NayeraB 1 day, 3 hours ago

But option A doesn't address the need for the application to pull the batch jobs from the new storage, also is the use of EFS needed here? In terms of it being a shared storage and whatnot..

upvoted 1 times

osmk 2 days, 10 hours ago

A>>>>>

upvoted 1 times

osmk 2 days, 10 hours ago

The service is designed to be highly scalable, highly available, and highly durable. Amazon EFS offers the following file system types to meet your availability and durability needs
->https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html
Amazon S3 achieves high availability by replicating data across multiple servers within AWS data centers->https://docs.aws.amazon.com/AmazonS3/latest/userguide/Welcome.html

upvoted 1 times

Andy_09 2 weeks ago

Option D

upvoted 2 times

Question #754

A company has users all around the world accessing its HTTP-based application deployed on Amazon EC2 instances in multiple AWS Regions. The company wants to improve the availability and performance of the application. The company also wants to protect the application against common web exploits that may affect availability, compromise security, or consume excessive resources. Static IP addresses are required.

What should a solutions architect recommend to accomplish this?

A. Put the EC2 instances behind Network Load Balancers (NLBs) in each Region. Deploy AWS WAF on the NLBs. Create an accelerator using AWS Global Accelerator and register the NLBs as endpoints.

B. Put the EC2 instances behind Application Load Balancers (ALBs) in each Region. Deploy AWS WAF on the ALBs. Create an accelerator using AWS Global Accelerator and register the ALBs as endpoints.

C. Put the EC2 instances behind Network Load Balancers (NLBs) in each Region. Deploy AWS WAF on the NLBs. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the NLBs.

D. Put the EC2 instances behind Application Load Balancers (ALBs) in each Region. Create an Amazon CloudFront distribution with an origin that uses Amazon Route 53 latency-based routing to route requests to the ALBs. Deploy AWS WAF on the CloudFront distribution.

Correct Answer: *B*

Community vote distribution

B (100%)

ogerber 2 days, 3 hours ago

Selected Answer: B

HTTP based application so ALB is required.
because static IP addresses are required, we should use global accelerator:
"By default, Global Accelerator provides you with static IP addresses that you associate with your accelerator."

upvoted 1 times

osmk 2 days, 9 hours ago

Selected Answer: B

Network Load Balancer (NLB): NLB operates at layer 4 and does not support AWS WAF directly
https://docs.aws.amazon.com/elasticloadbalancing/latest/application/introduction.html

upvoted 1 times

osmk 2 days, 9 hours ago

The company wants to improve the availability and performance of the application

upvoted 1 times

jaswantn 3 days, 13 hours ago

Static IP addresses are required, so option B....global accelerator with ALB

upvoted 1 times

Dhokal 1 week, 3 days ago

B is correct

upvoted 2 times

Andy_09 2 weeks ago

Option D

upvoted 4 times

Typewriter101 4 days, 4 hours ago

Why D cause i think global accelerator will do a better job an cloudfront to increase availability and performance

upvoted 1 times

Typewriter101 4 days, 4 hours ago

than cloudfront*

upvoted 1 times

Question #755

A company's data platform uses an Amazon Aurora MySQL database. The database has multiple read replicas and multiple DB instances across different Availability Zones. Users have recently reported errors from the database that indicate that there are too many connections. The company wants to reduce the failover time by 20% when a read replica is promoted to primary writer.

Which solution will meet this requirement?

A. Switch from Aurora to Amazon RDS with Multi-AZ cluster deployment.

B. Use Amazon RDS Proxy in front of the Aurora database.

C. Switch to Amazon DynamoDB with DynamoDB Accelerator (DAX) for read connections.

D. Switch to Amazon Redshift with relocation capability.

Correct Answer: *B*

Community vote distribution

B (100%)

osmk 2 days, 8 hours ago

Selected Answer: B

By using Amazon RDS Proxy, your applications can pool and share database connections. This pooling improves scalability by allowing multiple application instances to reuse existing connections.
It also makes your applications more resilient to database failures. When a primary database instance fails, RDS Proxy automatically connects to a standby DB instance while preserving application connections. =>https://docs.aws.amazon.com/AmazonRDS/latest/AuroraUserGuide/rds-proxy.html

upvoted 1 times

Umuntu 1 week, 4 days ago

Option B

upvoted 3 times

Andy_09 2 weeks ago

Option B

upvoted 3 times

Question #756

A company stores text files in Amazon S3. The text files include customer chat messages, date and time information, and customer personally identifiable information (PII).

The company needs a solution to provide samples of the conversations to an external service provider for quality control. The external service provider needs to randomly pick sample conversations up to the most recent conversation. The company must not share the customer PII with the external service provider. The solution must scale when the number of customer conversations increases.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Object Lambda Access Point. Create an AWS Lambda function that redacts the PII when the function reads the file. Instruct the external service provider to access the Object Lambda Access Point.

B. Create a batch process on an Amazon EC2 instance that regularly reads all new files, redacts the PII from the files, and writes the redacted files to a different S3 bucket. Instruct the external service provider to access the bucket that does not contain the PII.

B. Create a web application on an Amazon EC2 instance that presents a list of the files, redacts the PII from the files, and allows the external service provider to download new versions of the files that have the PII redacted.

D. Create an Amazon DynamoDB table. Create an AWS Lambda function that reads only the data in the files that does not contain PII. Configure the Lambda function to store the non-PII data in the DynamoDB table when a new file is written to Amazon S3. Grant the external service provider access to the DynamoDB table.

Correct Answer: *A*

Community vote distribution

A (100%)

osmk 2 days, 8 hours ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/tutorial-s3-object-lambda-redact-pii.html

upvoted 2 times

Vlad 1 week, 1 day ago

A is the correct choice.

upvoted 2 times

Umuntu 1 week, 4 days ago

A is the best choice

upvoted 2 times

Andy_09 2 weeks ago

Option A

upvoted 4 times

Question #757

A company is running a legacy system on an Amazon EC2 instance. The application code cannot be modified, and the system cannot run on more than one instance. A solutions architect must design a resilient solution that can improve the recovery time for the system.

What should the solutions architect recommend to meet these requirements?

A. Enable termination protection for the EC2 instance.

B. Configure the EC2 instance for Multi-AZ deployment.

C. Create an Amazon CloudWatch alarm to recover the EC2 instance in case of failure.

D. Launch the EC2 instance with two Amazon Elastic Block Store (Amazon EBS) volumes that use RAID configurations for storage redundancy.

Correct Answer: D

Community vote distribution

D (100%)

osmk 2 days, 8 hours ago

Selected Answer: D

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/raid-config.html

upvoted 1 times

Andy_09 1 week, 6 days ago

Option C

upvoted 4 times

Typewriter101 4 days, 3 hours ago

i think D is the answer.
Cause the question asks for a resilient solution and EBS with RAID config can balance between the performance and redundancy. EBS can also help with faster launch.

upvoted 2 times

Question #758

A company wants to deploy its containerized application workloads to a VPC across three Availability Zones. The company needs a solution that is highly available across Availability Zones. The solution must require minimal changes to the application.

Which solution will meet these requirements with the LEAST operational overhead?

A. Use Amazon Elastic Container Service (Amazon ECS). Configure Amazon ECS Service Auto Scaling to use target tracking scaling. Set the minimum capacity to 3. Set the task placement strategy type to spread with an Availability Zone attribute.

B. Use Amazon Elastic Kubernetes Service (Amazon EKS) self-managed nodes. Configure Application Auto Scaling to use target tracking scaling. Set the minimum capacity to 3.

C. Use Amazon EC2 Reserved Instances. Launch three EC2 instances in a spread placement group. Configure an Auto Scaling group to use target tracking scaling. Set the minimum capacity to 3.

D. Use an AWS Lambda function. Configure the Lambda function to connect to a VPC. Configure Application Auto Scaling to use Lambda as a scalable target. Set the minimum capacity to 3.

Correct Answer: *A*

Community vote distribution

A (100%)

osmk 2 days, 7 hours ago

Selected Answer: A

Amazon EKS self-managed nodes require you to manually install and configure the Kubernetes node components, such as kubelet, kube-proxy, and Docker, on your Amazon EC2 instances. You also need to manage the security group, IAM role, and subnet for your node group. Amazon ECS handles these tasks for you when you use the Amazon EC2 launch type .

upvoted 2 times

Andy_09 2 weeks ago

Option A

upvoted 2 times

Question #759

A media company stores movies in Amazon S3. Each movie is stored in a single video file that ranges from 1 GB to 10 GB in size.

The company must be able to provide the streaming content of a movie within 5 minutes of a user purchase. There is higher demand for movies that are less than 20 years old than for movies that are more than 20 years old. The company wants to minimize hosting service costs based on demand.

Which solution will meet these requirements?

A. Store all media content in Amazon S3. Use S3 Lifecycle policies to move media data into the Infrequent Access tier when the demand for a movie decreases.

B. Store newer movie video files in S3 Standard. Store older movie video files in S3 Standard-infrequent Access (S3 Standard-IA). When a user orders an older movie, retrieve the video file by using standard retrieval.

C. Store newer movie video files in S3 Intelligent-Tiering. Store older movie video files in S3 Glacier Flexible Retrieval. When a user orders an older movie, retrieve the video file by using expedited retrieval.

D. Store newer movie video files in S3 Standard. Store older movie video files in S3 Glacier Flexible Retrieval. When a user orders an older movie, retrieve the video file by using bulk retrieval.

Correct Answer: *C*

Community vote distribution

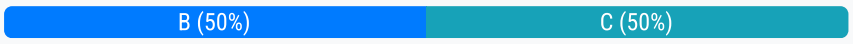

haci 2 days, 4 hours ago

Selected Answer: C

Expedited retrievals is typically made available within 1–5 minutes. Each unit of capacity provides that at least three Expedited retrievals can be performed every 5 minutes and provides up to 150 megabytes per second (MBps) of retrieval throughput.

There are some limitations but the bottom line is 5 minutes and I believe this leads us to Expedited retrievals.

https://docs.aws.amazon.com/amazonglacier/latest/dev/downloading-an-archive-two-steps.html#api-downloading-an-archive-two-steps-retrieval-expedited-capacity

upvoted 1 times

osmk 2 days, 7 hours ago

Selected Answer: B

S3 Standard-IA is for data that is accessed less frequently, but requires rapid access when needed. S3 Standard-IA offers the high durability, high throughput, and low latency of S3 Standard, with a low per GB storage price and per GB retrieval chargehttps://aws.amazon.com/s3/storage-classes/

upvoted 1 times

jaswantn 3 days, 11 hours ago

option B

upvoted 2 times

Freddie26 4 days, 11 hours ago

Technically, expedited retrieval for files is not guaranteed within 1-5 minutes for files larger than 250 MB+. See https://docs.aws.amazon.com/AmazonS3/latest/userguide/restoring-objects-retrieval-options.html.

upvoted 3 times

Andy_09 2 weeks ago

Option C

upvoted 3 times

Question #760

A solutions architect needs to design the architecture for an application that a vendor provides as a Docker container image. The container needs 50 GB of storage available for temporary files. The infrastructure must be serverless.

Which solution meets these requirements with the LEAST operational overhead?

A. Create an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume that has more than 50 GB of space.

B. Create an AWS Lambda function that uses the Docker container image with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space.

C. Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the AWS Fargate launch type. Create a task definition for the container image with an Amazon Elastic File System (Amazon EFS) volume. Create a service with that task definition.

D. Create an Amazon Elastic Container Service (Amazon ECS) cluster that uses the Amazon EC2 launch type with an Amazon Elastic Block Store (Amazon EBS) volume that has more than 50 GB of space. Create a task definition for the container image. Create a service with that task definition.

Correct Answer: *C*

Andy_09 2 weeks ago

Option C

upvoted 4 times

nj1999 1 week, 4 days ago

Why C and not B?

upvoted 1 times

hajra313 1 week, 2 days ago

the infrastructure must be serverless

upvoted 1 times

Cali182 1 week, 2 days ago

Creating an AWS Lambda function that uses the Docker container image with an Amazon S3 mounted volume might not be suitable because Lambda functions have limitations on execution duration (15 minutes) and storage size (maximum 512 MB in the /tmp directory).

upvoted 1 times

Question #761

A company needs to use its on-premises LDAP directory service to authenticate its users to the AWS Management Console. The directory service is not compatible with Security Assertion Markup Language (SAML).

Which solution meets these requirements?

A. Enable AWS IAM Identity Center (AWS Single Sign-On) between AWS and the on-premises LDAP.

B. Create an IAM policy that uses AWS credentials, and integrate the policy into LDAP.

C. Set up a process that rotates the IAM credentials whenever LDAP credentials are updated.

D. Develop an on-premises custom identity broker application or process that uses AWS Security Token Service (AWS STS) to get short-lived credentials.

Correct Answer: D

Community vote distribution

D (100%)

jaswantn 3 days, 11 hours ago

If your identity store is not compatible with SAML 2.0, then you can build a custom identity broker application to perform a similar function. .....option D

upvoted 1 times

kempes 1 week, 5 days ago

Selected Answer: D

The solution that best meets the requirements. This approach provides a pathway for authenticating LDAP users to AWS without requiring direct LDAP to AWS IAM Identity Center integration or SAML compatibility, offering a flexible and secure method to extend on-premises authentication mechanisms to AWS services.

upvoted 2 times

Question #762

A company stores multiple Amazon Machine Images (AMIs) in an AWS account to launch its Amazon EC2 instances. The AMIs contain critical data and configurations that are necessary for the company's operations. The company wants to implement a solution that will recover accidentally deleted AMIs quickly and efficiently.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create Amazon Elastic Block Store (Amazon EBS) snapshots of the AMIs. Store the snapshots in a separate AWS account.

B. Copy all AMIs to another AWS account periodically.

C. Create a retention rule in Recycle Bin.

D. Upload the AMIs to an Amazon S3 bucket that has Cross-Region Replication.

Correct Answer: C

Community vote distribution

C (100%)

Freddie26 4 days, 8 hours ago

Option C is correct. Recycling bin is a new feature to protect snaps and AMIs from accidental or malicious deleting. Inside the recycling bin, set a retention policy, and then your images or snapshots are protected.

upvoted 1 times

mestule 1 week, 5 days ago

Selected Answer: C

https://aws.amazon.com/about-aws/whats-new/2022/02/amazon-ec2-recycle-bin-machine-images/

upvoted 3 times

Andy_09 2 weeks ago

Option C

upvoted 3 times

Question #763

A company has 150 TB of archived image data stored on-premises that needs to be moved to the AWS Cloud within the next month. The company's current network connection allows up to 100 Mbps uploads for this purpose during the night only.

What is the MOST cost-effective mechanism to move this data and meet the migration deadline?

A. Use AWS Snowmobile to ship the data to AWS.

B. Order multiple AWS Snowball devices to ship the data to AWS.

C. Enable Amazon S3 Transfer Acceleration and securely upload the data.

D. Create an Amazon S3 VPC endpoint and establish a VPN to upload the data.

Correct Answer: B

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 11 times

Question #764

A company wants to migrate its three-tier application from on premises to AWS. The web tier and the application tier are running on third-party virtual machines (VMs). The database tier is running on MySQL.

The company needs to migrate the application by making the fewest possible changes to the architecture. The company also needs a database solution that can restore data to a specific point in time.

Which solution will meet these requirements with the LEAST operational overhead?

A. Migrate the web tier and the application tier to Amazon EC2 instances in private subnets. Migrate the database tier to Amazon RDS for MySQL in private subnets.

B. Migrate the web tier to Amazon EC2 instances in public subnets. Migrate the application tier to EC2 instances in private subnets. Migrate the database tier to Amazon Aurora MySQL in private subnets.

C. Migrate the web tier to Amazon EC2 instances in public subnets. Migrate the application tier to EC2 instances in private subnets. Migrate the database tier to Amazon RDS for MySQL in private subnets.

D. Migrate the web tier and the application tier to Amazon EC2 instances in public subnets. Migrate the database tier to Amazon Aurora MySQL in public subnets.

Correct Answer: *B*

Community vote distribution

B (100%)

haci 2 days, 3 hours ago

Selected Answer: B

I'm between B and C. Since RDS requires an additional configuration for PTR, it adds an operational overhead. So I will go with B.

Aurora provides automated backup and point-in-time recovery, simplifying backup management and data protection. Continuous incremental backups are taken automatically and stored in Amazon S3, and data retention periods can be specified to meet compliance requirements.

RDS provides the same but first, the users should set a retention period for these backups, allowing historical data recovery in case of accidental data loss or corruption, and point-in-time recovery (PITR) allows users to restore the database to any specific moment within the set retention period.

upvoted 1 times

hajra313 1 week, 4 days ago

C. This option aligns with the requirements by keeping the web tier in public subnets, migrating the application tier to EC2 instances in private subnets to enhance security, and using Amazon RDS for MySQL in private subnets to meet the database requirements with minimal operational overhead. option A:While migrating the web tier and application tier to EC2 instances in private subnets minimizes exposure to the internet. option B:. Migrating the database tier to Amazon Aurora MySQL introduces changes to the database engine, which might require additional testing and adjustments to the application. Additionally, Aurora MySQL does not directly support point-in-time recovery; instead, it uses continuous backups and snapshots for data recovery.

upvoted 2 times

Andy_09 1 week, 6 days ago

Option A works better

upvoted 2 times

Andy_09 2 weeks ago

Option B

upvoted 2 times

Question #765

A development team is collaborating with another company to create an integrated product. The other company needs to access an Amazon Simple Queue Service (Amazon SQS) queue that is contained in the development team's account. The other company wants to poll the queue without giving up its own account permissions to do so.

How should a solutions architect provide access to the SQS queue?

A. Create an instance profile that provides the other company access to the SQS queue.

B. Create an IAM policy that provides the other company access to the SQS queue.

C. Create an SQS access policy that provides the other company access to the SQS queue.

D. Create an Amazon Simple Notification Service (Amazon SNS) access policy that provides the other company access to the SQS queue.

Correct Answer: *C*

Community vote distribution

C (100%)

NayeraB 1 day, 4 hours ago

Selected Answer: C

https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/sqs-overview-of-managing-access.html

upvoted 1 times

hajra313 1 week, 4 days ago

option A: Instance profiles are used to grant permissions to EC2 instances, not for granting access to other AWS services like SQS queues. Option B.:AM policies are applied to IAM users, groups, or roles within the same AWS account. They are not directly applicable to granting access to resources in other AWS accounts. option C:SQS access policies allow you to grant cross-account access to SQS resources. You can specify the necessary permissions in the policy and attach it directly to the SQS queue. This way, you can give the other company's AWS account the necessary permissions to poll the queue without compromising their account permissions. option D. Amazon SNS access policies are used to manage access to SNS topics, not SQS queues

upvoted 2 times

kempes 1 week, 5 days ago

Selected Answer: C

Option C

upvoted 2 times

Andy_09 2 weeks ago

Option B

upvoted 1 times

Question #766

A company's developers want a secure way to gain SSH access on the company's Amazon EC2 instances that run the latest version of Amazon Linux. The developers work remotely and in the corporate office.

The company wants to use AWS services as a part of the solution. The EC2 instances are hosted in a VPC private subnet and access the internet through a NAT gateway that is deployed in a public subnet.

What should a solutions architect do to meet these requirements MOST cost-effectively?

A. Create a bastion host in the same subnet as the EC2 instances. Grant the ec2:CreateVpnConnection IAM permission to the developers. Install EC2 Instance Connect so that the developers can connect to the EC2 instances.

B. Create an AWS Site-to-Site VPN connection between the corporate network and the VPC. Instruct the developers to use the Site-to-Site VPN connection to access the EC2 instances when the developers are on the corporate network. Instruct the developers to set up another VPN connection for access when they work remotely.

C. Create a bastion host in the public subnet of the VPConfigure the security groups and SSH keys of the bastion host to only allow connections and SSH authentication from the developers' corporate and remote networks. Instruct the developers to connect through the bastion host by using SSH to reach the EC2 instances.

D. Attach the AmazonSSMManagedInstanceCore IAM policy to an IAM role that is associated with the EC2 instances. Instruct the developers to use AWS Systems Manager Session Manager to access the EC2 instances.

Correct Answer: *D*

Community vote distribution

D (100%)

kempes 1 week, 5 days ago

Selected Answer: D

Option D

upvoted 2 times

Andy_09 2 weeks ago

Option D

upvoted 4 times

Question #767

A pharmaceutical company is developing a new drug. The volume of data that the company generates has grown exponentially over the past few months. The company's researchers regularly require a subset of the entire dataset to be immediately available with minimal lag. However, the entire dataset does not need to be accessed on a daily basis. All the data currently resides in on-premises storage arrays, and the company wants to reduce ongoing capital expenses.

Which storage solution should a solutions architect recommend to meet these requirements?

A. Run AWS DataSync as a scheduled cron job to migrate the data to an Amazon S3 bucket on an ongoing basis.

B. Deploy an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance.

C. Deploy an AWS Storage Gateway volume gateway with cached volumes with an Amazon S3 bucket as the target storage. Migrate the data to the Storage Gateway appliance.

D. Configure an AWS Site-to-Site VPN connection from the on-premises environment to AWS. Migrate data to an Amazon Elastic File System (Amazon EFS) file system.

Correct Answer: *B*

Andy_09 Highly Voted 2 weeks ago

Option C

upvoted 7 times

hajra313 Most Recent 1 week, 4 days ago

B. Deploying an AWS Storage Gateway file gateway with an Amazon S3 bucket as the target storage would require the entire dataset to be stored in Amazon S3, which might not be cost-effective considering that only a subset of the data needs to be accessed regularly. Additionally, accessing data directly from S3 might introduce latency. so correct option is C bcz AWS Storage Gateway volume gateway with cached volumes allows the company to keep frequently accessed data locally on-premises while storing the entire dataset in Amazon S3. This solution provides immediate access to the subset of data with minimal lag, as frequently accessed data is cached locally. It also reduces ongoing capital expenses as it leverages Amazon S3 storage, which is cost-effective.

upvoted 4 times

Question #768

A company has a business-critical application that runs on Amazon EC2 instances. The application stores data in an Amazon DynamoDB table.

The company must be able to revert the table to any point within the last 24 hours.

Which solution meets these requirements with the LEAST operational overhead?

A. Configure point-in-time recovery for the table.

B. Use AWS Backup for the table.

C. Use an AWS Lambda function to make an on-demand backup of the table every hour.

D. Turn on streams on the table to capture a log of all changes to the table in the last 24 hours. Store a copy of the stream in an Amazon S3 bucket.

Correct Answer: *A*

Andy_09 Highly Voted 2 weeks ago

Option A

upvoted 7 times

Question #769

A company hosts an application used to upload files to an Amazon S3 bucket. Once uploaded, the files are processed to extract metadata, which takes less than 5 seconds. The volume and frequency of the uploads varies from a few files each hour to hundreds of concurrent uploads. The company has asked a solutions architect to design a cost-effective architecture that will meet these requirements.

What should the solutions architect recommend?

A. Configure AWS CloudTrail trails to log S3 API calls. Use AWS AppSync to process the files.

B. Configure an object-created event notification within the S3 bucket to invoke an AWS Lambda function to process the files.

C. Configure Amazon Kinesis Data Streams to process and send data to Amazon S3. Invoke an AWS Lambda function to process the files.

D. Configure an Amazon Simple Notification Service (Amazon SNS) topic to process the files uploaded to Amazon S3. Invoke an AWS Lambda function to process the files.

Correct Answer: *B*

Community vote distribution

B (100%)

jaswantn 3 days, 10 hours ago

option B

upvoted 1 times

hajra313 1 week, 4 days ago

option b bcz option c is WS AppSync is not the most appropriate solution for file processing.
option d While Amazon Simple Notification Service (SNS) can be used to trigger actions based on S3 events, it's not directly involved in processing files .option c :Kinesis is typically used for real-time data streaming and analytics, which may not be needed for simple file processing tasks such as extracting metadata.

upvoted 2 times

kempes 1 week, 5 days ago

Option D

upvoted 2 times

mestule 1 week, 5 days ago

Selected Answer: B

B seems to be make most sense to me.

upvoted 4 times

Andy_09 2 weeks ago

Option D

upvoted 1 times

Question #770

A company's application is deployed on Amazon EC2 instances and uses AWS Lambda functions for an event-driven architecture. The company uses nonproduction development environments in a different AWS account to test new features before the company deploys the features to production.

The production instances show constant usage because of customers in different time zones. The company uses nonproduction instances only during business hours on weekdays. The company does not use the nonproduction instances on the weekends. The company wants to optimize the costs to run its application on AWS.

Which solution will meet these requirements MOST cost-effectively?

A. Use On-Demand Instances for the production instances. Use Dedicated Hosts for the nonproduction instances on weekends only.

B. Use Reserved Instances for the production instances and the nonproduction instances. Shut down the nonproduction instances when not in use.

C. Use Compute Savings Plans for the production instances. Use On-Demand Instances for the nonproduction instances. Shut down the nonproduction instances when not in use.

D. Use Dedicated Hosts for the production instances. Use EC2 Instance Savings Plans for the nonproduction instances.

Correct Answer: C

Andy_09 Highly Voted 2 weeks ago

Option C

upvoted 7 times

Question #771

A company stores data in an on-premises Oracle relational database. The company needs to make the data available in Amazon Aurora PostgreSQL for analysis. The company uses an AWS Site-to-Site VPN connection to connect its on-premises network to AWS.

The company must capture the changes that occur to the source database during the migration to Aurora PostgreSQL.

Which solution will meet these requirements?

A. Use the AWS Schema Conversion Tool (AWS SCT) to convert the Oracle schema to Aurora PostgreSQL schema. Use the AWS Database Migration Service (AWS DMS) full-load migration task to migrate the data.

B. Use AWS DataSync to migrate the data to an Amazon S3 bucket. Import the S3 data to Aurora PostgreSQL by using the Aurora PostgreSQL aws_s3 extension.

C. Use the AWS Schema Conversion Tool (AWS SCT) to convert the Oracle schema to Aurora PostgreSQL schema. Use AWS Database Migration Service (AWS DMS) to migrate the existing data and replicate the ongoing changes.

D. Use an AWS Snowball device to migrate the data to an Amazon S3 bucket. Import the S3 data to Aurora PostgreSQL by using the Aurora PostgreSQL aws_s3 extension.

Correct Answer: *C*

Community vote distribution

C (100%)

kempes Highly Voted 1 week, 5 days ago

Selected Answer: C

Option C

upvoted 6 times

Andy_09 Highly Voted 2 weeks ago

Option C

upvoted 6 times

Question #772

A company built an application with Docker containers and needs to run the application in the AWS Cloud. The company wants to use a managed service to host the application.

The solution must scale in and out appropriately according to demand on the individual container services. The solution also must not result in additional operational overhead or infrastructure to manage.

Which solutions will meet these requirements? (Choose two.)

A. Use Amazon Elastic Container Service (Amazon ECS) with AWS Fargate.

B. Use Amazon Elastic Kubernetes Service (Amazon EKS) with AWS Fargate.

C. Provision an Amazon API Gateway API. Connect the API to AWS Lambda to run the containers.

D. Use Amazon Elastic Container Service (Amazon ECS) with Amazon EC2 worker nodes.

E. Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes.

Correct Answer: *AB*

NayeraB 1 day, 5 hours ago

Are people picking A&B as alternate solutions? Is the question asking for alternates?? Am I missing something? Somebody explain please I'm super confused.

upvoted 1 times

kempes 1 week, 5 days ago

Option AB

upvoted 1 times

Andy_09 2 weeks ago

Option AB

upvoted 2 times

Question #773

An ecommerce company is running a seasonal online sale. The company hosts its website on Amazon EC2 instances spanning multiple Availability Zones. The company wants its website to manage sudden traffic increases during the sale.

Which solution will meet these requirements MOST cost-effectively?

A. Create an Auto Scaling group that is large enough to handle peak traffic load. Stop half of the Amazon EC2 instances. Configure the Auto Scaling group to use the stopped instances to scale out when traffic increases.

B. Create an Auto Scaling group for the website. Set the minimum size of the Auto Scaling group so that it can handle high traffic volumes without the need to scale out.

C. Use Amazon CloudFront and Amazon ElastiCache to cache dynamic content with an Auto Scaling group set as the origin. Configure the Auto Scaling group with the instances necessary to populate CloudFront and ElastiCache. Scale in after the cache is fully populated.

D. Configure an Auto Scaling group to scale out as traffic increases. Create a launch template to start new instances from a preconfigured Amazon Machine Image (AMI).

Correct Answer: *D*

Andy_09 Highly Voted 2 weeks ago

Option D

upvoted 6 times

Question #774

A solutions architect must provide an automated solution for a company's compliance policy that states security groups cannot include a rule that allows SSH from 0.0.0.0/0. The company needs to be notified if there is any breach in the policy. A solution is needed as soon as possible.

What should the solutions architect do to meet these requirements with the LEAST operational overhead?

A. Write an AWS Lambda script that monitors security groups for SSH being open to 0.0.0.0/0 addresses and creates a notification every time it finds one.

B. Enable the restricted-ssh AWS Config managed rule and generate an Amazon Simple Notification Service (Amazon SNS) notification when a noncompliant rule is created.

C. Create an IAM role with permissions to globally open security groups and network ACLs. Create an Amazon Simple Notification Service (Amazon SNS) topic to generate a notification every time the role is assumed by a user.

D. Configure a service control policy (SCP) that prevents non-administrative users from creating or editing security groups. Create a notification in the ticketing system when a user requests a rule that needs administrator permissions.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 6 times

hajra313 Most Recent 1 week, 4 days ago

option b

upvoted 2 times

kempes 1 week, 5 days ago

Selected Answer: B

Option B

upvoted 3 times

Question #775

Use Amazon Elastic Kubernetes Service (Amazon EKS) with Amazon EC2 worker nodes.

A company has deployed an application in an AWS account. The application consists of microservices that run on AWS Lambda and Amazon Elastic Kubernetes Service (Amazon EKS). A separate team supports each microservice. The company has multiple AWS accounts and wants to give each team its own account for its microservices.

A solutions architect needs to design a solution that will provide service-to-service communication over HTTPS (port 443). The solution also must provide a service registry for service discovery.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Create an inspection VPC. Deploy an AWS Network Firewall firewall to the inspection VPC. Attach the inspection VPC to a new transit gateway. Route VPC-to-VPC traffic to the inspection VPC. Apply firewall rules to allow only HTTPS communication.

B. Create a VPC Lattice service network. Associate the microservices with the service network. Define HTTPS listeners for each service. Register microservice compute resources as targets. Identify VPCs that need to communicate with the services. Associate those VPCs with the service network.

C. Create a Network Load Balancer (NLB) with an HTTPS listener and target groups for each microservice. Create an AWS PrivateLink endpoint service for each microservice. Create an interface VPC endpoint in each VPC that needs to consume that microservice.

D. Create peering connections between VPCs that contain microservices. Create a prefix list for each service that requires a connection to a client. Create route tables to route traffic to the appropriate VPC. Create security groups to allow only HTTPS communication.

Correct Answer: *B*

Andy_09 2 weeks ago
Option B
upvoted 2 times

Question #776

A company has a mobile game that reads most of its metadata from an Amazon RDS DB instance. As the game increased in popularity, developers noticed slowdowns related to the game's metadata load times. Performance metrics indicate that simply scaling the database will not help. A solutions architect must explore all options that include capabilities for snapshots, replication, and sub-millisecond response times.

What should the solutions architect recommend to solve these issues?

A. Migrate the database to Amazon Aurora with Aurora Replicas.

B. Migrate the database to Amazon DynamoDB with global tables.

C. Add an Amazon ElastiCache for Redis layer in front of the database.

D. Add an Amazon ElastiCache for Memcached layer in front of the database.

Correct Answer: *C*

Andy_09 Highly Voted 1 week, 6 days ago
Option C is better as we need replication and snapshots
upvoted 8 times

Andy_09 Most Recent 2 weeks ago
Option D
upvoted 1 times

Question #777

A company uses AWS Organizations for its multi-account AWS setup. The security organizational unit (OU) of the company needs to share approved Amazon Machine Images (AMIs) with the development OU. The AMIs are created by using AWS Key Management Service (AWS KMS) encrypted snapshots.

Which solution will meet these requirements? (Choose two.)

A. Add the development team's OU Amazon Resource Name (ARN) to the launch permission list for the AMIs.

B. Add the Organizations root Amazon Resource Name (ARN) to the launch permission list for the AMIs.

C. Update the key policy to allow the development team's OU to use the AWS KMS keys that are used to decrypt the snapshots.

D. Add the development team’s account Amazon Resource Name (ARN) to the launch permission list for the AMIs.

E. Recreate the AWS KMS key. Add a key policy to allow the Organizations root Amazon Resource Name (ARN) to use the AWS KMS key.

Correct Answer: *AC*

Andy_09 Highly Voted 1 week, 6 days ago

Changing to options AC

upvoted 7 times

Andy_09 Most Recent 2 weeks ago

Option CD

upvoted 1 times

Question #778

A data analytics company has 80 offices that are distributed globally. Each office hosts 1 PB of data and has between 1 and 2 Gbps of internet bandwidth.

The company needs to perform a one-time migration of a large amount of data from its offices to Amazon S3. The company must complete the migration within 4 weeks.

Which solution will meet these requirements MOST cost-effectively?

A. Establish a new 10 Gbps AWS Direct Connect connection to each office. Transfer the data to Amazon S3.

B. Use multiple AWS Snowball Edge storage-optimized devices to store and transfer the data to Amazon S3.

C. Use an AWS Snowmobile to store and transfer the data to Amazon S3.

D. Set up an AWS Storage Gateway Volume Gateway to transfer the data to Amazon S3.

Correct Answer: *B*

Community vote distribution

B (100%)

mestule Highly Voted 1 week, 5 days ago

Selected Answer: B

B because too many offices that are geographically separated.

"data analytics company has 80 offices that are distributed globally."

upvoted 5 times

Andy_09 1 week, 4 days ago

Nice spot...completely missed that part !!

upvoted 1 times

chefKC Most Recent 1 week, 5 days ago

option B

upvoted 1 times

Andy_09 2 weeks ago

Option C looks good, as option B would lead to usage of too many snowball devices.

upvoted 2 times

Question #779

A company has an Amazon Elastic File System (Amazon EFS) file system that contains a reference dataset. The company has applications on Amazon EC2 instances that need to read the dataset. However, the applications must not be able to change the dataset. The company wants to use IAM access control to prevent the applications from being able to modify or delete the dataset.

Which solution will meet these requirements?

A. Mount the EFS file system in read-only mode from within the EC2 instances.

B. Create a resource policy for the EFS file system that denies the elasticfilesystem:ClientWrite action to the IAM roles that are attached to the EC2 instances.

C. Create an identity policy for the EFS file system that denies the elasticfilesystem:ClientWrite action on the EFS file system.

D. Create an EFS access point for each application. Use Portable Operating System Interface (POSIX) file permissions to allow read-only access to files in the root directory.

Correct Answer: *C*

Community vote distribution

C (100%)

Oo_Cc 1 week, 1 day ago

Selected Answer: C

"The company wasn't to use IAM access control". Yes, it would deny writing action to everything .. but it's still the only one that uses IAM.

upvoted 1 times

hajra313 1 week, 4 days ago

Create an EFS access point for each application. Use Portable Operating System Interface (POSIX) file permissions to allow read-only access to files in the root directory.

Explanation:
By creating an EFS access point for each application and configuring POSIX file permissions to allow read-only access, you can enforce the desired access control. This approach restricts write and delete actions on the dataset while allowing read access, aligning with the company's requirements.

upvoted 3 times

Andy_09 2 weeks ago

Option B

upvoted 3 times

Question #780

A company has hired an external vendor to perform work in the company's AWS account. The vendor uses an automated tool that is hosted in an AWS account that the vendor owns. The vendor does not have IAM access to the company's AWS account. The company needs to grant the vendor access to the company's AWS account.

Which solution will meet these requirements MOST securely?

A. Create an IAM role in the company's account to delegate access to the vendor's IAM role. Attach the appropriate IAM policies to the role for the permissions that the vendor requires.

B. Create an IAM user in the company's account with a password that meets the password complexity requirements. Attach the appropriate IAM policies to the user for the permissions that the vendor requires.

C. Create an IAM group in the company's account. Add the automated tool's IAM user from the vendor account to the group. Attach the appropriate IAM policies to the group for the permissions that the vendor requires.

D. Create an IAM user in the company's account that has a permission boundary that allows the vendor's account. Attach the appropriate IAM policies to the user for the permissions that the vendor requires.

Correct Answer: *A*

Andy_09 2 weeks ago

Option A looks ok

upvoted 3 times

Question #781

A company wants to run its experimental workloads in the AWS Cloud. The company has a budget for cloud spending. The company's CFO is concerned about cloud spending accountability for each department. The CFO wants to receive notification when the spending threshold reaches 60% of the budget.

Which solution will meet these requirements?

A. Use cost allocation tags on AWS resources to label owners. Create usage budgets in AWS Budgets. Add an alert threshold to receive notification when spending exceeds 60% of the budget.

B. Use AWS Cost Explorer forecasts to determine resource owners. Use AWS Cost Anomaly Detection to create alert threshold notifications when spending exceeds 60% of the budget.

C. Use cost allocation tags on AWS resources to label owners. Use AWS Support API on AWS Trusted Advisor to create alert threshold notifications when spending exceeds 60% of the budget.

D. Use AWS Cost Explorer forecasts to determine resource owners. Create usage budgets in AWS Budgets. Add an alert threshold to receive notification when spending exceeds 60% of the budget.

Correct Answer: *A*

Community vote distribution

A (100%)

NayeraB 1 day, 7 hours ago

Selected Answer: A

Nothing with cost explorer in it, and I don't want to be Captain Obvious but we need to set the budget alerts through AWS Budgets, so A

upvoted 1 times

Andy_09 2 weeks ago

Option A

upvoted 3 times

Question #782

A company wants to deploy an internal web application on AWS. The web application must be accessible only from the company's office. The company needs to download security patches for the web application from the internet.

The company has created a VPC and has configured an AWS Site-to-Site VPN connection to the company's office. A solutions architect must design a secure architecture for the web application.

Which solution will meet these requirements?

A. Deploy the web application on Amazon EC2 instances in public subnets behind a public Application Load Balancer (ALB). Attach an internet gateway to the VPC. Set the inbound source of the ALB's security group to 0.0.0.0/0.

B. Deploy the web application on Amazon EC2 instances in private subnets behind an internal Application Load Balancer (ALB). Deploy NAT gateways in public subnets. Attach an internet gateway to the VPC. Set the inbound source of the ALB's security group to the company's office network CIDR block.

C. Deploy the web application on Amazon EC2 instances in public subnets behind an internal Application Load Balancer (ALB). Deploy NAT gateways in private subnets. Attach an internet gateway to the VPSet the outbound destination of the ALB’s security group to the company's office network CIDR block.

D. Deploy the web application on Amazon EC2 instances in private subnets behind a public Application Load Balancer (ALB). Attach an internet gateway to the VPC. Set the outbound destination of the ALB’s security group to 0.0.0.0/0.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 6 times

NayeraB Most Recent 1 day, 7 hours ago

Selected Answer: B

B is well structured

upvoted 1 times

ogerber 2 days, 2 hours ago

To my opinion, with only having inbound of the companys CIDR block, it will not include access for the patches available online.
i would go for D

upvoted 1 times

kempes 1 week, 5 days ago

Selected Answer: B

Option B

upvoted 4 times

Question #783

A company maintains its accounting records in a custom application that runs on Amazon EC2 instances. The company needs to migrate the data to an AWS managed service for development and maintenance of the application data. The solution must require minimal operational support and provide immutable, cryptographically verifiable logs of data changes.

Which solution will meet these requirements MOST cost-effectively?

A. Copy the records from the application into an Amazon Redshift cluster.

B. Copy the records from the application into an Amazon Neptune cluster.

C. Copy the records from the application into an Amazon Timestream database.

D. Copy the records from the application into an Amazon Quantum Ledger Database (Amazon QLDB) ledger.

Correct Answer: *D*

Community vote distribution

Andy_09 Highly Voted 2 weeks ago

Option D

upvoted 5 times

NayeraB Most Recent 1 day, 7 hours ago

Selected Answer: D

D is correct

upvoted 1 times

Question #784

A company's marketing data is uploaded from multiple sources to an Amazon S3 bucket. A series of data preparation jobs aggregate the data for reporting. The data preparation jobs need to run at regular intervals in parallel. A few jobs need to run in a specific order later.

The company wants to remove the operational overhead of job error handling, retry logic, and state management.

Which solution will meet these requirements?

A. Use an AWS Lambda function to process the data as soon as the data is uploaded to the S3 bucket. Invoke other Lambda functions at regularly scheduled intervals.

B. Use Amazon Athena to process the data. Use Amazon EventBridge Scheduler to invoke Athena on a regular internal.

C. Use AWS Glue DataBrew to process the data. Use an AWS Step Functions state machine to run the DataBrew data preparation jobs.

D. Use AWS Data Pipeline to process the data. Schedule Data Pipeline to process the data once at midnight.

Correct Answer: *C*

Andy_09 2 weeks ago

Option C

upvoted 4 times

Question #785

A solutions architect is designing a payment processing application that runs on AWS Lambda in private subnets across multiple Availability Zones. The application uses multiple Lambda functions and processes millions of transactions each day.

The architecture must ensure that the application does not process duplicate payments.

Which solution will meet these requirements?

A. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon S3 bucket. Configure the S3 bucket with an event notification to invoke another Lambda function to process the due payments.

B. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon Simple Queue Service (Amazon SQS) queue. Configure another Lambda function to poll the SQS queue and to process the due payments.

C. Use Lambda to retrieve all due payments. Publish the due payments to an Amazon Simple Queue Service (Amazon SQS) FIFO queue. Configure another Lambda function to poll the FIFO queue and to process the due payments.

D. Use Lambda to retrieve all due payments. Store the due payments in an Amazon DynamoDB table. Configure streams on the DynamoDB table to invoke another Lambda function to process the due payments.

Correct Answer: *C*

Community vote distribution

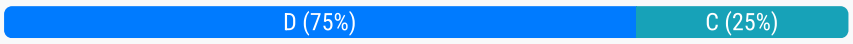

hajra313 Highly Voted 1 week, 6 days ago

Standard queues provide at-least-once delivery, which means that each message is delivered at least once.

FIFO queues provide exactly-once processing , which means that each message is delivered once and remains available until a consumer processes it and deletes it. Duplicates are not introduced into the queue. OPTION C

upvoted 8 times

Mikado211 Most Recent 14 hours, 42 minutes ago

SQS can have duplicate messages in case of problems with the timeout window.

upvoted 1 times

haci 1 day, 15 hours ago

Selected Answer: C

"The application does not process duplicate payments" is the key point, which leads us directly to SQS FIFO

upvoted 1 times

Cali182 1 week, 2 days ago

Selected Answer: D

Option D
DynamoDB Streams helps ensure the following:

Each stream record appears exactly once in the stream.

For each item that is modified in a DynamoDB table, the stream records appear in the same sequence as the actual modifications to the item.

DynamoDB Streams writes stream records in near-real time so that you can build applications that consume these streams and take action based on the contents.

upvoted 3 times

jaswantn 3 days, 9 hours ago

Option D...If you need to handle millions of transactions each day, you might need to consider other approach instead of SQS FIFO. And amongst the given options, we have DynmamoDB that maintains order in the streams.

upvoted 1 times

NayeraB 1 day, 6 hours ago

I'm not sure if the answer is DynamoDB as well, but answering your question, SQS Fifo can handle 300 messages/second without batching, 3,000 messages/second with batching. Assuming we're using the 300/sec option, with 86,400 seconds in a day, that gives you 25,920,000 messages, so in short, yes SQS can handle millions of requests each day.
Not to mention DynamoDB doesn't provide the exactly-once processing the SQS offer and clearly requested in the question. That's just my train of thought, I'm happy to be corrected.

upvoted 1 times

kempes 1 week, 5 days ago

Option c
upvoted 2 times

Andy_09 2 weeks ago
Option B
upvoted 1 times

Question #786

A company runs multiple workloads in its on-premises data center. The company's data center cannot scale fast enough to meet the company's expanding business needs. The company wants to collect usage and configuration data about the on-premises servers and workloads to plan a migration to AWS.

Which solution will meet these requirements?

A. Set the home AWS Region in AWS Migration Hub. Use AWS Systems Manager to collect data about the on-premises servers.

B. Set the home AWS Region in AWS Migration Hub. Use AWS Application Discovery Service to collect data about the on-premises servers.

C. Use the AWS Schema Conversion Tool (AWS SCT) to create the relevant templates. Use AWS Trusted Advisor to collect data about the on-premises servers.

D. Use the AWS Schema Conversion Tool (AWS SCT) to create the relevant templates. Use AWS Database Migration Service (AWS DMS) to collect data about the on-premises servers.

Correct Answer: *B*

Andy_09 2 weeks ago
Option B
upvoted 2 times

Question #787

A company has an organization in AWS Organizations that has all features enabled. The company requires that all API calls and logins in any existing or new AWS account must be audited. The company needs a managed solution to prevent additional work and to minimize costs. The company also needs to know when any AWS account is not compliant with the AWS Foundational Security Best Practices (FSBP) standard.

Which solution will meet these requirements with the LEAST operational overhead?

A. Deploy an AWS Control Tower environment in the Organizations management account. Enable AWS Security Hub and AWS Control Tower Account Factory in the environment.

B. Deploy an AWS Control Tower environment in a dedicated Organizations member account. Enable AWS Security Hub and AWS Control Tower Account Factory in the environment.

C. Use AWS Managed Services (AMS) Accelerate to build a multi-account landing zone (MALZ). Submit an RFC to self-service provision Amazon GuardDuty in the MALZ.

D. Use AWS Managed Services (AMS) Accelerate to build a multi-account landing zone (MALZ). Submit an RFC to self-service provision AWS Security Hub in the MALZ.

Correct Answer: *A*

Community vote distribution

A (100%)

kempes 1 week, 5 days ago

Selected Answer: A

Option A

upvoted 3 times

Andy_09 2 weeks ago

Option A

upvoted 2 times

Question #788

A company has stored 10 TB of log files in Apache Parquet format in an Amazon S3 bucket. The company occasionally needs to use SQL to analyze the log files.

Which solution will meet these requirements MOST cost-effectively?

A. Create an Amazon Aurora MySQL database. Migrate the data from the S3 bucket into Aurora by using AWS Database Migration Service (AWS DMS). Issue SQL statements to the Aurora database.

B. Create an Amazon Redshift cluster. Use Redshift Spectrum to run SQL statements directly on the data in the S3 bucket.

C. Create an AWS Glue crawler to store and retrieve table metadata from the S3 bucket. Use Amazon Athena to run SQL statements directly on the data in the S3 bucket.

D. Create an Amazon EMR cluster. Use Apache Spark SQL to run SQL statements directly on the data in the S3 bucket.

Correct Answer: *C*

Community vote distribution

C (100%)

kempes 1 week, 5 days ago

Selected Answer: C

Option C

upvoted 3 times

Andy_09 2 weeks ago

Option C

upvoted 3 times

Question #789

A company needs a solution to prevent AWS CloudFormation stacks from deploying AWS Identity and Access Management (IAM) resources that include an inline policy or “*” in the statement. The solution must also prohibit deployment of Amazon EC2 instances with public IP addresses. The company has AWS Control Tower enabled in its organization in AWS Organizations.

Which solution will meet these requirements?

A. Use AWS Control Tower proactive controls to block deployment of EC2 instances with public IP addresses and inline policies with elevated access or “*”.

B. Use AWS Control Tower detective controls to block deployment of EC2 instances with public IP addresses and inline policies with elevated access or “*”.

C. Use AWS Config to create rules for EC2 and IAM compliance. Configure the rules to run an AWS Systems Manager Session Manager automation to delete a resource when it is not compliant.

D. Use a service control policy (SCP) to block actions for the EC2 instances and IAM resources if the actions lead to noncompliance.

Correct Answer: *D*

Community vote distribution

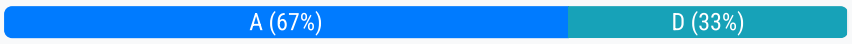

NayeraB 1 day, 6 hours ago

Selected Answer: A

A would provide a proactive solution, also I'm not sure if SCP are made for granular details like creation of EC2 instances with public IP addresses or IAM resources with certain inline policies.

upvoted 2 times

jaswantn 3 days, 8 hours ago

Selected Answer: D

Option D... This is preventive control of Control Tower where we use SCP to disallow actions that lead to policy violation.

upvoted 1 times

Andy_09 2 weeks ago

Option D

upvoted 1 times

Question #790

A company's web application that is hosted in the AWS Cloud recently increased in popularity. The web application currently exists on a single Amazon EC2 instance in a single public subnet. The web application has not been able to meet the demand of the increased web traffic.

The company needs a solution that will provide high availability and scalability to meet the increased user demand without rewriting the web application.

Which combination of steps will meet these requirements? (Choose two.)

A. Replace the EC2 instance with a larger compute optimized instance.

B. Configure Amazon EC2 Auto Scaling with multiple Availability Zones in private subnets.

C. Configure a NAT gateway in a public subnet to handle web requests.

D. Replace the EC2 instance with a larger memory optimized instance.

E. Configure an Application Load Balancer in a public subnet to distribute web traffic.

Correct Answer: *BE*

Community vote distribution

BE (100%)

chefKC 1 week, 5 days ago

Option B & E

upvoted 1 times

kempes 1 week, 5 days ago

Selected Answer: BE

Option BE

upvoted 3 times

Andy_09 2 weeks ago

Option BE

upvoted 3 times

Question #791

A company has AWS Lambda functions that use environment variables. The company does not want its developers to see environment variables in plaintext.

Which solution will meet these requirements?

A. Deploy code to Amazon EC2 instances instead of using Lambda functions.

B. Configure SSL encryption on the Lambda functions to use AWS CloudHSM to store and encrypt the environment variables.

C. Create a certificate in AWS Certificate Manager (ACM). Configure the Lambda functions to use the certificate to encrypt the environment variables.

D. Create an AWS Key Management Service (AWS KMS) key. Enable encryption helpers on the Lambda functions to use the KMS key to store and encrypt the environment variables.

Correct Answer: *D*

Andy_09 2 weeks ago

Option D

upvoted 3 times

Question #792

An analytics company uses Amazon VPC to run its multi-tier services. The company wants to use RESTful APIs to offer a web analytics service to millions of users. Users must be verified by using an authentication service to access the APIs.

Which solution will meet these requirements with the MOST operational efficiency?

A. Configure an Amazon Cognito user pool for user authentication. Implement Amazon API Gateway REST APIs with a Cognito authorizer.

B. Configure an Amazon Cognito identity pool for user authentication. Implement Amazon API Gateway HTTP APIs with a Cognito authorizer.

C. Configure an AWS Lambda function to handle user authentication. Implement Amazon API Gateway REST APIs with a Lambda authorizer.

D. Configure an IAM user to handle user authentication. Implement Amazon API Gateway HTTP APIs with an IAM authorizer.

Correct Answer: *A*

Community vote distribution

A (50%) B (50%)

stephensimudemy 21 hours, 8 minutes ago

Selected Answer: A

User pools is for Authentication and user management

upvoted 1 times

NayeraB 1 day, 15 hours ago

Selected Answer: B

B offers more operational efficiency imo

upvoted 1 times

chefKC 1 week, 5 days ago

Answer is A

upvoted 1 times

Andy_09 2 weeks ago

Option A

upvoted 4 times

Question #793

A company has a mobile app for customers. The app's data is sensitive and must be encrypted at rest. The company uses AWS Key Management Service (AWS KMS).

The company needs a solution that prevents the accidental deletion of KMS keys. The solution must use Amazon Simple Notification Service (Amazon SNS) to send an email notification to administrators when a user attempts to delete a KMS key.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create an Amazon EventBridge rule that reacts when a user tries to delete a KMS key. Configure an AWS Config rule that cancels any deletion of a KMS key. Add the AWS Config rule as a target of the EventBridge rule. Create an SNS topic that notifies the administrators.

B. Create an AWS Lambda function that has custom logic to prevent KMS key deletion. Create an Amazon CloudWatch alarm that is activated when a user tries to delete a KMS key. Create an Amazon EventBridge rule that invokes the Lambda function when the DeleteKey operation is performed. Create an SNS topic. Configure the EventBridge rule to publish an SNS message that notifies the administrators.

C. Create an Amazon EventBridge rule that reacts when the KMS DeleteKey operation is performed. Configure the rule to initiate an AWS Systems Manager Automation runbook. Configure the runbook to cancel the deletion of the KMS key. Create an SNS topic. Configure the EventBridge rule to publish an SNS message that notifies the administrators.

D. Create an AWS CloudTrail trail. Configure the trail to deliver logs to a new Amazon CloudWatch log group. Create a CloudWatch alarm based on the metric filter for the CloudWatch log group. Configure the alarm to use Amazon SNS to notify the administrators when the KMS DeleteKey operation is performed.

Correct Answer: *C*

Community vote distribution

C (100%)

NayeraB 1 day, 15 hours ago

Selected Answer: C

I agree with hajra313

upvoted 1 times

hajra313 1 week, 4 days ago

option c bcz Option C emerges as the clear winner due to its:

Direct event monitoring for the DeleteKey operation

Pre-built automation using Systems Manager Automation runbooks

Efficient notification via Amazon SNS

Minimal code development and operational overhead

Reduced risk of accidental deletion with faster response times

upvoted 2 times

Andy_09 2 weeks ago

Option C
https://docs.aws.amazon.com/prescriptive-guidance/latest/patterns/monitor-and-remediate-scheduled-deletion-of-aws-kms-keys.html

upvoted 3 times

Question #794

A company wants to analyze and generate reports to track the usage of its mobile app. The app is popular and has a global user base. The company uses a custom report building program to analyze application usage.

The program generates multiple reports during the last week of each month. The program takes less than 10 minutes to produce each report. The company rarely uses the program to generate reports outside of the last week of each month The company wants to generate reports in the least amount of time when the reports are requested.

Which solution will meet these requirements MOST cost-effectively?

A. Run the program by using Amazon EC2 On-Demand Instances. Create an Amazon EventBridge rule to start the EC2 instances when reports are requested. Run the EC2 instances continuously during the last week of each month.

B. Run the program in AWS Lambda. Create an Amazon EventBridge rule to run a Lambda function when reports are requested.

C. Run the program in Amazon Elastic Container Service (Amazon ECS). Schedule Amazon ECS to run the program when reports are requested.

D. Run the program by using Amazon EC2 Spot Instances. Create an Amazon EventBndge rule to start the EC2 instances when reports are requested. Run the EC2 instances continuously during the last week of each month.

Correct Answer: *B*

Community vote distribution

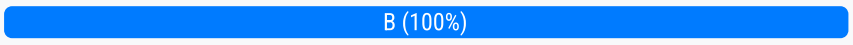

NayeraB 1 day, 7 hours ago

Selected Answer: B

B..maybe?

upvoted 1 times

Andy_09 2 weeks ago

Option B

upvoted 4 times

ogerber 2 days, 1 hour ago

not sure because it says that the program producues several reports , and each takes less than 10 min. i am voting for option A

upvoted 1 times

Question #795

A company is designing a tightly coupled high performance computing (HPC) environment in the AWS Cloud. The company needs to include features that will optimize the HPC environment for networking and storage.

Which combination of solutions will meet these requirements? (Choose two.)

A. Create an accelerator in AWS Global Accelerator. Configure custom routing for the accelerator.

B. Create an Amazon FSx for Lustre file system. Configure the file system with scratch storage.

C. Create an Amazon CloudFront distribution. Configure the viewer protocol policy to be HTTP and HTTPS.

D. Launch Amazon EC2 instances. Attach an Elastic Fabric Adapter (EFA) to the instances.

E. Create an AWS Elastic Beanstalk deployment to manage the environment.

Correct Answer: *BD*

Andy_09 2 weeks ago

Options BD

upvoted 4 times

Question #796

A company needs a solution to prevent photos with unwanted content from being uploaded to the company's web application. The solution must not involve training a machine learning (ML) model.

Which solution will meet these requirements?

A. Create and deploy a model by using Amazon SageMaker Autopilot. Create a real-time endpoint that the web application invokes when new photos are uploaded.

B. Create an AWS Lambda function that uses Amazon Rekognition to detect unwanted content. Create a Lambda function URL that the web application invokes when new photos are uploaded.

C. Create an Amazon CloudFront function that uses Amazon Comprehend to detect unwanted content. Associate the function with the web application.

D. Create an AWS Lambda function that uses Amazon Rekognition Video to detect unwanted content. Create a Lambda function URL that the web application invokes when new photos are uploaded.

Correct Answer: *B*

Community vote distribution

B (100%)

Andy_09 Highly Voted 2 weeks ago

Option B

upvoted 6 times

NayeraB Most Recent 1 day, 7 hours ago

Selected Answer: B

B is correct

upvoted 2 times

Question #797

A company uses AWS to run its ecommerce platform. The platform is critical to the company's operations and has a high volume of traffic and transactions. The company configures a multi-factor authentication (MFA) device to secure its AWS account root user credentials. The company wants to ensure that it will not lose access to the root user account if the MFA device is lost.

Which solution will meet these requirements?

A. Set up a backup administrator account that the company can use to log in if the company loses the MFA device.

B. Add multiple MFA devices for the root user account to handle the disaster scenario.

C. Create a new administrator account when the company cannot access the root account.

D. Attach the administrator policy to another IAM user when the company cannot access the root account.

Correct Answer: *B*

Community vote distribution

B (100%)

NayeraB 1 day, 7 hours ago

Selected Answer: B

I'd go for B

upvoted 2 times

hajra313 1 week, 6 days ago

B. Add multiple MFA devices for the root user account to handle the disaster scenario.

By adding multiple MFA devices for the root user account, the company ensures that it can still access the account even if one MFA device is lost. This approach provides a backup for authentication, addressing the concern of losing access to the root user account if the MFA device is lost.

upvoted 2 times

Andy_09 2 weeks ago

Option B

upvoted 3 times

Question #798

A social media company is creating a rewards program website for its users. The company gives users points when users create and upload videos to the website. Users redeem their points for gifts or discounts from the company's affiliated partners. A unique ID identifies users. The partners refer to this ID to verify user eligibility for rewards.

The partners want to receive notification of user IDs through an HTTP endpoint when the company gives users points. Hundreds of vendors are interested in becoming affiliated partners every day. The company wants to design an architecture that gives the website the ability to add partners rapidly in a scalable way.

Which solution will meet these requirements with the LEAST implementation effort?

A. Create an Amazon Timestream database to keep a list of affiliated partners. Implement an AWS Lambda function to read the list. Configure the Lambda function to send user IDs to each partner when the company gives users points.

B. Create an Amazon Simple Notification Service (Amazon SNS) topic. Choose an endpoint protocol. Subscribe the partners to the topic. Publish user IDs to the topic when the company gives users points.

C. Create an AWS Step Functions state machine. Create a task for every affiliated partner. Invoke the state machine with user IDs as input when the company gives users points.

D. Create a data stream in Amazon Kinesis Data Streams. Implement producer and consumer applications. Store a list of affiliated partners in the data stream. Send user IDs when the company gives users points.

Correct Answer: *B*

Community vote distribution

B (100%)

kempes Highly Voted 1 week, 5 days ago

Selected Answer: B

SNS is designed for precisely this kind of use case. It allows you to publish messages to a topic, which can then be delivered to multiple subscribers. Partners can subscribe to the SNS topic using an HTTP endpoint as the protocol, which meets the requirement to notify partners via an HTTP endpoint. This approach is highly scalable and requires the least implementation effort because it leverages managed services without the need for custom logic to manage subscriptions or deliver notifications.

upvoted 7 times

NayeraB Most Recent 1 day, 7 hours ago

Selected Answer: B

This is a perfect SNS use case

upvoted 2 times

jjcode 1 week, 3 days ago

The answer is B, create an SNS topic one subscriptions you can make is HTTP, This completely addresses the question objective.

upvoted 1 times

hajra313 1 week, 6 days ago

Option A involves creating an Amazon Timestream database to store affiliated partners and implementing an AWS Lambda function to read the list and send user IDs to each partner. While this approach can work, it involves more implementation effort than the Amazon SNS solution. It requires setting up and managing a database, as well as configuring the Lambda function to send notifications to partners. The Amazon SNS solution provides a simpler and more scalable approach for rapidly adding partners and notifying them when users receive points. so answer is B

upvoted 4 times

Andy_09 2 weeks ago

Option A

upvoted 1 times