Question #883

A company has deployed a multi-account strategy on AWS by using AWS Control Tower. The company has provided individual AWS accounts to each of its developers. The company wants to implement controls to limit AWS resource costs that the developers incur.

Which solution will meet these requirements with the LEAST operational overhead?

A. Instruct each developer to tag all their resources with a tag that has a key of CostCenter and a value of the developer's name. Use the required-tags AWS Config managed rule to check for the tag. Create an AWS Lambda function to terminate resources that do not have the tag. Configure AWS Cost Explorer to send a daily report to each developer to monitor their spending.

B. Use AWS Budgets to establish budgets for each developer account. Set up budget alerts for actual and forecast values to notify developers when they exceed or expect to exceed their assigned budget. Use AWS Budgets actions to apply a DenyAll policy to the developer's IAM role to prevent additional resources from being launched when the assigned budget is reached.

C. Use AWS Cost Explorer to monitor and report on costs for each developer account. Configure Cost Explorer to send a daily report to each developer to monitor their spending. Use AWS Cost Anomaly Detection to detect anomalous spending and provide alerts.

D. Use AWS Service Catalog to allow developers to launch resources within a limited cost range. Create AWS Lambda functions in each AWS account to stop running resources at the end of each work day. Configure the Lambda functions to resume the resources at the start of each work day.

Correct Answer: *C*

Community vote distribution

C (60%) B (40%)

f07ed8f 3 days, 1 hour ago

Selected Answer: C

Seem AWS Budgets does not have DenyAll function but only

Apply a custom Deny IAM policy that restricts the ability for a user, group, or role to provision additional Amazon EC2 resources

upvoted 1 times

BBR01 3 weeks, 1 day ago

Selected Answer: C

B and D are too aggressive.
A - "Instruct each developer", nope, too much operational work.

upvoted 2 times

sandordini 3 weeks, 3 days ago

Selected Answer: B

My first instinct says B, but Im concerned about the central management abilities of AWS Budgets. It seems that even though it is not planned to be used primarily to control other accounts its still possible:
"You can use actions to define an explicit response that you want to take when a budget exceeds its action threshold. You can trigger these alerts on actual or forecasted cost and usage budgets.
1. The management account sets the budget and threshold for the member account using budget filters.
2. When the budget threshold is breached, a budget action applies a restrictive SCP on the OU.

So hopefully B :D

upvoted 2 times

Question #884

A solutions architect is designing a three-tier web application. The architecture consists of an internet-facing Application Load Balancer (ALB) and a web tier that is hosted on Amazon EC2 instances in private subnets. The application tier with the business logic runs on EC2 instances in private subnets. The database tier consists of Microsoft SQL Server that runs on EC2 instances in private subnets. Security is a high priority for the company.

Which combination of security group configurations should the solutions architect use? (Choose three.)

A. Configure the security group for the web tier to allow inbound HTTPS traffic from the security group for the ALB.

B. Configure the security group for the web tier to allow outbound HTTPS traffic to 0.0.0.0/0.

C. Configure the security group for the database tier to allow inbound Microsoft SQL Server traffic from the security group for the application tier.

D. Configure the security group for the database tier to allow outbound HTTPS traffic and Microsoft SQL Server traffic to the security group for the web tier.

E. Configure the security group for the application tier to allow inbound HTTPS traffic from the security group for the web tier.

F. Configure the security group for the application tier to allow outbound HTTPS traffic and Microsoft SQL Server traffic to the security group for the web tier.

Correct Answer: *ACE*

Community vote distribution

ACE (100%)

sandordini 3 weeks, 3 days ago

Selected Answer: ACE

ALB >>HTTPS>> WEB tier >>HTTPS>> Application >>SQL traffic>> SQL DB

upvoted 2 times

Question #885

A company has released a new version of its production application. The company's workload uses Amazon EC2, AWS Lambda, AWS Fargate, and Amazon SageMaker.

The company wants to cost optimize the workload now that usage is at a steady state. The company wants to cover the most services with the fewest savings plans.

Which combination of savings plans will meet these requirements? (Choose two.)

A. Purchase an EC2 Instance Savings Plan for Amazon EC2 and SageMaker.

B. Purchase a Compute Savings Plan for Amazon EC2, Lambda, and SageMaker.

C. Purchase a SageMaker Savings Plan.

D. Purchase a Compute Savings Plan for Lambda, Fargate, and Amazon EC2.

E. Purchase an EC2 Instance Savings Plan for Amazon EC2 and Fargate.

Correct Answer: *CD*

Community vote distribution

CD (100%)

sandordini 3 weeks, 3 days ago

Selected Answer: CD

It's pretty obvious, although it's called: Machine Learning Savings Plans for Amazon SageMaker (C)
For the compute workloads we need a compute savings plan, that covers all the 3 compute options we use here (EC2, Lambda and Fargate) (D)

upvoted 2 times

Question #886

A company uses a Microsoft SQL Server database. The company's applications are connected to the database. The company wants to migrate to an Amazon Aurora PostgreSQL database with minimal changes to the application code.

Which combination of steps will meet these requirements? (Choose two.)

A. Use the AWS Schema Conversion Tool (AWS SCT) to rewrite the SQL queries in the applications.

B. Enable Babelfish on Aurora PostgreSQL to run the SQL queries from the applications.

C. Migrate the database schema and data by using the AWS Schema Conversion Tool (AWS SCT) and AWS Database Migration Service (AWS DMS).

D. Use Amazon RDS Proxy to connect the applications to Aurora PostgreSQL.

E. Use AWS Database Migration Service (AWS DMS) to rewrite the SQL queries in the applications.

Correct Answer: *BC*

Community vote distribution

BC (100%)

pranavsharma1604 1 week, 2 days ago

Selected Answer: BC

https://aws.amazon.com/rds/aurora/babelfish/

upvoted 1 times

sandordini 3 weeks, 3 days ago

Selected Answer: BC

B: Babelfish for Aurora PostgreSQL is a new capability for Amazon Aurora PostgreSQL-Compatible Edition that enables Aurora to understand commands from applications written for Microsoft SQL Server.
C: Is just obvious: Use Data Migration Tool for the migration, Schema Conversion tool for the Schema conversion.

upvoted 3 times

pranavsharma1604 1 week, 2 days ago

https://aws.amazon.com/rds/aurora/babelfish/

upvoted 1 times

Question #887

A company plans to rehost an application to Amazon EC2 instances that use Amazon Elastic Block Store (Amazon EBS) as the attached storage.

A solutions architect must design a solution to ensure that all newly created Amazon EBS volumes are encrypted by default. The solution must also prevent the creation of unencrypted EBS volumes.

Which solution will meet these requirements?

A. Configure the EC2 account attributes to always encrypt new EBS volumes.

B. Use AWS Config. Configure the encrypted-volumes identifier. Apply the default AWS Key Management Service (AWS KMS) key.

C. Configure AWS Systems Manager to create encrypted copies of the EBS volumes. Reconfigure the EC2 instances to use the encrypted volumes.

D. Create a customer managed key in AWS Key Management Service (AWS KMS). Configure AWS Migration Hub to use the key when the company migrates workloads.

Correct Answer: *B*

Community vote distribution

B (100%)

0bdf3af 5 days, 21 hours ago

A. https://repost.aws/knowledge-center/ebs-automatic-encryption

upvoted 2 times

lsomas 1 week, 5 days ago

Selected Answer: B

As it needs to prevent creation of Unencrypted EBS volume

upvoted 1 times

viejito 2 weeks ago

B es correcto , AWS Config para identificar automáticamente los volúmenes de EBS no cifrados y aplicar una acción correctiva.A,C,D : incorrectas , no cumplen con el cifrado automático

upvoted 2 times

Question #888

An ecommerce company wants to collect user clickstream data from the company's website for real-time analysis. The website experiences fluctuating traffic patterns throughout the day. The company needs a scalable solution that can adapt to varying levels of traffic.

Which solution will meet these requirements?

A. Use a data stream in Amazon Kinesis Data Streams in on-demand mode to capture the clickstream data. Use AWS Lambda to process the data in real time.

B. Use Amazon Kinesis Data Firehose to capture the clickstream data. Use AWS Glue to process the data in real time.

C. Use Amazon Kinesis Video Streams to capture the clickstream data. Use AWS Glue to process the data in real time.

D. Use Amazon Managed Service for Apache Flink (previously known as Amazon Kinesis Data Analytics) to capture the clickstream data. Use AWS Lambda to process the data in real time.

Correct Answer: *A*

Community vote distribution

f07ed8f 1 day, 1 hour ago

Selected Answer: B

Both Kinesis Data Streams and Firehose are scalable but Firehose offers automated scaling. I vote fore B

upvoted 1 times

sandordini 3 weeks, 3 days ago

Selected Answer: A

I think Apache Flink (previously known as Amazon Kinesis Data Analytics) would also be fine, but as here it wants to combine it with Lambda, I would rather opt for Kinesis Data Streams + Lambda, so A, because of the figure on this page:
https://aws.amazon.com/kinesis/

upvoted 1 times

Question #889

A global company runs its workloads on AWS. The company's application uses Amazon S3 buckets across AWS Regions for sensitive data storage and analysis. The company stores millions of objects in multiple S3 buckets daily. The company wants to identify all S3 buckets that are not versioning-enabled.

Which solution will meet these requirements?

A. Set up an AWS CloudTrail event that has a rule to identify all S3 buckets that are not versioning-enabled across Regions.

B. Use Amazon S3 Storage Lens to identify all S3 buckets that are not versioning-enabled across Regions.

C. Enable IAM Access Analyzer for S3 to identify all S3 buckets that are not versioning-enabled across Regions.

D. Create an S3 Multi-Region Access Point to identify all S3 buckets that are not versioning-enabled across Regions.

Correct Answer: *B*

Community vote distribution

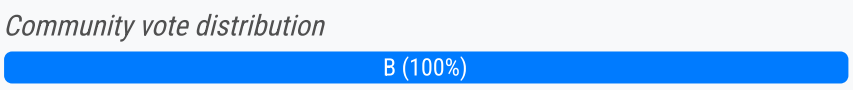

sandordini 3 weeks, 3 days ago

Selected Answer: B

You can use the Versioning-enabled bucket count metric to see which buckets use S3 Versioning. Then, you can take action in the S3 console to enable S3 Versioning for other buckets.

upvoted 1 times

Question #890

A company needs to optimize its Amazon S3 storage costs for an application that generates many files that cannot be recreated. Each file is approximately 5 MB and is stored in Amazon S3 Standard storage.

The company must store the files for 4 years before the files can be deleted. The files must be immediately accessible. The files are frequently accessed in the first 30 days of object creation, but they are rarely accessed after the first 30 days.

Which solution will meet these requirements MOST cost-effectively?

A. Create an S3 Lifecycle policy to move the files to S3 Glacier Instant Retrieval 30 days after object creation. Delete the files 4 years after object creation.

B. Create an S3 Lifecycle policy to move the files to S3 One Zone-Infrequent Access (S3 One Zone-IA) 30 days after object creation. Delete the files 4 years after object creation.

C. Create an S3 Lifecycle policy to move the files to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days after object creation. Delete the files 4 years after object creation.

D. Create an S3 Lifecycle policy to move the files to S3 Standard-Infrequent Access (S3 Standard-IA) 30 days after object creation. Move the files to S3 Glacier Flexible Retrieval 4 years after object creation.

Correct Answer: *A*

Community vote distribution

A (75%) C (25%)

th1002 1 week, 1 day ago

Selected Answer: C

why do we need one zone, glacier instant for 30 days ? or why do we need to move to glacier after 4 years ?
I think C is correct

upvoted 1 times

Karls 3 weeks, 3 days ago

B. Create an S3 Lifecycle policy to move the files to S3 One Zone-Infrequent Access (S3 One Zone-IA) 30 days after object creation. Delete the files 4 years after object creation.

This option leverages S3 One Zone-IA, which offers a lower cost compared to S3 Standard-IA, while ensuring that files are immediately accessible during the first 30 days of their creation. Then, after this period, the files are moved to S3 One Zone-IA for less frequent access storage, further reducing costs. Finally, the files are deleted after 4 years, meeting the requirement for long-term retention.

upvoted 2 times

sandordini 3 weeks, 3 days ago

Selected Answer: A

Although it's not stated what is meant by 'rarely accessed', this scenario would primarily be a candidate for the Glacier Instant Retrieval tier as the storage price would be more than 3 times lower compared to Standard IA. In the specific case of files being more frequently retrieved than quarterly, it can qualify for consideration of Standard IA.
Actually, we don't have the required info, so we have to guess what they are thinking.. which is pretty lame, to be honest..

upvoted 3 times

Question #891

A company runs its critical storage application in the AWS Cloud. The application uses Amazon S3 in two AWS Regions. The company wants the application to send remote user data to the nearest S3 bucket with no public network congestion. The company also wants the application to fail over with the least amount of management of Amazon S3.

Which solution will meet these requirements?

A. Implement an active-active design between the two Regions. Configure the application to use the regional S3 endpoints closest to the user.

B. Use an active-passive configuration with S3 Multi-Region Access Points. Create a global endpoint for each of the Regions.

C. Send user data to the regional S3 endpoints closest to the user. Configure an S3 cross-account replication rule to keep the S3 buckets synchronized.

D. Set up Amazon S3 to use Multi-Region Access Points in an active-active configuration with a single global endpoint. Configure S3 Cross-Region Replication.

Correct Answer: *D*

Community vote distribution

D (100%)

sandordini Highly Voted 3 weeks, 3 days ago

Selected Answer: D

Using a Multi-region Accesspoint in an Active-Active setup will send data to the closest Region, without accessing the internet: "send remote user data to the nearest S3 bucket with no public network congestion"
Not very easy to read and understand but it's all there: https://docs.aws.amazon.com/AmazonS3/latest/userguide/MultiRegionAccessPoints.html

upvoted 7 times

Scheldon Most Recent 23 hours, 53 minutes ago

Selected Answer: D

Answer D

https://docs.aws.amazon.com/AmazonS3/latest/userguide/MultiRegionAccessPoints.html

When you create a Multi-Region Access Point, you specify a set of AWS Regions where you want to store data to be served through that Multi-Region Access Point. You can use S3 Cross-Region Replication (CRR) to synchronize data among buckets in those Regions. You can then request or write data through the Multi-Region Access Point global endpoint. Amazon S3 automatically serves requests to the replicated dataset from the closest available Region. Multi-Region Access Points are also compatible with applications that are running in Amazon virtual private clouds (VPCs) including those that are using AWS PrivateLink for Amazon S3.

https://docs.aws.amazon.com/AmazonS3/latest/userguide/replication.html

upvoted 1 times

1223d0e 3 weeks, 4 days ago

To me it looks like C, the requirement is to send the request to the closest region

upvoted 1 times

Question #892

A company is migrating a data center from its on-premises location to AWS. The company has several legacy applications that are hosted on individual virtual servers. Changes to the application designs cannot be made.

Each individual virtual server currently runs as its own EC2 instance. A solutions architect needs to ensure that the applications are reliable and fault tolerant after migration to AWS. The applications will run on Amazon EC2 instances.

Which solution will meet these requirements?

A. Create an Auto Scaling group that has a minimum of one and a maximum of one. Create an Amazon Machine Image (AMI) of each application instance. Use the AMI to create EC2 instances in the Auto Scaling group Configure an Application Load Balancer in front of the Auto Scaling group.

B. Use AWS Backup to create an hourly backup of the EC2 instance that hosts each application. Store the backup in Amazon S3 in a separate Availability Zone. Configure a disaster recovery process to restore the EC2 instance for each application from its most recent backup.

C. Create an Amazon Machine Image (AMI) of each application instance. Launch two new EC2 instances from the AMI. Place each EC2 instance in a separate Availability Zone. Configure a Network Load Balancer that has the EC2 instances as targets.

D. Use AWS Mitigation Hub Refactor Spaces to migrate each application off the EC2 instance. Break down functionality from each application into individual components. Host each application on Amazon Elastic Container Service (Amazon ECS) with an AWS Fargate launch type.

Correct Answer: *C*

Community vote distribution

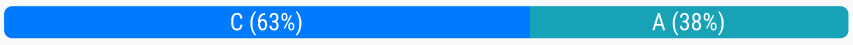

sandordini Highly Voted 3 weeks, 3 days ago

Selected Answer: C

NOT A: Autoscaling with Maximum of 1 EC2 :D
NOT B: Hourly backup... RPO 1hr
C: AMI, Multi-AZ -> Fault tolerant
NOT D: ECS with Fargate, but it needs to run on EC2..

upvoted 5 times

Scheldon Most Recent 1 day ago

Selected Answer: A

Answer A

It is possible to set Min and Max to 1 which will automatically bring up server when it will crash. Taking into consideration that we cannot change application design and load-balancing between regions would probably need that (no information if applications are statefull or stateless) i would go for solution in answer A
https://docs.aws.amazon.com/autoscaling/ec2/userguide/asg-capacity-limits.html

upvoted 1 times

Vasiliy 1 week, 1 day ago

Selected Answer: A

Autoscaling with max=1 is what is needed to keep only one instance at a time - it will still fail, but it will spawn exactly one instance in case of failure (we are not allowed to change the design of the app)
Having single instances in different AZ will not help - if one of the AZs is down, the app will still be affected

upvoted 2 times

Question #893

A company wants to isolate its workloads by creating an AWS account for each workload. The company needs a solution that centrally manages networking components for the workloads. The solution also must create accounts with automatic security controls (guardrails).

Which solution will meet these requirements with the LEAST operational overhead?

A. Use AWS Control Tower to deploy accounts. Create a networking account that has a VPC with private subnets and public subnets. Use AWS Resource Access Manager (AWS RAM) to share the subnets with the workload accounts.

B. Use AWS Organizations to deploy accounts. Create a networking account that has a VPC with private subnets and public subnets. Use AWS Resource Access Manager (AWS RAM) to share the subnets with the workload accounts.

C. Use AWS Control Tower to deploy accounts. Deploy a VPC in each workload account. Configure each VPC to route through an inspection VPC by using a transit gateway attachment.

D. Use AWS Organizations to deploy accounts. Deploy a VPC in each workload account. Configure each VPC to route through an inspection VPC by using a transit gateway attachment.

Correct Answer: *A*

Community vote distribution

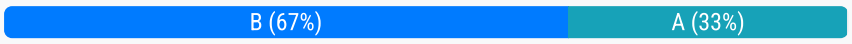

Tomrr 1 day, 17 hours ago

Selected Answer: A

Anser is A, Control Tower has guardrails

AWS Audit Manager provides an AWS Control Tower Guardrails framework to assist you with your audit preparation.

upvoted 1 times

Scheldon 1 day, 21 hours ago

Selected Answer: A

Taking into consideration that AWS Control Tower is Orchestrator for AWS Organization which applies guardrails, I think A is a good choose.

https://docs.aws.amazon.com/controltower/latest/userguide/what-is-control-tower.html

upvoted 1 times

sandordini 3 weeks, 3 days ago

Selected Answer: B

It's a hard one. I'd go for B
Several accounts in an org, with central mgmt > AWS Organization
Sharing resources among accounts > AWS RAM
AWS Organizations and RAM typically work well together...
Happy to be challenged, of course.

upvoted 4 times

sandordini 3 weeks, 3 days ago

Although automatic security control could be a hint for AWS Control Tower
(set up and operate your multi-account AWS environment with prescriptive controls)

upvoted 1 times

1223d0e 3 weeks, 4 days ago

Please explain why the answer is option A

upvoted 1 times

jackey_feng 1 week, 5 days ago

I prefer B which is free. A may cause fee for sevice used while I am not sure about it.

upvoted 1 times

Question #894

A company hosts a website on Amazon EC2 instances behind an Application Load Balancer (ALB). The website serves static content. Website traffic is increasing. The company wants to minimize the website hosting costs.

Which solution will meet these requirements?

A. Move the website to an Amazon S3 bucket. Configure an Amazon CloudFront distribution for the S3 bucket.

B. Move the website to an Amazon S3 bucket. Configure an Amazon ElastiCache cluster for the S3 bucket.

C. Move the website to AWS Amplify. Configure an ALB to resolve to the Amplify website.

D. Move the website to AWS Amplify. Configure EC2 instances to cache the website.

Correct Answer: *A*

Community vote distribution

A (100%)

trinh_le 3 weeks, 3 days ago

Selected Answer: A

static content -> S3

upvoted 1 times

Question #895

A company is implementing a shared storage solution for a media application that the company hosts on AWS. The company needs the ability to use SMB clients to access stored data.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Create an AWS Storage Gateway Volume Gateway. Create a file share that uses the required client protocol. Connect the application server to the file share.

B. Create an AWS Storage Gateway Tape Gateway. Configure tapes to use Amazon S3. Connect the application server to the Tape Gateway.

C. Create an Amazon EC2 Windows instance. Install and configure a Windows file share role on the instance. Connect the application server to the file share.

D. Create an Amazon FSx for Windows File Server file system. Connect the application server to the file system.

Correct Answer: *D*

Community vote distribution

D (100%)

trinh_le 3 weeks, 3 days ago

Selected Answer: D

SMB protocol -> FSx windows

upvoted 1 times

Question #896

A company is designing its production application's disaster recovery (DR) strategy. The application is backed by a MySQL database on an Amazon Aurora cluster in the us-east-1 Region. The company has chosen the us-west-1 Region as its DR Region.

The company's target recovery point objective (RPO) is 5 minutes and the target recovery time objective (RTO) is 20 minutes. The company wants to minimize configuration changes.

Which solution will meet these requirements with the MOST operational efficiency?

A. Create an Aurora read replica in us-west-1 similar in size to the production application's Aurora MySQL cluster writer instance.

B. Convert the Aurora cluster to an Aurora global database. Configure managed failover.

C. Create a new Aurora cluster in us-west-1 that has Cross-Region Replication.

D. Create a new Aurora cluster in us-west-1. Use AWS Database Migration Service (AWS DMS) to sync both clusters.

Correct Answer: *B*

Community vote distribution

B (100%)

sandordini 3 weeks, 3 days ago

Selected Answer: B

Aurora Global Database: allowing a single Amazon Aurora database to span multiple AWS Regions. It replicates your data with no impact on database performance, enables fast local reads with low latency in each Region, and provides disaster recovery from Region-wide outages.

upvoted 2 times

Question #897

A company runs a critical data analysis job each week before the first day of the work week. The job requires at least 1 hour to complete the analysis. The job is stateful and cannot tolerate interruptions. The company needs a solution to run the job on AWS.

Which solution will meet these requirements?

A. Create a container for the job. Schedule the job to run as an AWS Fargate task on an Amazon Elastic Container Service (Amazon ECS) cluster by using Amazon EventBridge Scheduler.

B. Configure the job to run in an AWS Lambda function. Create a scheduled rule in Amazon EventBridge to invoke the Lambda function.

C. Configure an Auto Scaling group of Amazon EC2 Spot Instances that run Amazon Linux. Configure a crontab entry on the instances to run the analysis.

D. Configure an AWS DataSync task to run the job. Configure a cron expression to run the task on a schedule.

Correct Answer: *A*

Question #898

A company runs workloads in the AWS Cloud. The company wants to centrally collect security data to assess security across the entire company and to improve workload protection.

Which solution will meet these requirements with the LEAST development effort?

A. Configure a data lake in AWS Lake Formation. Use AWS Glue crawlers to ingest the security data into the data lake.

B. Configure an AWS Lambda function to collect the security data in .csv format. Upload the data to an Amazon S3 bucket.

C. Configure a data lake in Amazon Security Lake to collect the security data. Upload the data to an Amazon S3 bucket.

D. Configure an AWS Database Migration Service (AWS DMS) replication instance to load the security data into an Amazon RDS cluster.

Correct Answer: *C*

Community vote distribution

C (100%)

sandordini Highly Voted 3 weeks, 3 days ago

Selected Answer: C

A, B, D are senseless +
Amazon Security Lake automatically centralizes security data from AWS environments, you can get a more complete understanding of your securit data across your entire organization. You can also improve the protection.

upvoted 5 times

Question #899

A company is migrating five on-premises applications to VPCs in the AWS Cloud. Each application is currently deployed in isolated virtual networks on premises and should be deployed similarly in the AWS Cloud. The applications need to reach a shared services VPC. All the applications must be able to communicate with each other.

If the migration is successful, the company will repeat the migration process for more than 100 applications.

Which solution will meet these requirements with the LEAST administrative overhead?

A. Deploy software VPN tunnels between the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets to the shared services VPC.

B. Deploy VPC peering connections between the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets to the shared services VPC through the peering connection.

C. Deploy an AWS Direct Connect connection between the application VPCs and the shared services VPAdd routes from the application VPCs in their subnets to the shared services VPC and the applications VPCs. Add routes from the shared services VPC subnets to the applications VPCs.

D. Deploy a transit gateway with associations between the transit gateway and the application VPCs and the shared services VPC. Add routes between the application VPCs in their subnets and the application VPCs to the shared services VPC through the transit gateway.

Correct Answer: *D*

0bdf3af 2 days, 2 hours ago

D. https://docs.aws.amazon.com/whitepapers/latest/building-scalable-secure-multi-vpc-network-infrastructure/transit-gateway.html

upvoted 1 times

Question #900

A company wants to use Amazon Elastic Container Service (Amazon ECS) to run its on-premises application in a hybrid environment. The application currently runs on containers on premises.

The company needs a single container solution that can scale in an on-premises, hybrid, or cloud environment. The company must run new application containers in the AWS Cloud and must use a load balancer for HTTP traffic.

Which combination of actions will meet these requirements? (Choose two.)

A. Set up an ECS cluster that uses the AWS Fargate launch type for the cloud application containers. Use an Amazon ECS Anywhere external launch type for the on-premises application containers.

B. Set up an Application Load Balancer for cloud ECS services.

C. Set up a Network Load Balancer for cloud ECS services.

D. Set up an ECS cluster that uses the AWS Fargate launch type. Use Fargate for the cloud application containers and the on-premises application containers.

E. Set up an ECS cluster that uses the Amazon EC2 launch type for the cloud application containers. Use Amazon ECS Anywhere with an AWS Fargate launch type for the on-premises application containers.

Correct Answer: *AB*

Scheldon 1 day ago

Answer: AB
We need to load-balance HTTP traffic hence Application Load Balancer is needed. Because Customer want to use container solution we need to us ECS with Fargate which will lunch cloud applications. To run on-premises applications in containers we need to use ECS Anyware.

https://docs.aws.amazon.com/AmazonECS/latest/developerguide/AWS_Fargate.html
https://docs.aws.amazon.com/AmazonECS/latest/developerguide/Welcome.html

upvoted 1 times

Scheldon 1 day ago

Amazon ECS services on AWS Fargate support the Application Load Balancer and Network Load Balancer load balancer types. Application Load Balancers are used to route HTTP/HTTPS (or layer 7) traffic.
Serverless (AWS Fargate (Fargate)) in the AWS cloud
Fargate is a serverless, pay-as-you-go compute engine. With Fargate you don't need to manage servers, handle capacity planning, or isolate container workloads for security.
On-premises virtual machines (VM) or servers
Amazon ECS Anywhere provides support for registering an external instance such as an on-premises server or virtual machine (VM), to your Amazon ECS cluster.

upvoted 1 times

0bdf3af 2 days, 2 hours ago

BD.

https://aws.amazon.com/blogs/aws/getting-started-with-amazon-ecs-anywhere-now-generally-available/

upvoted 1 times

Question #901

A company is migrating its workloads to AWS. The company has sensitive and critical data in on-premises relational databases that run on SQL Server instances.

The company wants to use the AWS Cloud to increase security and reduce operational overhead for the databases.

Which solution will meet these requirements?

A. Migrate the databases to Amazon EC2 instances. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.

B. Migrate the databases to a Multi-AZ Amazon RDS for SQL Server DB instance. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.

C. Migrate the data to an Amazon S3 bucket. Use Amazon Macie to ensure data security.

D. Migrate the databases to an Amazon DynamoDB table. Use Amazon CloudWatch Logs to ensure data security.

Correct Answer: *B*

Community vote distribution

B (100%)

trinh_le 3 weeks, 3 days ago

Selected Answer: B

Migrate the databases to a Multi-AZ Amazon RDS for SQL Server DB instance. Use an AWS Key Management Service (AWS KMS) AWS managed key for encryption.

upvoted 1 times

Question #902

A company wants to migrate an application to AWS. The company wants to increase the application's current availability. The company wants to use AWS WAF in the application's architecture.

Which solution will meet these requirements?

A. Create an Auto Scaling group that contains multiple Amazon EC2 instances that host the application across two Availability Zones. Configure an Application Load Balancer (ALB) and set the Auto Scaling group as the target. Connect a WAF to the ALB.

B. Create a cluster placement group that contains multiple Amazon EC2 instances that hosts the application. Configure an Application Load Balancer and set the EC2 instances as the targets. Connect a WAF to the placement group.

C. Create two Amazon EC2 instances that host the application across two Availability Zones. Configure the EC2 instances as the targets of an Application Load Balancer (ALB). Connect a WAF to the ALB.

D. Create an Auto Scaling group that contains multiple Amazon EC2 instances that host the application across two Availability Zones. Configure an Application Load Balancer (ALB) and set the Auto Scaling group as the target. Connect a WAF to the Auto Scaling group.

Correct Answer: *A*

Community vote distribution

A (100%)

sandordini 3 weeks, 3 days ago

A: EC2 - MultiAZ > ALB > WAF

upvoted 1 times

trinh_le 3 weeks, 3 days ago

Selected Answer: A

Not D because AWS WAF cannot be directly connected to an Auto Scaling Group, it should be associated with the ALB which managing the incoming web traffic

upvoted 1 times

Question #903

A company manages a data lake in an Amazon S3 bucket that numerous applications access. The S3 bucket contains a unique prefix for each application. The company wants to restrict each application to its specific prefix and to have granular control of the objects under each prefix.

Which solution will meet these requirements with the LEAST operational overhead?

A. Create dedicated S3 access points and access point policies for each application.

B. Create an S3 Batch Operations job to set the ACL permissions for each object in the S3 bucket.

C. Replicate the objects in the S3 bucket to new S3 buckets for each application. Create replication rules by prefix.

D. Replicate the objects in the S3 bucket to new S3 buckets for each application. Create dedicated S3 access points for each application.

Correct Answer: *B*

Community vote distribution

B (67%) A (33%)

Scheldon 16 hours, 6 minutes ago

Answer B

Taking into consideration that we have "numerous applications" (10,100,1000?) and we need meet requirements with the LEAST operational overhead I would go into authomatization of operations hence Batch Operations seems to be good choice.

https://aws.amazon.com/blogs/storage/updating-amazon-s3-object-acls-at-scale-with-s3-batch-operations/

upvoted 1 times

Tomrr 1 day, 16 hours ago

Selected Answer: A

https://docs.aws.amazon.com/AmazonS3/latest/userguide/access-points-policies.html

upvoted 1 times

trinh_le 3 weeks, 3 days ago

Selected Answer: B

Create an S3 Batch Operations job to set the ACL permissions for each object in the S3 bucket

upvoted 2 times

Question #904

A company has an application that customers use to upload images to an Amazon S3 bucket. Each night, the company launches an Amazon EC2 Spot Fleet that processes all the images that the company received that day. The processing for each image takes 2 minutes and requires 512 MB of memory.

A solutions architect needs to change the application to process the images when the images are uploaded.

Which change will meet these requirements MOST cost-effectively?

A. Use S3 Event Notifications to write a message with image details to an Amazon Simple Queue Service (Amazon SQS) queue. Configure an AWS Lambda function to read the messages from the queue and to process the images.

B. Use S3 Event Notifications to write a message with image details to an Amazon Simple Queue Service (Amazon SQS) queue. Configure an EC2 Reserved Instance to read the messages from the queue and to process the images.

C. Use S3 Event Notifications to publish a message with image details to an Amazon Simple Notification Service (Amazon SNS) topic. Configure a container instance in Amazon Elastic Container Service (Amazon ECS) to subscribe to the topic and to process the images.

D. Use S3 Event Notifications to publish a message with image details to an Amazon Simple Notification Service (Amazon SNS) topic. Configure an AWS Elastic Beanstalk application to subscribe to the topic and to process the images.

Correct Answer: *A*

Community vote distribution

A (100%)

Scheldon 15 hours, 37 minutes ago

Selected Answer: A

Answer A

I would go with Lambda and SQS.
when using SQS we will be sure that all images will be processed and hence to process we need 2 min and 512 MB of memory (Lambda is allowing upto 15 min and upto10K MB) Lambda should be perfect scalable solution which will allow for almost in real time image processing.

upvoted 1 times

Scheldon 15 hours, 37 minutes ago

and it is cost effective ;)

upvoted 1 times

trinh_le 3 weeks, 3 days ago

Selected Answer: A

less than 5 minutes -> use lambda

upvoted 2 times